


# COMP163

Database Management Systems

**Concurrency Control: Locking**  
**Chapter 22**

# Serializability Classes

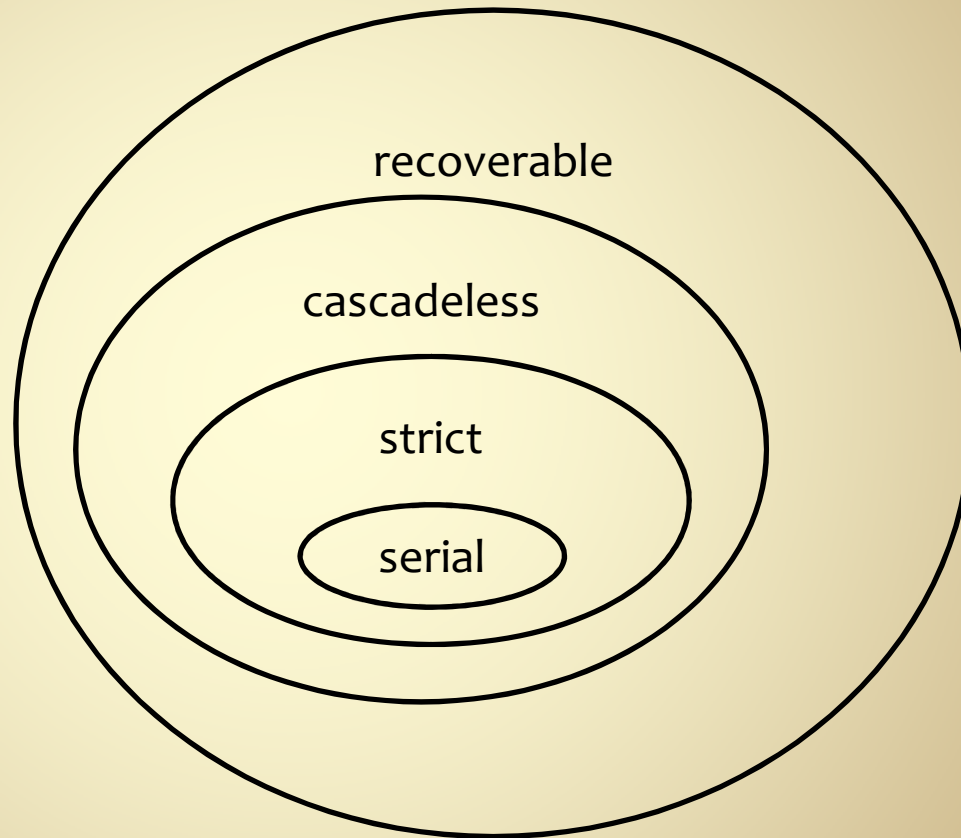
all possible schedules



REVIEW

# Recoverability Classes

all possible schedules



REVIEW

# Concurrency Control

- **Transaction theory** classifies possible schedules in terms of recoverability and correctness
- **Concurrency control** implements mechanisms to achieve specific policies
- **Pessimistic CC:**
  - prevent unwanted schedules from occurring
  - may reduce concurrency and stall transactions
- **Optimistic CC:**
  - allow any schedule
  - check at commit time and abort transactions contributing to unwanted schedules

REVIEW

**Locking**

# CC: Locking Protocols

- **Locking** is an operation that secures
  - permission to **read**, and/or permission to **write** a data item for a transaction
  - Example:
    - Lock(X): Data item X is locked on behalf of the requesting transaction
- **Unlocking** is an operation that removes these permissions from the data item.
  - Example:
    - Unlock(X): Data item X is made available to all other transactions
- Lock and Unlock are **atomic** operations
- Locking implements pessimistic CC

# SQL Isolation Levels

| isolation level  | prevents   | locking  |
|------------------|--|--|
| READ UNCOMMITTED |  | all locks released immediately following SQL statement execution             |
| READ COMMITTED   | dirty reads  | read locks released immediately<br>write locks held until end of transaction |
| REPEATABLE READ  | dirty reads<br>non-repeatable reads                    | all locks held until end of transaction (strict 2PL)                         |
| SERIALIZABLE     | dirty reads<br>non-repeatable reads<br>phantom records | requires index or table locks to prevent phantom reads                       |



# Recoverability Classes

- Recoverable schedule test:
  - *no transaction  $T$  commits until all transactions that wrote something that  $T$  reads have committed*
- Cascadeless test:
  - *every transaction only reads things written by committed transactions*
- Strict schedule test:
  - *no transaction can read or write anything that was written by an uncompleted transaction*



# Two Types of Locks

- Two locks modes:
  - shared (read)
  - exclusive (write)
- Shared lock:  $s(X)$ 
  - Multiple transactions can hold a shared lock on X
  - No exclusive lock can be applied on X while a shared lock is held on X
- Exclusive lock:  $x(X)$ 
  - Only one exclusive lock on X can exist at any time
  - No shared lock can be applied on X when an exclusive lock is held on X

# Lock Granting

locks held by other transactions

| requested<br>lock | none  | s(X)  | x(X) |
|-------------------|-------|-------|------|
| s(X)              | grant | grant | wait |
| x(X)              | grant | wait  | wait |

# Well-formed Transactions

- Locking assumes that all transactions are well-formed
- A transaction is well-formed if:
  - It properly locks a data item before it reads or writes to it
  - It does not lock an item locked by another transaction
  - It does not unlock an item that it does not hold a lock on
- More simply:  
Well-formed transactions obey locking rules

# Basic Lock/Unlock Algorithm

Lock(X) :

START:

if lock(X) = 0 then

lock(X)  $\leftarrow$  1

else

wait until (lock(X) = 0)

goto START

Unlock(X) :

lock(X)  $\leftarrow$  0 (\*unlock the item\*)

wake any transaction waiting for lock on X

# Shared-Lock Requests

START:

```
if lock(X) = "unlocked" then  
    lock(X) ← "shared-lock"  
    no_of_reads(X) ← 1
```

```
else if lock(X) = "shared-lock" then  
    no_of_reads(X) ← no_of_reads(X) + 1
```

```
else (* must be an exclusive lock *)  
    wait until (LOCK(X) = "unlocked")  
    go to START
```

# Exclusive-Lock Requests

START:

```
if lock(X) = "unlocked"
    lock(X) ← "exclusive-lock"
else
    wait until (lock(X) = "unlocked")
    go to START
```

# Unlocking

```
if LOCK(X) = "exclusive-lock"
    LOCK (X) ← "unlocked"
    wake up a waiting transactions (if any)

else if LOCK(X) ← "shared-lock"
    no_of_reads(X) ← no_of_reads(X)-1
    if no_of_reads(X) = 0
        LOCK(X) = "unlocked"
        wake up a waiting transactions (if any)
```



# Lock Conversions

- Lock upgrade: convert shared lock to exclusive lock

if T has the only shared lock on X

convert shared-lock(X) to exclusive-lock(X)

else

force T to wait until all other transactions unlock X

- Lock downgrade: convert exclusive lock to shared lock

if T has an exclusive-lock(X)

convert exclusive-lock(X) to shared-lock(X)

# Two Phase Locking

# Two-Phase Locking

- **Two Phases:**
  - Locking (Growing)
  - Unlocking (Shrinking)
- **Locking (Growing) Phase:**
  - A transaction applies locks (read or write) on desired data items one at a time
- **Unlocking (Shrinking) Phase:**
  - A transaction unlocks its locked data items one at a time
- **Requirement:**
  - Within any transaction these two phases must be mutually exclusive – once you start unlocking, you cannot request any more locks

# Two-Phase Locking

- Locking itself does not imply serializability
- 2PL guarantees serializability
  - improper ordering of operations is prevented
  - if 2PL is enforced, there is no need to test schedules for serializability
- 2PL limits concurrency
  - locks may need to be held longer than needed
- Basic 2PL may cause deadlock

# Locking Example

**T1**

|                |       |
|----------------|-------|
| read_lock (Y)  | s1(Y) |
| read_item (Y)  | r1(Y) |
| unlock (Y)     | u1(Y) |
| write_lock (X) | x1(X) |
| read_item (X)  | r1(X) |
| X:=X+Y         |       |
| write_item (X) | w1(X) |
| unlock (X)     | u1(X) |

**T2**

|                |       |
|----------------|-------|
| read_lock (X)  | s2(X) |
| read_item (X)  | r2(X) |
| unlock (X)     | u2(X) |
| write_lock (Y) | x2(Y) |
| read_item (Y)  | r2(Y) |
| Y:=X+Y         |       |
| write_item (Y) | w2(Y) |
| unlock (Y)     | u2(Y) |

Initial values: X=20; Y=30

Result of serial execution, T1 followed by T2: X=50, Y=80

Result of serial execution, T2 followed by T1: X=70, Y=50

# Locking Example

**T1**

|                |       |
|----------------|-------|
| read_lock (Y)  | s1(Y) |
| read_item (Y)  | r1(Y) |
| unlock (Y)     | u1(Y) |
| write_lock (X) | x1(X) |
| read_item (X)  | r1(X) |
| X:=X+Y         |       |
| write_item (X) | w1(X) |
| unlock (X)     | u1(X) |

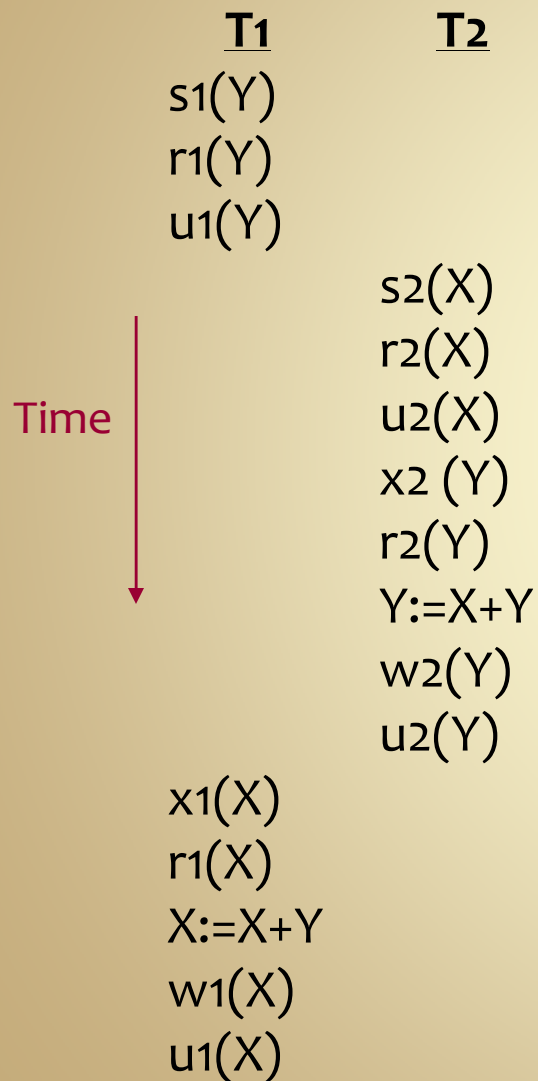
**T2**

|                |       |
|----------------|-------|
| read_lock (X)  | s2(X) |
| read_item (X)  | r2(X) |
| unlock (X)     | u2(X) |
| write_lock (Y) | x2(Y) |
| read_item (Y)  | r2(Y) |
| Y:=X+Y         |       |
| write_item (Y) | w2(Y) |
| unlock (Y)     | u2(Y) |

Both transactions obey basic locking protocols, since they hold appropriate locks when reading or writing data items.

Neither transaction obeys 2PL.

# Locking Example



Result: X=50; Y=50

This schedule is legal in that it obeys locking rules, but it is not serializable (and not correct)

violates two-phase policy



# 2PL Example

| <u>T1</u> | <u>T2</u> |                 |
|-----------|-----------|-----------------|
| s1(Y)     | s2(X)     |                 |
| r1 (Y)    | r2(X)     |                 |
| x1 (X)    | x2(Y)     | growing phase   |
| u1(Y)     | u2(X)     |                 |
| r1(X)     | r2(Y)     | shrinking phase |
| X:=X+Y    | Y:=X+Y    |                 |
| w1(X)     | w2(Y)     |                 |
| u1(X)     | u2(Y)     |                 |

Both transactions obey 2PL.

It is not possible to interleave them in a manner that results in a non-serializable schedule.

# Basic 2PL

- Basic 2PL requires that no locks are requested after the first unlock
- Guarantees serializability
  - transactions that request operations that violate serializability are delayed while waiting on locks
- Reduces concurrency, since locks must be held until all needed locks have been acquired
- May cause deadlock

# Conservative 2PL

- Conservative 2PL requires that all locks must be acquired at start of transaction
- Prevents deadlock, since all locks are acquired as a block
  - No transaction can be waiting on one lock while it holds another lock
- Further restricts concurrency, since transaction must request strongest lock that *might* be needed

# Strict 2PL

- Strict 2PL requires that all locks must be held until end of transaction
- Deadlock is possible
- Guarantees strict schedules
- May require holding locks longer than necessary
- Most commonly used algorithm

# Deadlock

# 2PL: Deadlock

T1

s1(Y)

r1 (Y)

x1 (X)

u1(Y)

r1(X)

X:=X+Y

w1(X)

u1(X)

T2

s2(X)

r2(X)

x2(Y)

u2(X)

r2(Y)

Y:=X+Y

w2(Y)

u2(Y)

T1 cannot proceed until T2 releases lock on X.  
T2 cannot proceed until T1 releases lock on Y.  
→ DEADLOCK

# Conservative 2PL: ~~Deadlock~~

T1

s1(Y), x1  
(X)

r1 (Y)

u1(Y)

r1(X)

X:=X+Y

w1(X)

u1(X)

T2

s2(X), x2(Y)

r2(X)

u2(X)

r2(Y)

Y:=X+Y

w2(Y)

u2(Y)

*In this example, the only possible schedules are serial schedules.*

Locks must be acquired as a unit at beginning of transaction.  
*Transaction cannot be holding locks while waiting on locks.*  
Deadlock is not possible.



# Conservative 2PL: ~~Deadlock~~

T1

$s1(Y), x1$   
 $(Z)$

$r1(Y)$

$u1(Y)$

$r1(Z)$

$Z := Z + Y$

$w1(Z)$

$u1(Z)$

T2

$s2(X), x2(Y)$

$r2(X)$

$u2(X)$

$r2(Y)$

$Y := X + Y$

$w2(Y)$

$u2(Y)$

T2 can proceed as soon as T1 releases lock on Y.

Concurrency is still possible under conservative 2PL.

# Deadlock: Detection/Resolution

- Let deadlocks happen, then resolve the problem

- **Wait-for graph**

- scheduler maintains a *wait-for graph*
  - arc from Tx to Ty indicates Tx is waiting for a lock held by Ty
- when a transaction is blocked, it is added to the graph
- a cycle in the wait-for-graph indicates deadlock
- one transaction involved in the cycle is selected (victim) and rolled-back

- **Timeout**

- abort any transaction that has been waiting for some set amount of time
- simple solution, but may be abort a transaction that could eventually proceed

# Deadlock: Prevention

- **Locking policy**

- Implement a CC policy that never allows deadlock to occur
- Example: conservative 2PL

- **Waits-for cycle avoidance**

- Use wait-for graph, but do not allow cycles to occur
- Example: any transaction that would create a cycle is aborted
- Other algorithms use timestamps to chose victim

# Deadlock Victim Selection

- T1 tries to lock X, T2 holds lock on X

- **wait-die:**  
if T1 is older than T2, T1 waits  
otherwise, T1 aborts
- **wound-wait:**  
if T1 is older than T2, T2 aborts,  
otherwise, T1 waits
- **no-waiting:**  
T1 aborts
- **cautious waiting:**  
if T2 is waiting, T1 aborts,  
otherwise T1 waits

# Starvation

- **Starvation**

- A particular transaction consistently waits or gets restarted and never gets a chance to complete
- Caused by deadlock victim selection policy
- Inherent in all priority based scheduling mechanisms
- Example: Wound-Wait  
a younger transaction may always be aborted by a long running older transaction,

# Livelock

- States of the processes involved are changing, but no process is progressing
- example: two people meet in a narrow corridor, and each tries to be polite by moving aside to let the other pass, but they end up swaying from side to side without making any progress
- Livelock can result from deadlock detection/recovery  
If more than one process takes action, the deadlock detection algorithm can be repeatedly triggered.
  - avoided by ensuring that only one process takes action

source: Wikipedia





# SQL Isolation Levels

| isolation level  | prevents   | locking  |
|------------------|--|--|
| READ UNCOMMITTED |  | all locks released immediately following SQL statement execution             |
| READ COMMITTED   | dirty reads  | read locks released immediately<br>write locks held until end of transaction |
| REPEATABLE READ  | dirty reads<br>non-repeatable reads                    | all locks held until end of transaction (strict 2PL)                         |
| SERIALIZABLE     | dirty reads<br>non-repeatable reads<br>phantom records | requires index or table locks to prevent phantom reads                       |