

COMP163

Database Management Systems

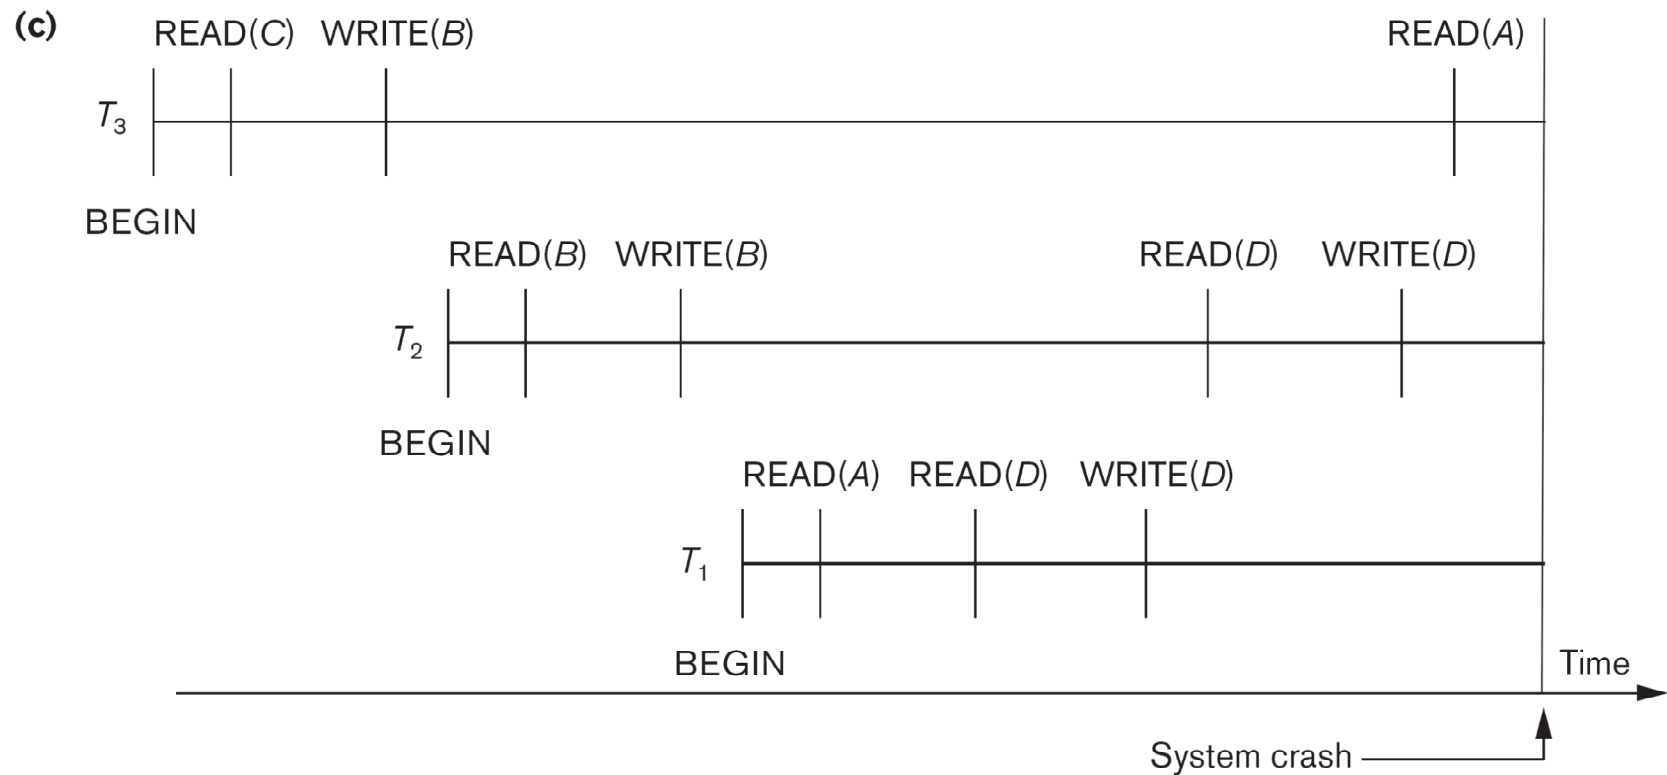
Database Recovery Techniques
Chapter 23

Discussion Notes

- We started the discussion with notes on recover in Oracle and SQL Server
 - SQL Server:
<http://msdn.microsoft.com/en-us/library/ms189275.aspx>
 - Oracle:
http://docs.oracle.com/cd/B10501_01/server.920/a96519/strategy.htm

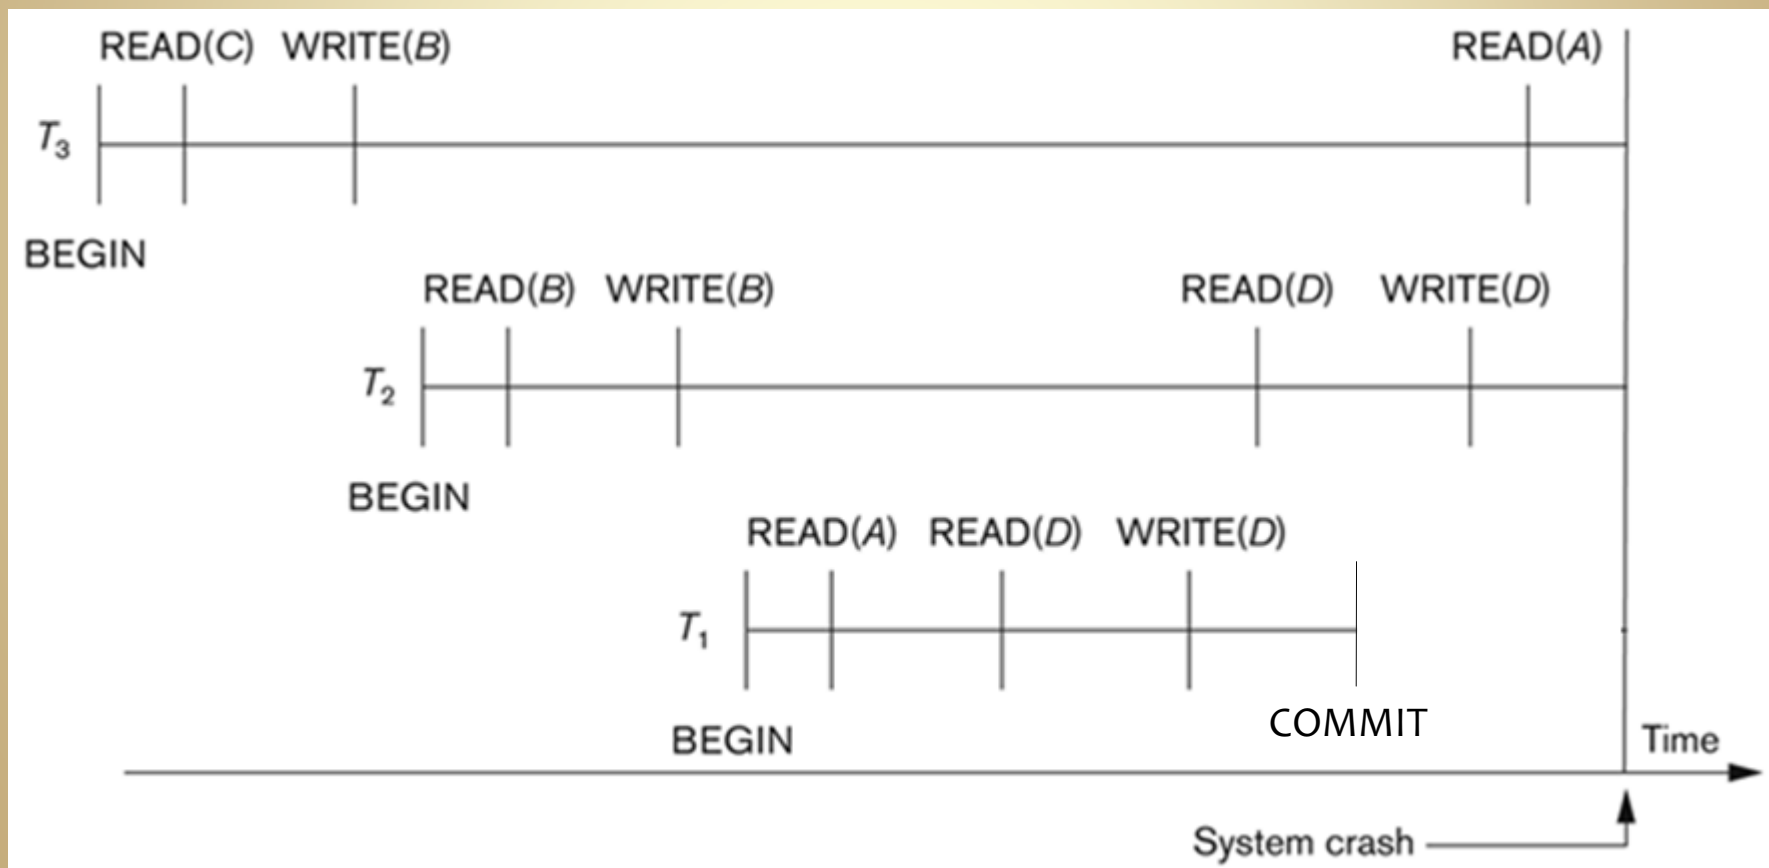
Recovery Example

Three concurrent transactions and timeline before system crash



Recovery Example

Three concurrent transactions and timeline before system crash



Purpose of Database Recovery

- Bring the database into the ***most recent consistent state*** that existed prior to a failure.
- Preserve transaction properties
 - Atomicity and Durability
- Example:
 - bank database crashes before a fund transfer transaction completes
 - either one or both accounts may have incorrect values
 - database must be restored to the state before the transaction modified any of the accounts

Types of Failure

The database may become unavailable due to

- **Transaction failure:**
Transactions may fail because of incorrect input, deadlock, incorrect synchronization
- **System failure:**
System may fail because of addressing error, application error, operating system fault, RAM failure, ...
- **Media failure:**
Disk head crash, power disruption, ...

Transaction Log

- Recovery from failures, requires
 - data values prior to modification: BFIM - BeFore Image
 - new value after modification: AFIM – AFter Image
- These values and other information are stored in a sequential file - a **transaction log**
- Sample log data:

T ID	Back P	Next P	Operation	Data item	BFIM	AFIM
T1	0	1	Begin			
T1	1	4	Write	X	X = 100	X = 200
T2	0	8	Begin			
T1	2	5	W	Y	Y = 50	Y = 100
T1	4	7	R	M	M = 200	M = 200
T3	0	9	R	N	N = 400	N = 400
T1	5	nil	End			

Data Update Options

- **Immediate Update:**
 - As soon as a data item is modified in cache, the disk copy is updated
- **Deferred Update:**
 - Modified data items in the cache are written to disk either after a transaction ends its execution, or after a fixed number of transactions have completed their execution

Deferred Update

- **Active Transactions:**
 - transaction updates stored in local workspace or in main memory cache
- **Partially Committed:**
 - updates are transferred to system log
- **Commit:**
 - updates written to disk
- This is a **NO-UNDO/REDO** algorithm
 - uncommitted transactions have not changed database
 - committed transactions are logged

Immediate Update

- **Active Transactions:**
 - updates may reach disk before transaction commits
 - updates must be logged before it is written to disk
- **Transaction Failure:**
 - updates must be removed from the disk
- This is an **UNDO/REDO** algorithm
- If we require all updates reach disk before commit, then REDO is not necessary

Data Caching

- Modified data items are first stored into a cache, and later flushed and written to the disk
- The flushing is controlled by **Dirty** and **Pin** bits (flags)
 - **Pin**: A pinned data item cannot be flushed from the cache
 - **Dirty (Modified)**: A data item has been modified and must eventually be flushed to disk

Cache Flushing

- **In-Place Update:**
Modified values in cache
replace actual values on disk
- **Shadow update:**
Modified version of a data item does not overwrite
disk copy but is written at a separate disk location

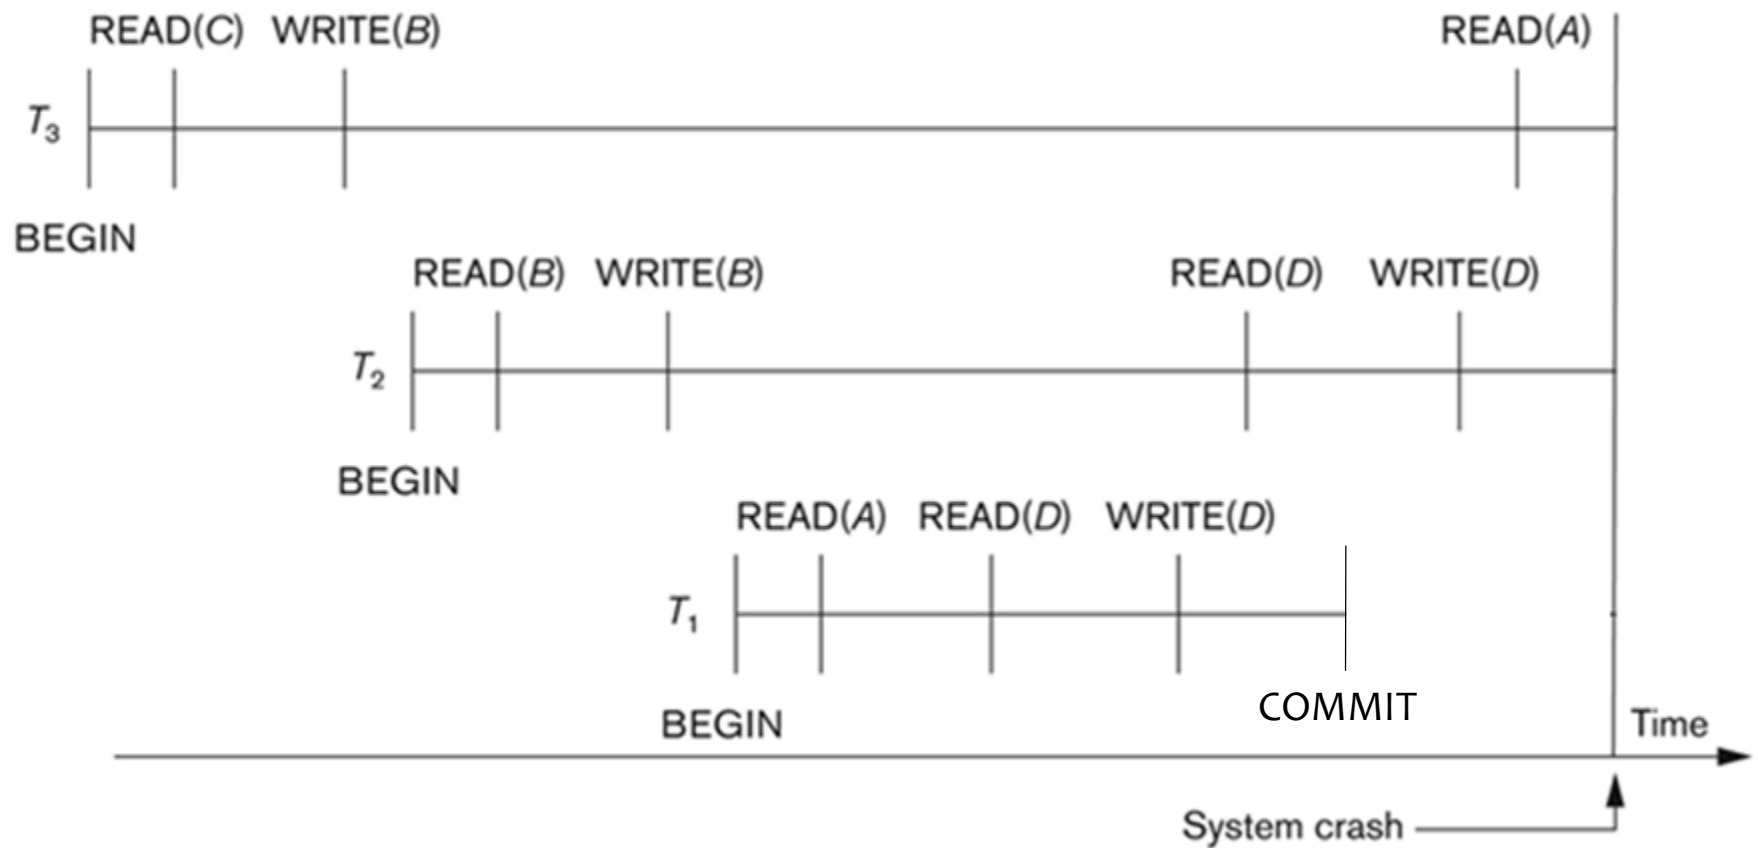
Undo and Redo

- To maintain atomicity,
a transaction's operations
may need to be redone or undone
- **Undo (roll-back):**
 - restore all BFIMs to disk (replace all AFIMs)
- **Redo (roll-forward):**
 - restore all AFIMs to disk

Roll-back Example

T_1	T_2	T_3
read_item(A)	read_item(B)	read_item(C)
read_item(D)	write_item(B)	write_item(B)
write_item(D)	read_item(D)	read_item(A)
	write_item(D)	write_item(A)

Three concurrent transactions and timeline before system crash



Roll-back Example

Transaction log
at time of crash

	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
	30	15	40	20
[start_transaction, T_3]				
[read_item, T_3, C]				
[write_item, $T_3, B, 15, 12$]		12		
[start_transaction, T_2]				
[read_item, T_2, B]				
[write_item, $T_2, B, 12, 18$]		18		
[start_transaction, T_1]				
[read_item, T_1, A]				
[read_item, T_1, D]				
[write_item, $T_1, D, 20, 25$]				25
[read_item, T_2, D]				
[write_item, $T_2, D, 25, 26$]				26
[read_item, T_3, A]				

[commit, T_1] →

← System crash

Roll-back Example

	A	B	C	D
	30	15	40	20
[start_transaction, T_3]				
[read_item, T_3 , C]				
[write_item, T_3 , B, 15, 12]		12		
[start_transaction, T_2]				
[read_item, T_2 , B]				
[write_item, T_2 , B, 12, 18]		18		
[start_transaction, T_1]				
[read_item, T_1 , A]				
[read_item, T_1 , D]				
[write_item, T_1 , D, 20, 25]				25
[read_item, T_2 , D]				
[commit_transaction, T_1]				
[write_item, T_2 , D, 25, 26]				26
[read_item, T_3 , A]				

← System crash

T_3 is rolled-back, since it has not yet committed

T_2 is also rolled-back, since it read values written by T_3

T_1 is has committed and is not dependent on other transaction, so it's updates should remain in database

Restored database state should be $\langle 30, 15, 40, 25 \rangle$

Write-Ahead Logging

- The **Write-Ahead Logging (WAL)** protocol insures that log is consistent with database state at the time of a crash
- WAL states that
 - **For Undo:** Before a data item's AFIM is flushed to the database disk (overwriting the BFIM) its BFIM must be written to the log
 - **For Redo:** Before a transaction executes its commit operation, all its AFIMs must be written to the log
 - In both cases, the log must be saved in stable storage, before the flush or commit is processed.

Cache Control Options

- Steal = no pinning
 - can flush data items to recover buffer space
 - smaller buffer space requirements
- No-steal = pinning
 - cannot flush pinned data items before xact commits
 - may require larger buffer space
- Force
 - dirty data items must be flushed when xact commits
- No-force
 - dirty data items do not have to be flushed at commit (but do need to be flushed eventually)

Recovery Schemes

- The force/no-force and steal/no-steal protocols used determine the recovery scheme:

Steal/No-Force → Undo/Redo

Steal/Force → Undo/No-redo

No-Steal/No-Force → No-undo/Redo

No-Steal/Force → No-undo/No-redo