

COMP163: Database Management Systems

Object-Oriented and Object-Relational Databases Chapter 11

Relational Data Model

- Pros:

- simple: one basic structure
- well-founded – solid theory
- efficient and optimizable

- Cons:

- simple: no complex structures
- does not fit data structures of modern programming paradigms

Objects + DBs: approaches

- Translation
 - objects in program \leftrightarrow relations in DB
- Augment Relational DB
 - add support for complex type within the RDM
 - object-relational databases
- Full OO
 - build database around OO data model
 - object-oriented databases

Translation: Objects \leftrightarrow Relations

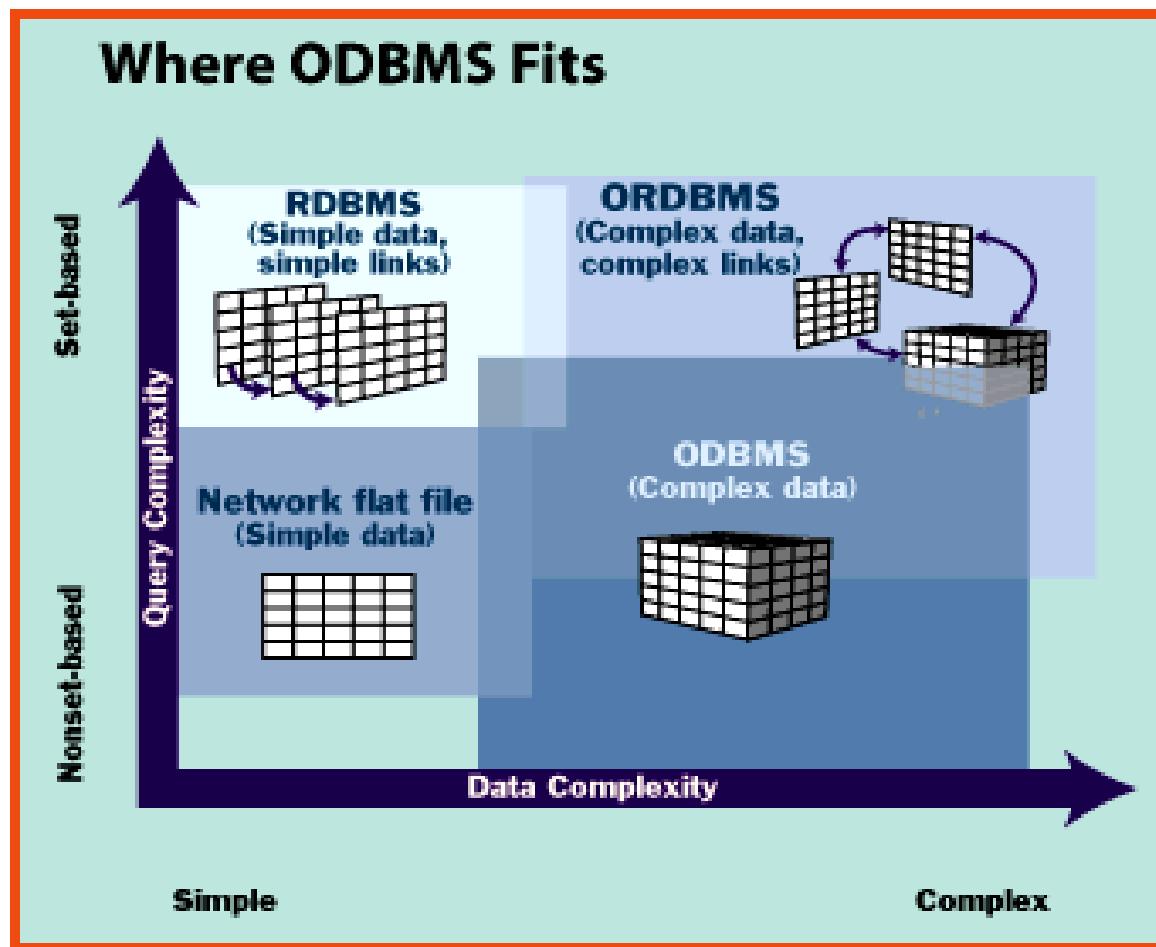
- Ad-hoc:
 - API provides SQL interface to database
 - programmer responsible for mapping relations into objects and back
 - common in data driven web applications (PHP)
- Systematic:
 - defined process for mapping
 - similar to serialization/marshalling/pickling

Mapping: relational – OO

relational	object-oriented
relations	classes
rows/tuples	objects
columns/attributes	attributes
primary keys	object identifiers
foreign keys / join attributes	pointers/references
	complex structures
	inheritance

Algorithms/implementations for this mapping exist.
Example: JDO (Java Data Objects)
gives automatic persistence to Java objects

Data Model Spectrum



Object-Relational DBs

- Allow User Defined Types (UDT)
- Put UDT instances in tables
 - breaks 0NF
- Allow method calls on UDT instances in SQL commands
- Add direct references between UDT instances

Object-oriented DBs

- DBMS built on OO data model
- Direct persistence of objects in application
- Issues:
 - reachability of objects
 - object identifiers
 - relationships
 - query optimization
 - joins vs. pointer chasing

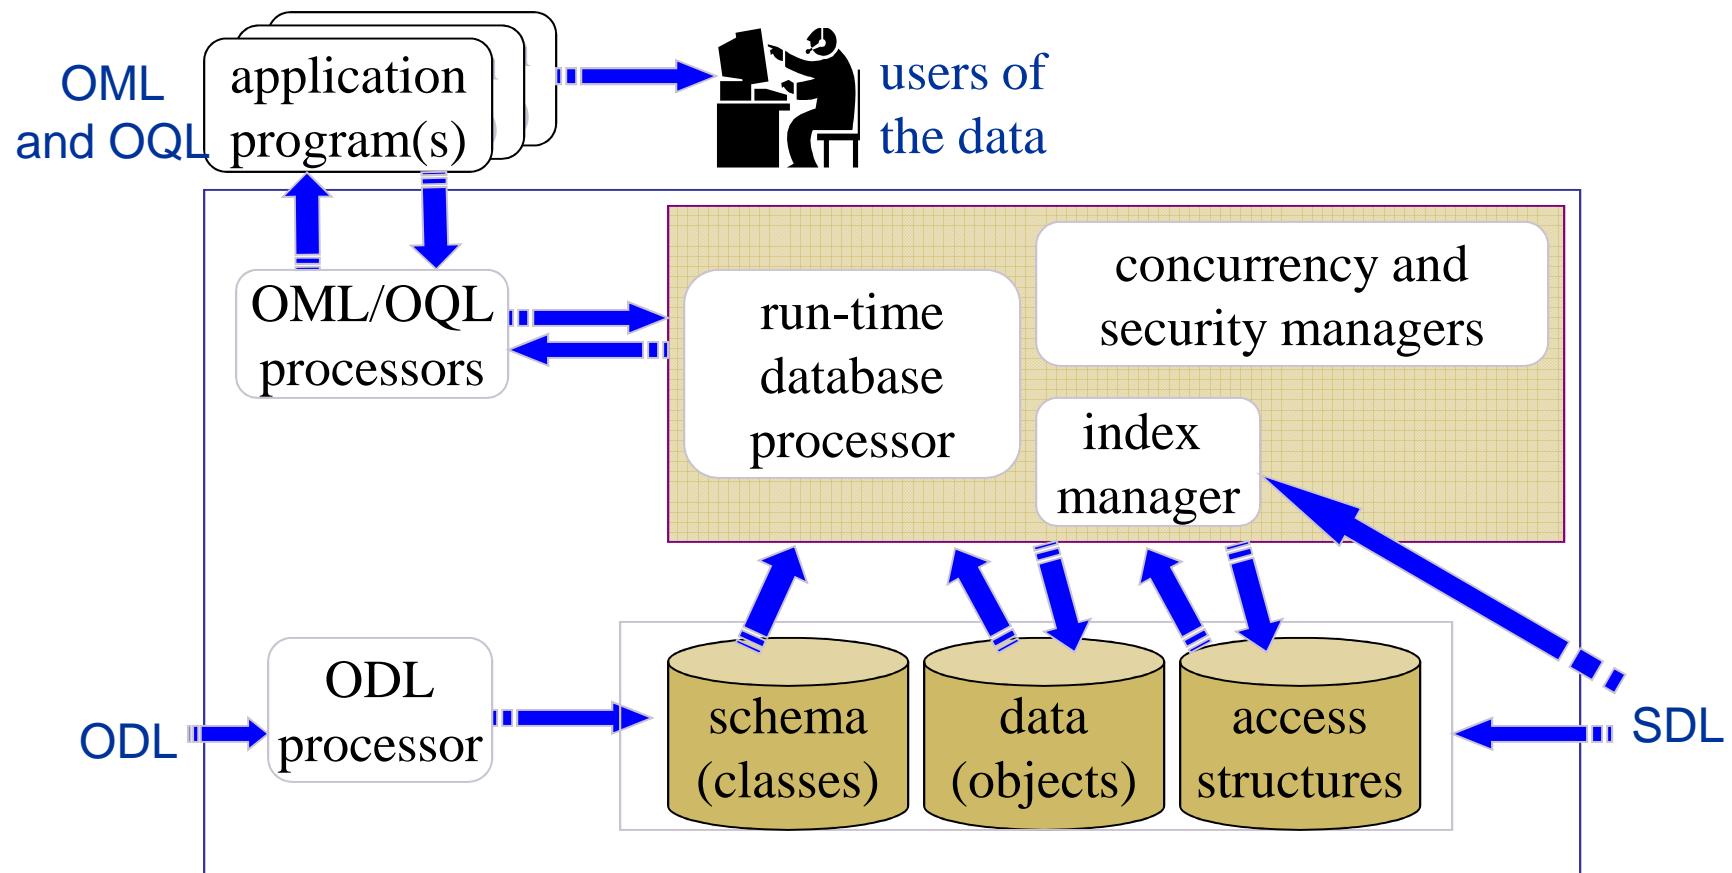
ODMG: The Object Database Management Group

- Early OODBMS vendors based their systems on slightly different object models:
 - Versant, ONTOS: persistent C++ objects
 - O2: object model based on complex value theory
 - Others: persistent Smalltalk objects, Objective-C, etc.
- ODMG standardization developed a common model for OODBs
 - similar to the standard relational model.
 - This allows for portability of applications and sharing of objects between systems.

Components of the ODMG Standard

- Object Model: defines the concepts available for defining an OO schema
 - usual OO things: classes, attributes, methods, inheritance
 - database things: relationships, extents, collections, transactions, DBMS control
- Languages:
 - Object Definition Language: ODL
 - Object Query Language: OQL
 - Object Manipulation Language: OML

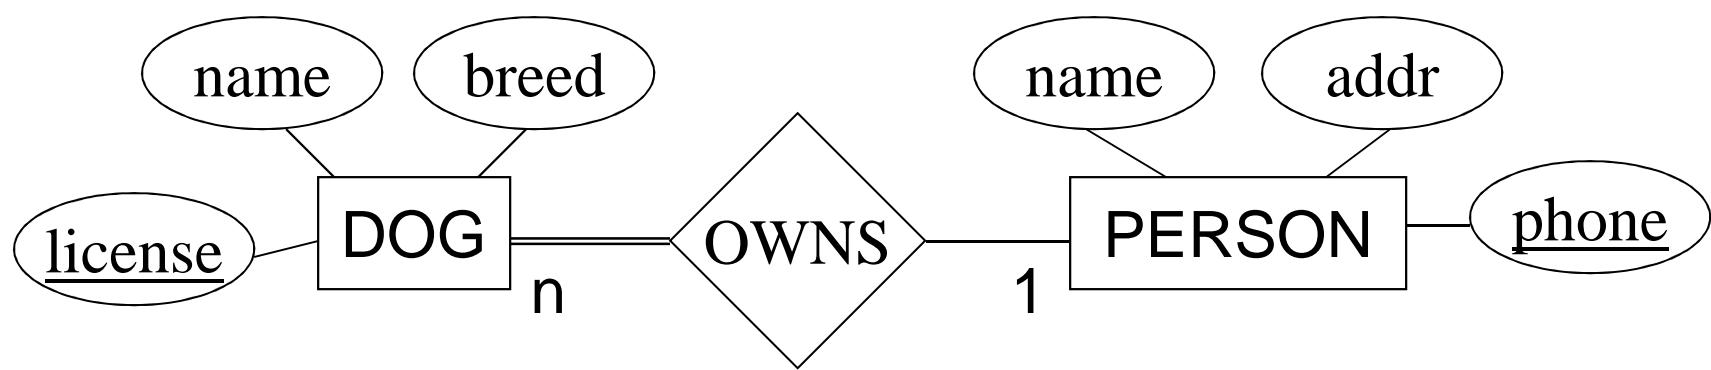
OO-DBMS Architecture



ODL: Object Definition

- ODL defines the syntax for implementing the object model
 - ODL is the language for defining an object schema.
- ODL is actually a family of languages:
 - the ODMG language neutral ODL
 - C++/ODL
 - Java/ODL
 - Smalltalk/ODL
- ODL consists of class declarations

Example Schema: ER



Example Schema: Relational

```
CREATE TABLE PERSON
(
    name    VARCHAR( 20 ) NOT NULL ,
    addr    VARCHAR( 50 ) NOT NULL ,
    phone   CHAR( 10 )      NOT NULL ;
    CONSTRAINT PERSON_PK PRIMARY KEY( phone )
) ;
```

Since phone is the primary key,
it becomes the thing that *identifies* a person.

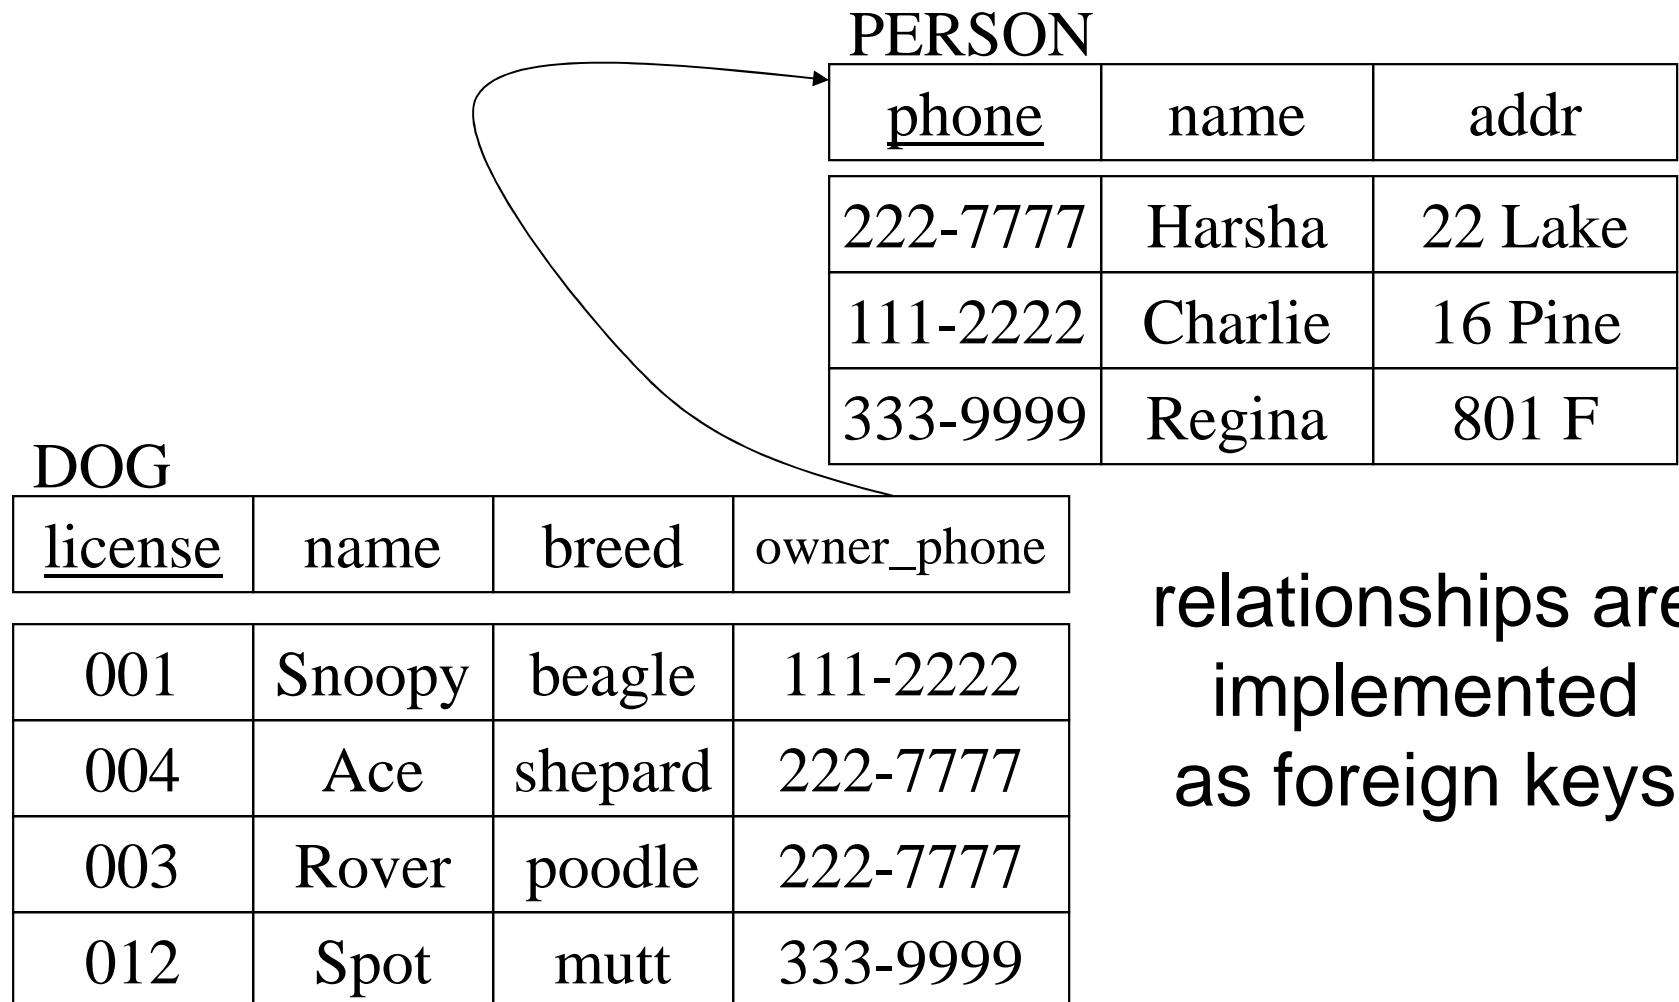
Example Schema: Relational

```
CREATE TABLE DOG
( name          VARCHAR( 20 ) NOT NULL ,
  breed         VARCHAR( 15 ) NOT NULL ,
  license       VARCHAR( 10 ) NOT NULL ;
  owner_phone   CHAR( 10 )      NOT NULL ,
  PRIMARY KEY(license) ,
  FOREIGN KEY (owner_phone)
    REFERENCES PERSON(phone)
) ;
```

A dog is identified by its license number.

A dog's owner is identified by his/her phone number,
since that is a person's primary key.

Relational Instance



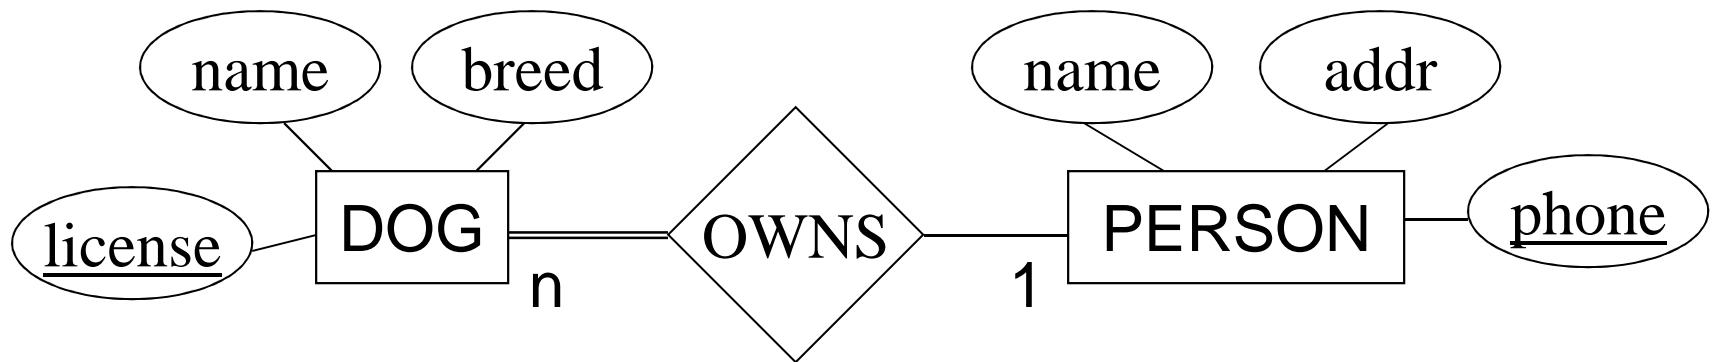
Relational Queries

```
SELECT P.NAME  
FROM PERSON AS P, DOG AS D  
WHERE D.NAME = "Snoopy"  
AND D.owner_phone = P.phone;
```

```
SELECT D.NAME  
FROM PERSON AS P, DOG AS D  
WHERE P.NAME = "Harsha"  
AND D.owner_phone = P.phone;
```

relationships are accessed through
joins over the foreign key

Object Schema

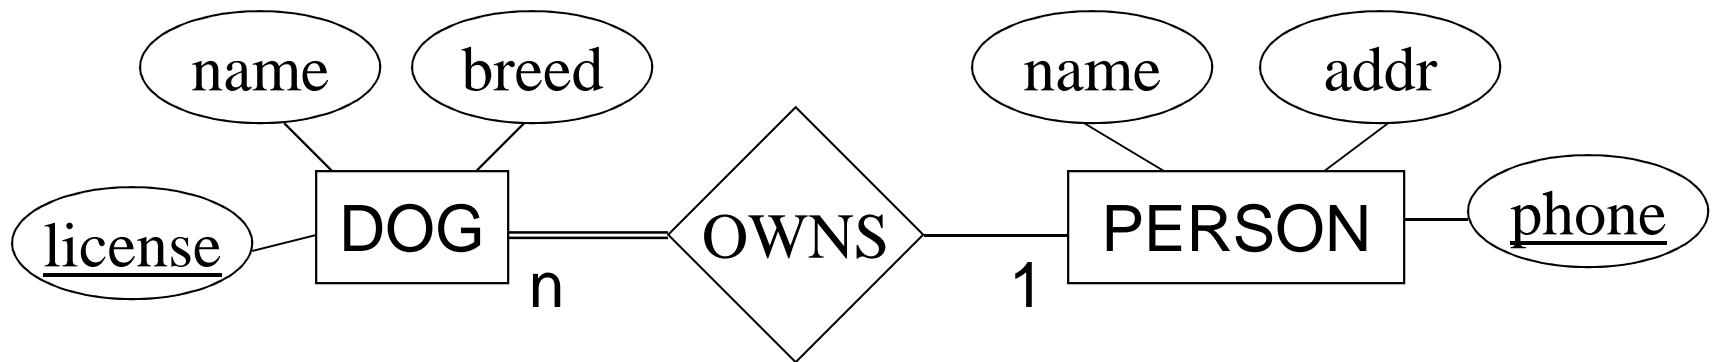


```
class PERSON
(extent people, key phone)
{
    attribute string name;
    attribute string phone;
    attribute string addr;
    relationship set<PET> owns
        inverse PET::owner;
};
```

extent: set of all instances

key: attribute is unique
among all instances
(not a primary key)

Object Schema



```
class DOG
(extent dogs, key license)
{
    attribute string name;
    attribute string breed;
    attribute string license;
    relationship PERSON owner
        inverse PERSON::owns;
};
```

relationships are
references to
other objects

the inverse relationship
implies a consistency
constraint

C++ Schema

```
class PERSON : public d_Object {  
private:  
    d_String name;  
    d_String addr;  
    d_String phone;  
    d_Rel_Set<DOG, "owner"> owns;  
};
```

```
d_Set<d_Ref<PERSON>> people;
```

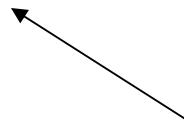
the class extent is a
set of references

all classes inherit
from d_Object
(database object)

one side of the
relationship,
defined as a set
of relationship
references

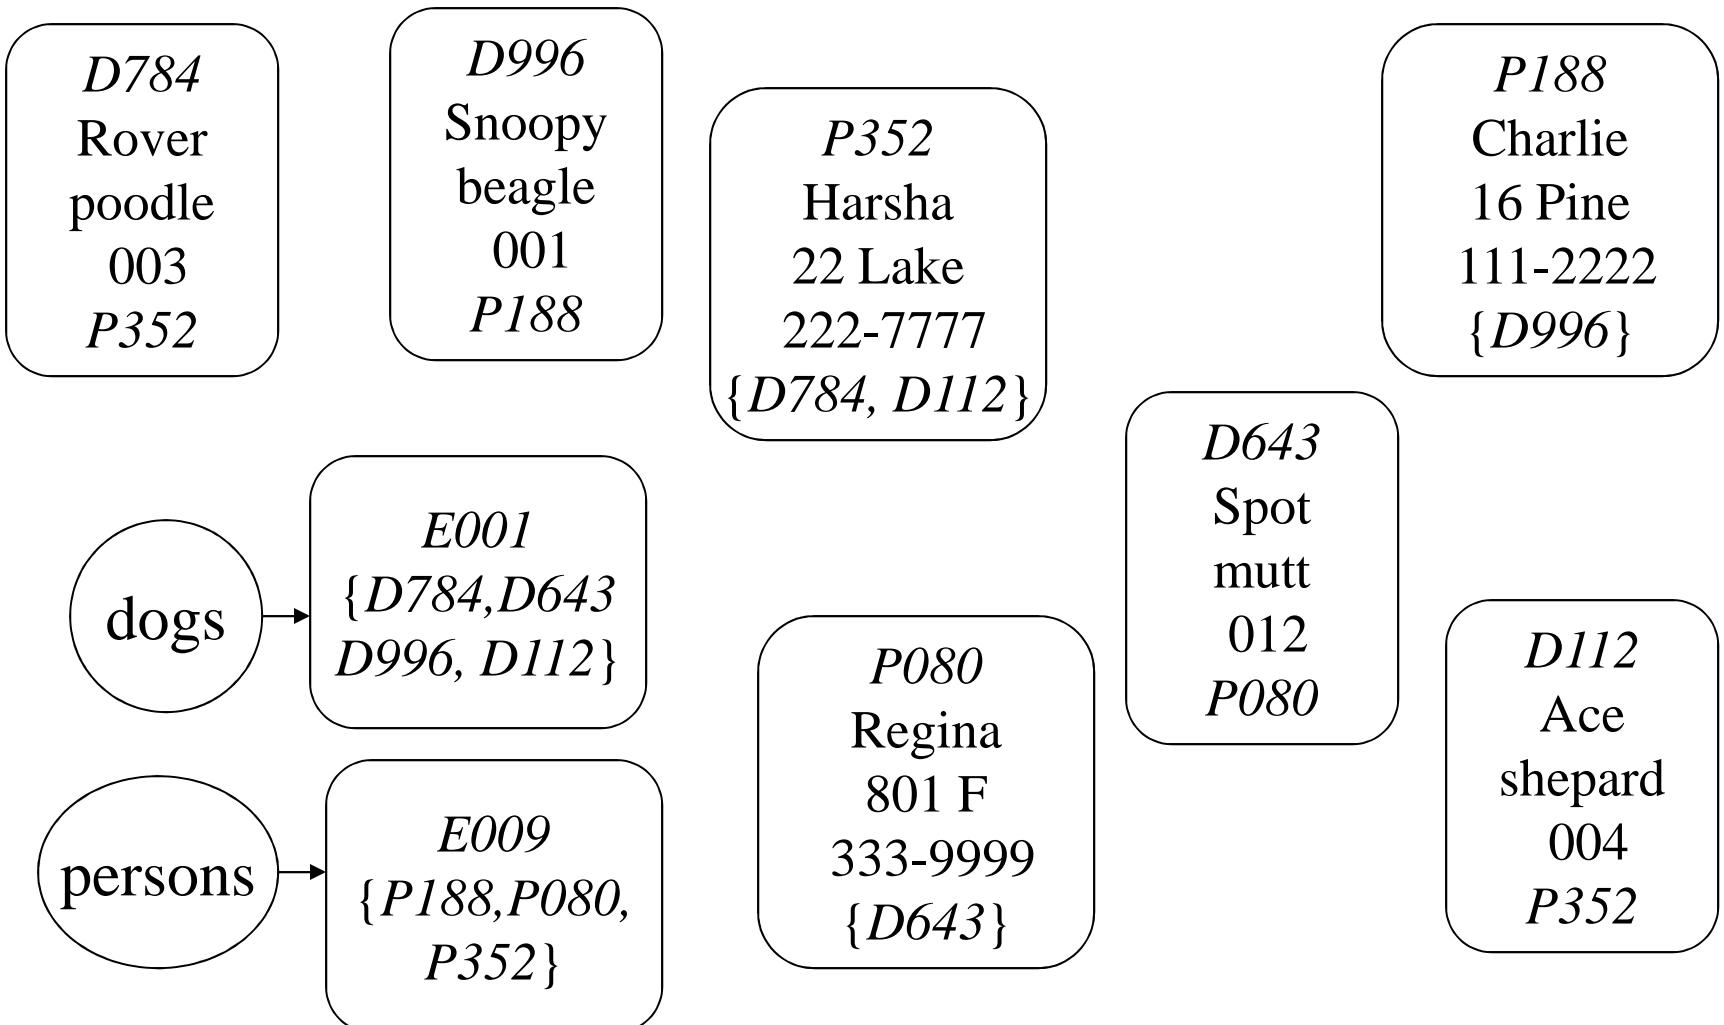
C++ Schema

```
class DOG : public d_Object {  
private:  
    d_String name;  
    d_String breed;  
    d_String license;  
    d_Ref_Ref<PERSON, "owns"> owner;  
};  
  
d_Set<d_Ref<DOG> > dogs;
```



other side of the
relationship,
defined as a single
relationship
reference

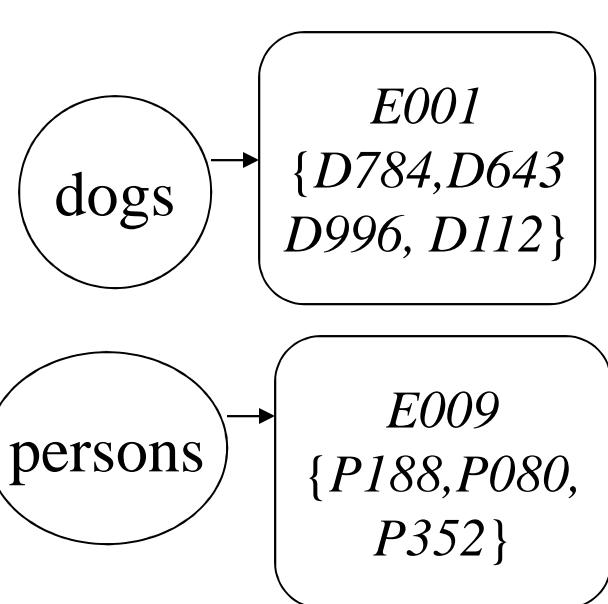
Object Instances



Object Instances

D784
Rover
poodle
003
P352

D996
Snoopy
beagle
001
P188



P188
extents have
persistent names

(notation: circles are names,
not pointers or objects)

extents are sets of
object identifiers (OIDS)

333-9999
{D643}

004
P352

Object Identifiers

- Questions:
 - How are objects identified in C++ or Java programs?
 - Will this work in a database context?

Object Identifiers

- All database objects are assigned unique object identifiers (OIDs)
 - inherited from d_Object
- An OID gives an object an immutable identity, apart from its value
 - In a relational database, the identity of tuples is determined by their value
- An OID identifies an object regardless of its location in memory, on disk or on network

Class Extents

- An extent is simply a set of references to all instances of a class
 - For an OODB, this can be a set of OIDs
- Implementation:
 - hack into all constructors and destructors
- Value:
 - Extent replaces the relation as the basic storage structure
 - Ensures reachability of all objects

Basic Implementation (C++)

```
typedef long OID;

class D_Object
{
public:
    D_Object()          { setOID(); addToExtent(); }
    virtual ~D_Object() { removeFromExtent(); }

protected:
    static set<OID> extent;

private:
    OID oid;

    static OID next_oid;

    void setOID()        { oid = next_oid++; }
    void addToExtent()   { extent.insert(oid); }
    void removeFromExtent() { extent.erase(oid); }
};
```

Basic Implementation (C++)

```
typedef long OID;

class D_Object
{
public:
    D_Object()
    virtual ~D_Object() { }

protected:
    static set<OID> extent;

private:
    OID oid;

    static OID next_oid;

    void setOID() { }
    void addToExtent() { }
    void removeFromExtent() { }
};
```

Additional practical issues:

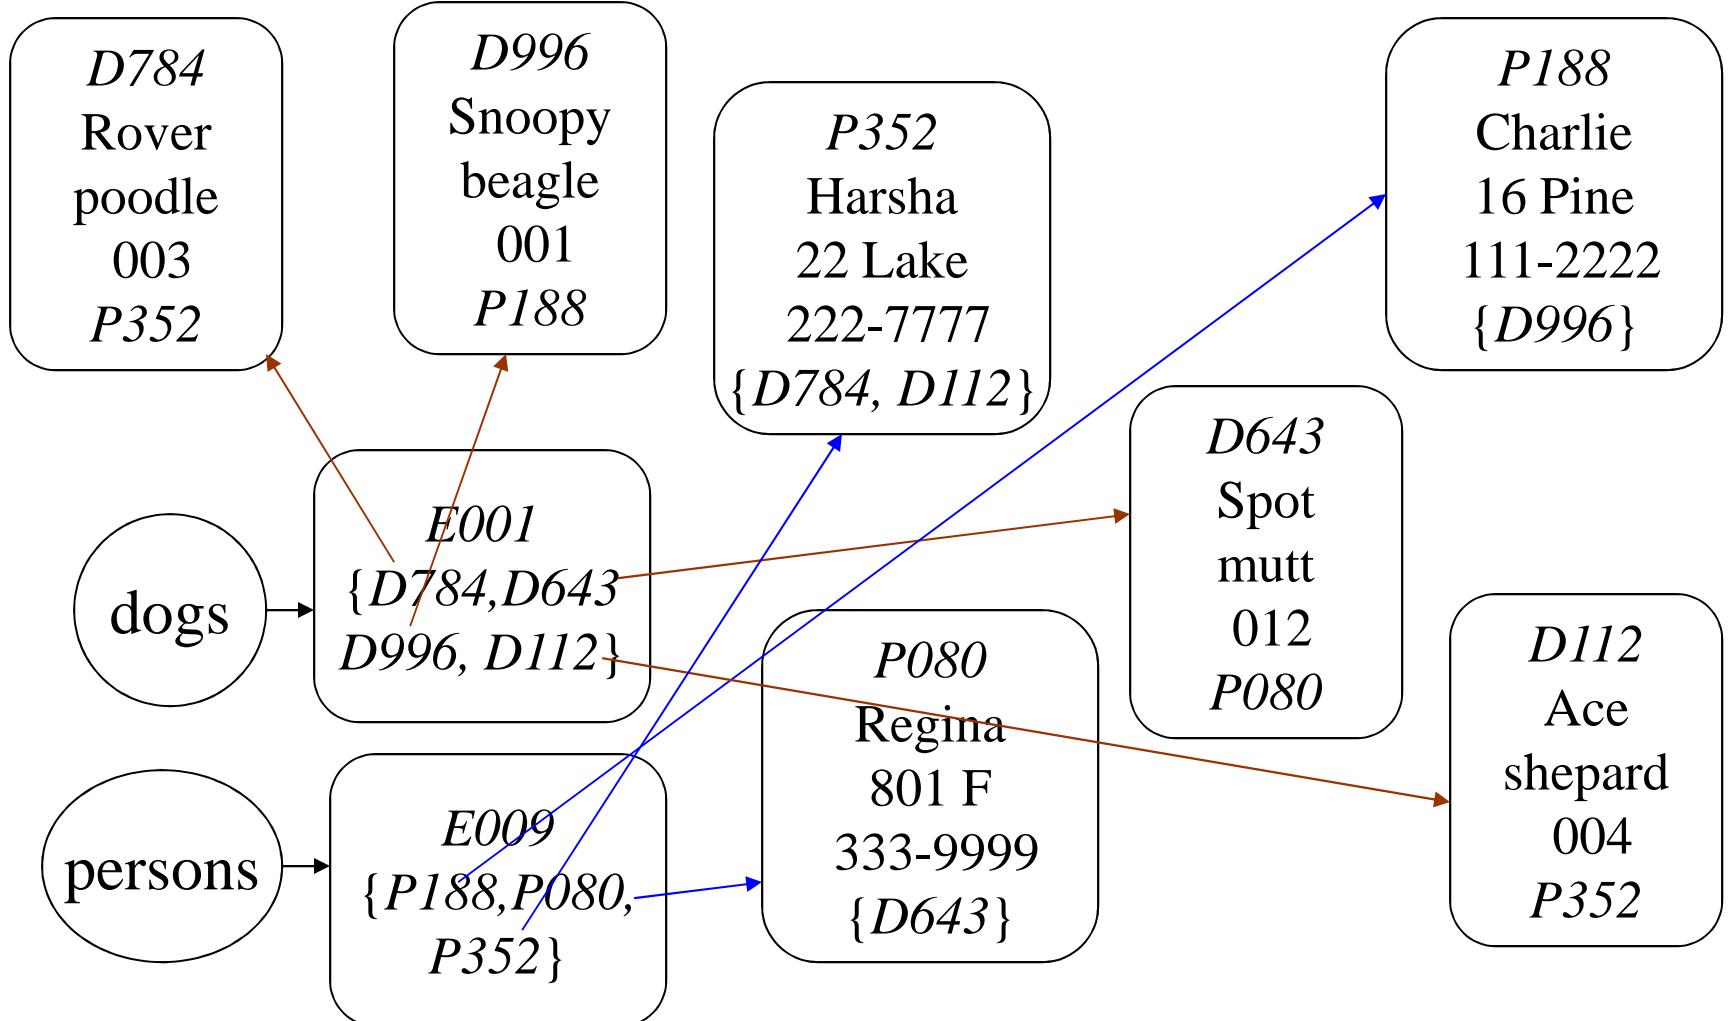
More complex OIDs require an OID generator

Need ability to map OIDs to object pointers/references

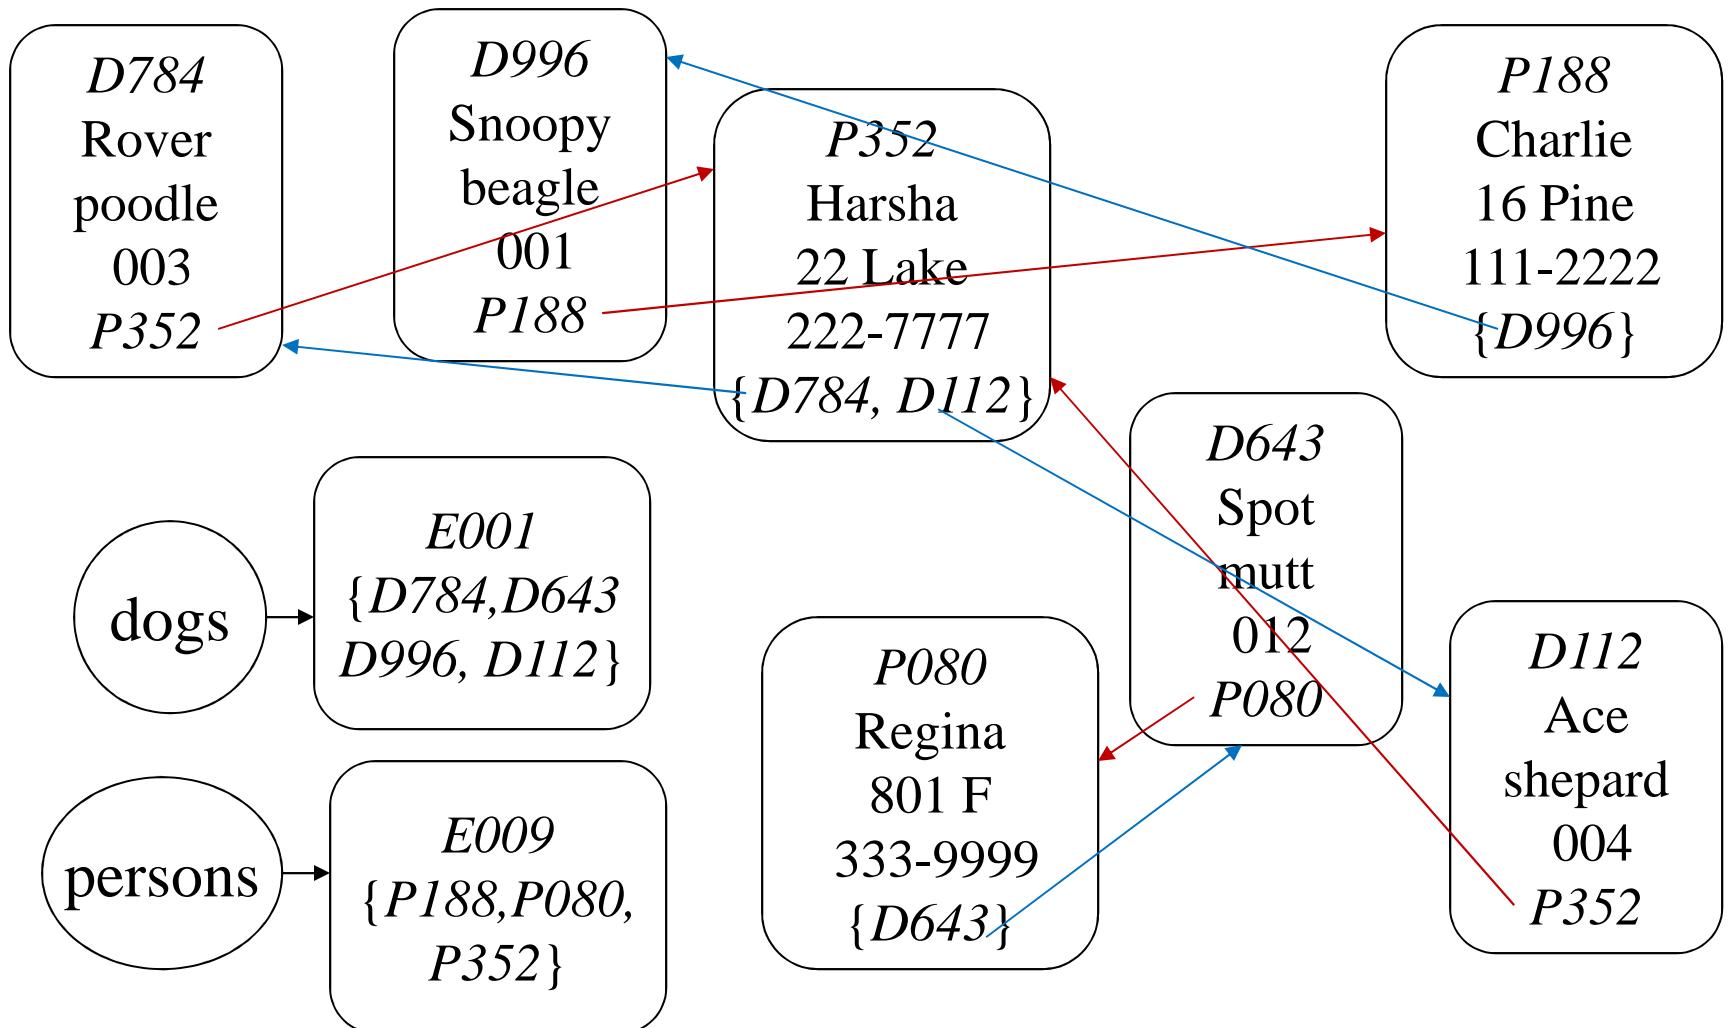
Need separate extents for each class

{
{ (This are all easily solvable)

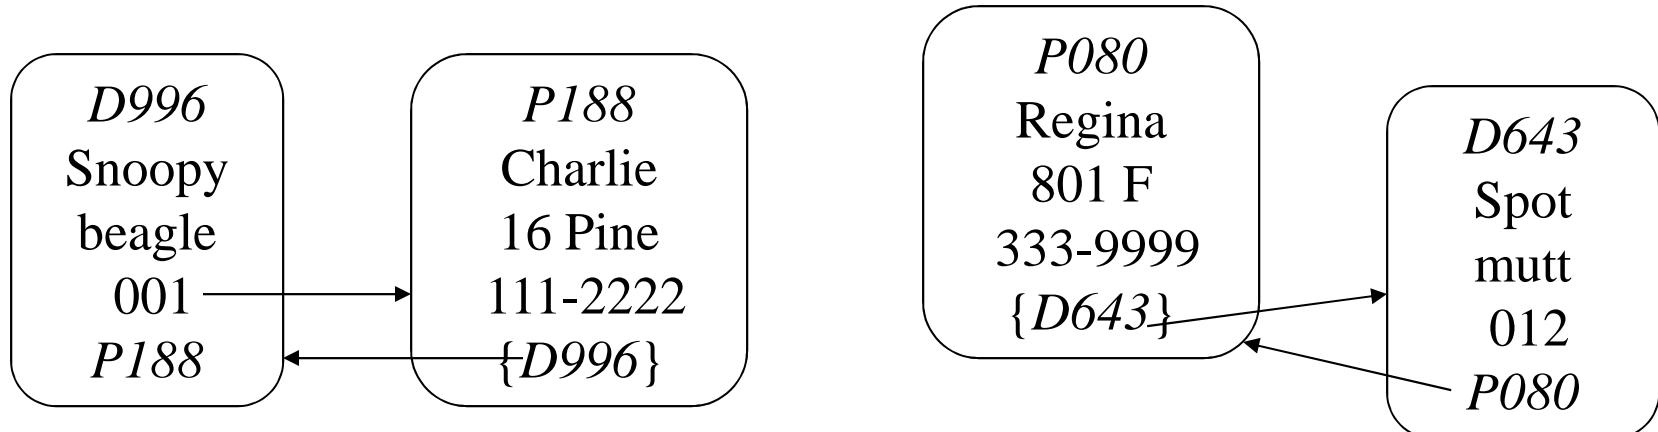
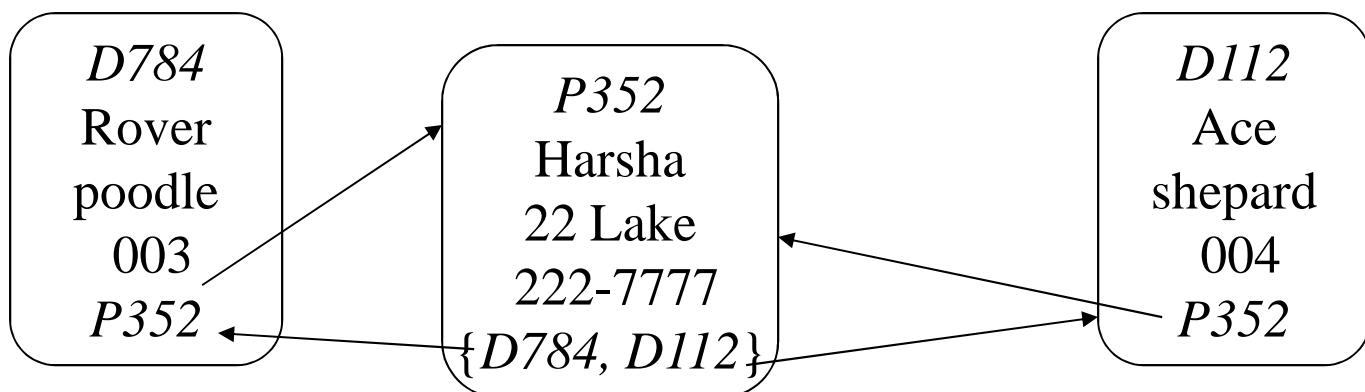
OIDs Are References



OIDs Are References



Relationships are Reciprocal References

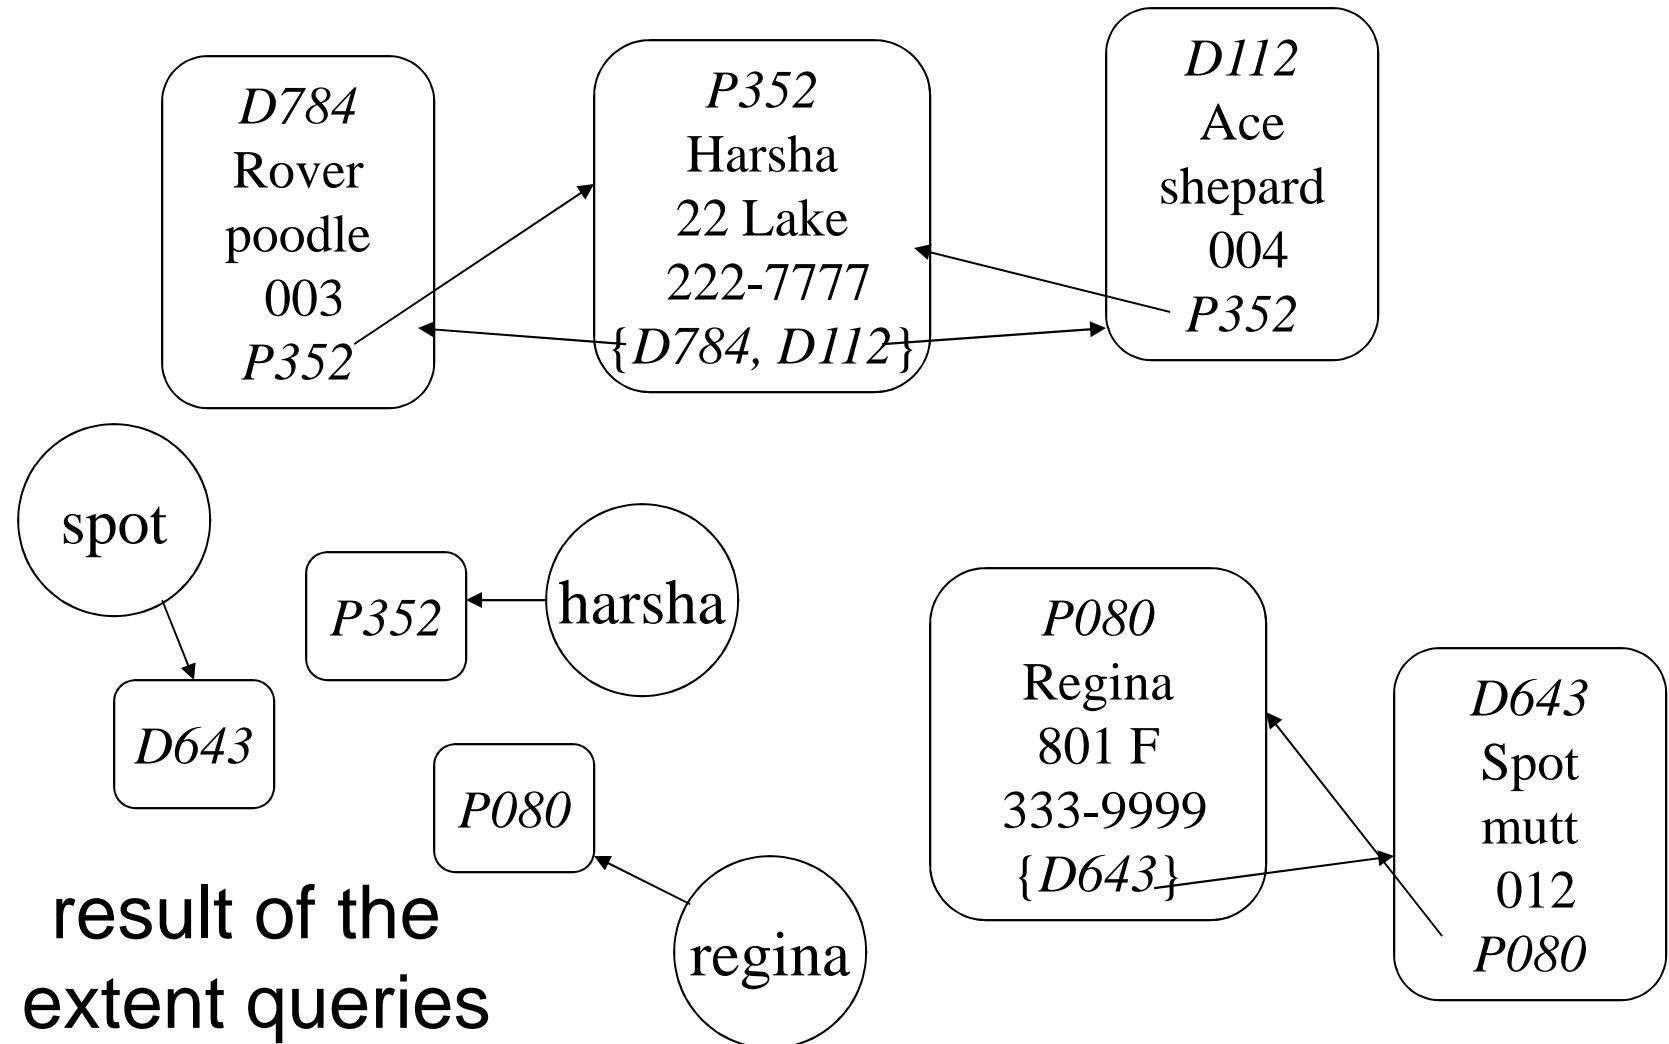


Relationship Maintenance

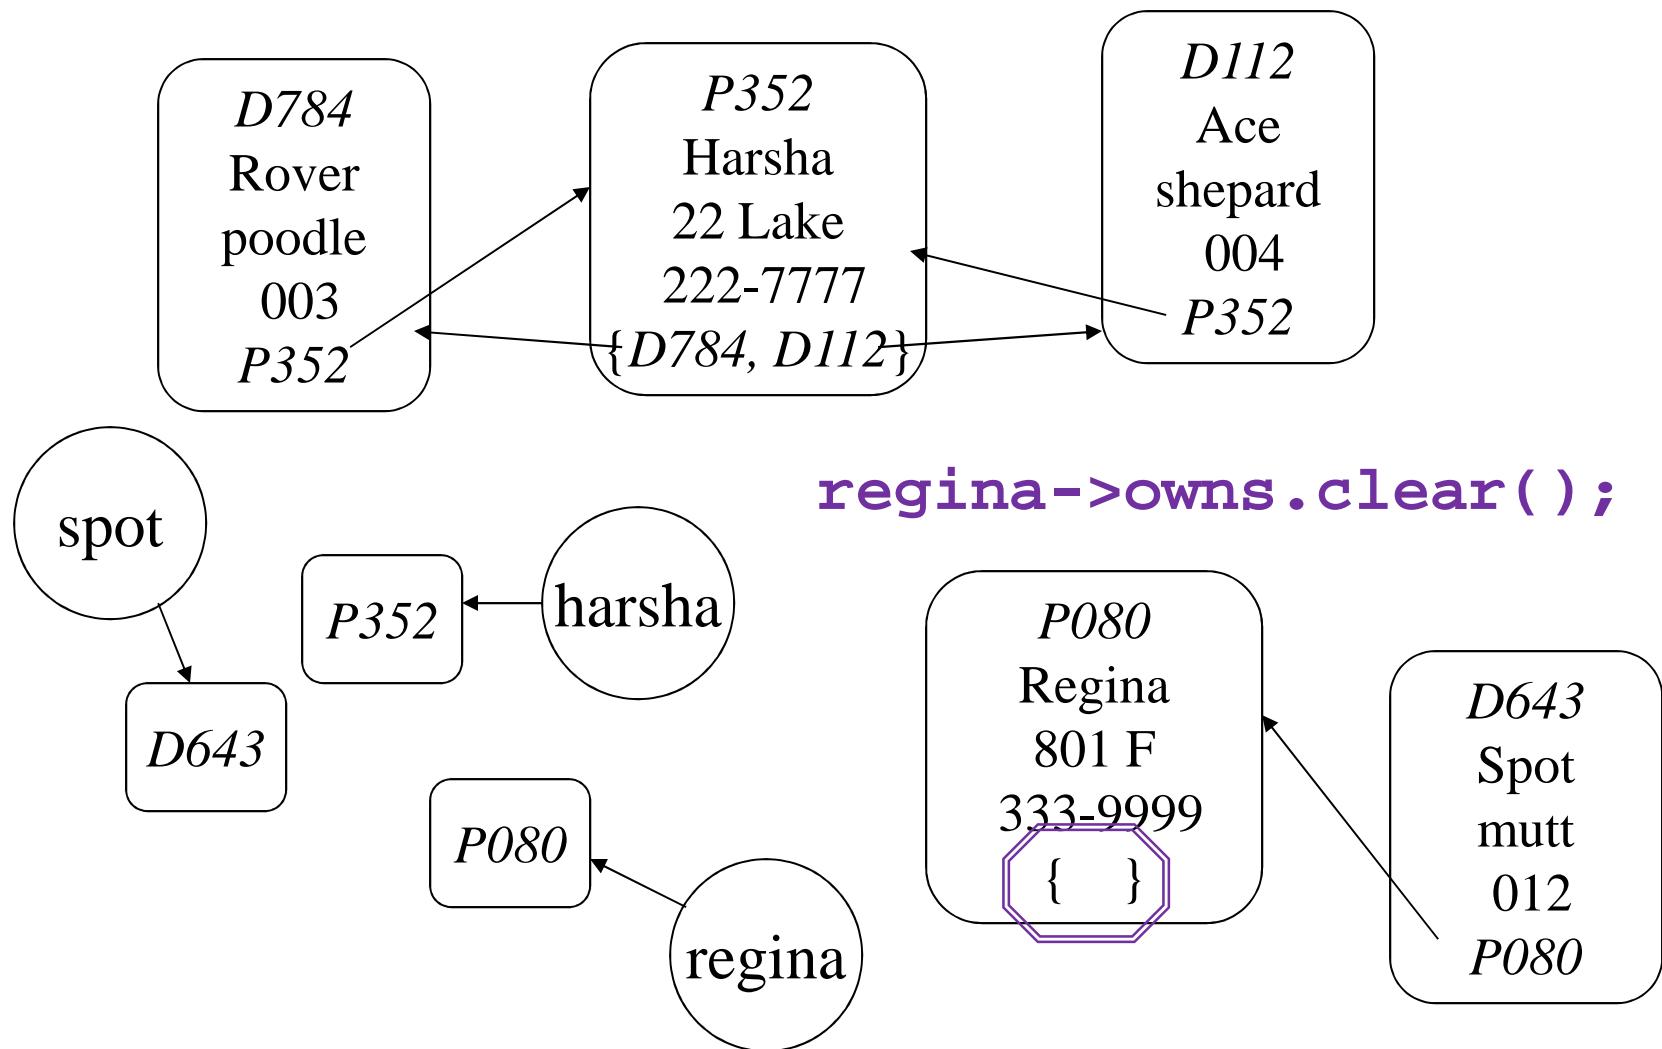
- The DBMS is responsible for maintaining the consistency of relationships
- Suppose Regina sells her dog to Harsha:

```
d_Ref<Person> regina =
    people.select_element( "name='Regina' " );
d_Ref<Person> harsha =
    people.select_element( "name='Harsha' " );
d_Ref<Dog> spot =
    regina->owns.select_element( "name='Spot' " );
regina->owns.clear();
harsha->owns.insert_element(&spot);
```

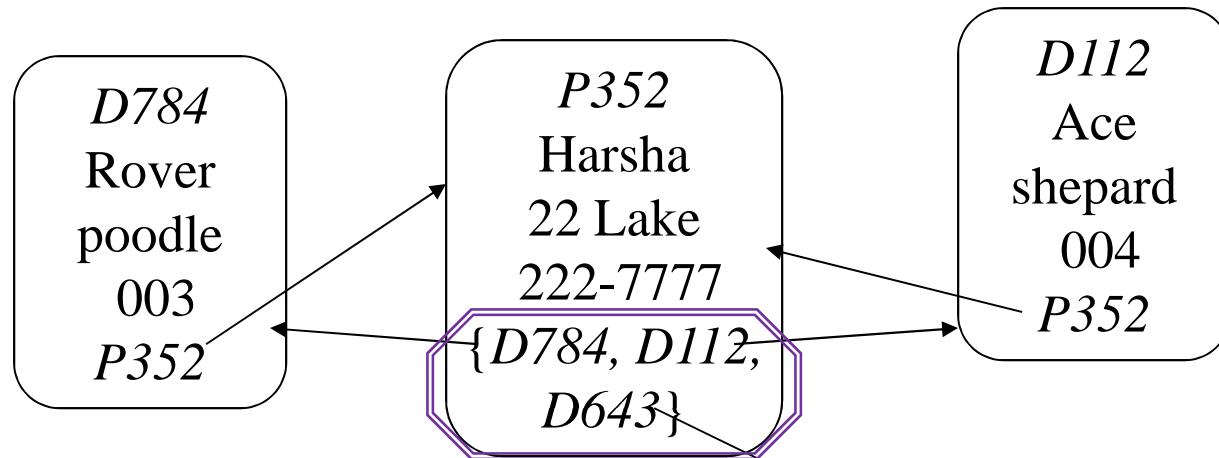
Relationship Maintenance



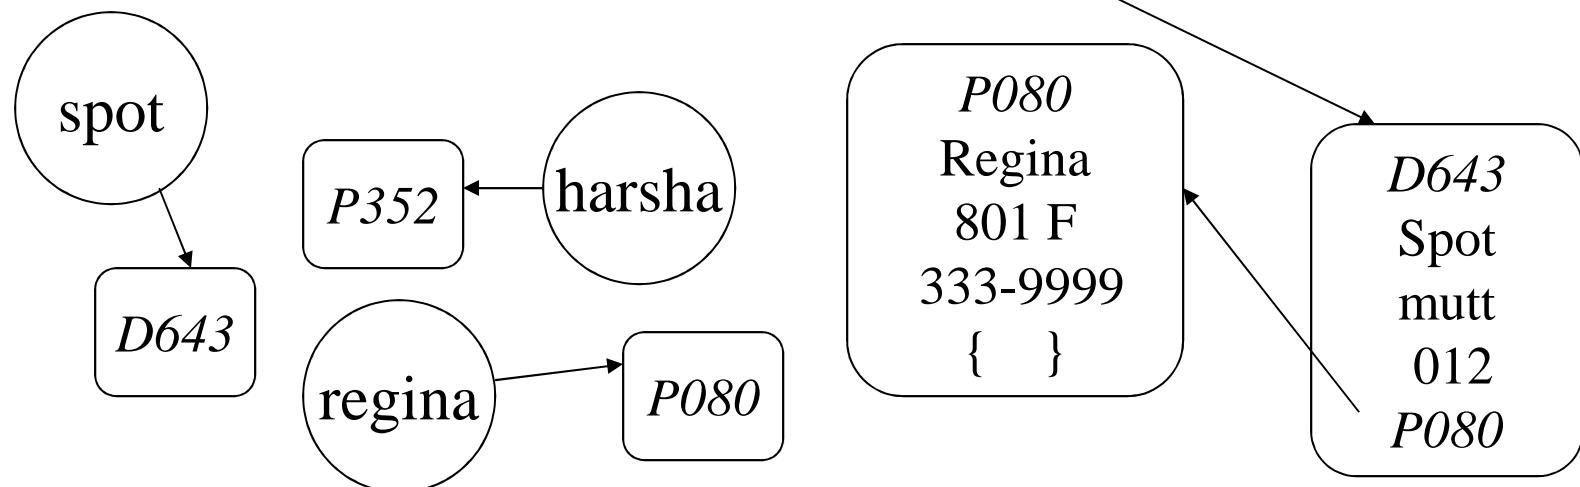
Relationship Maintenance



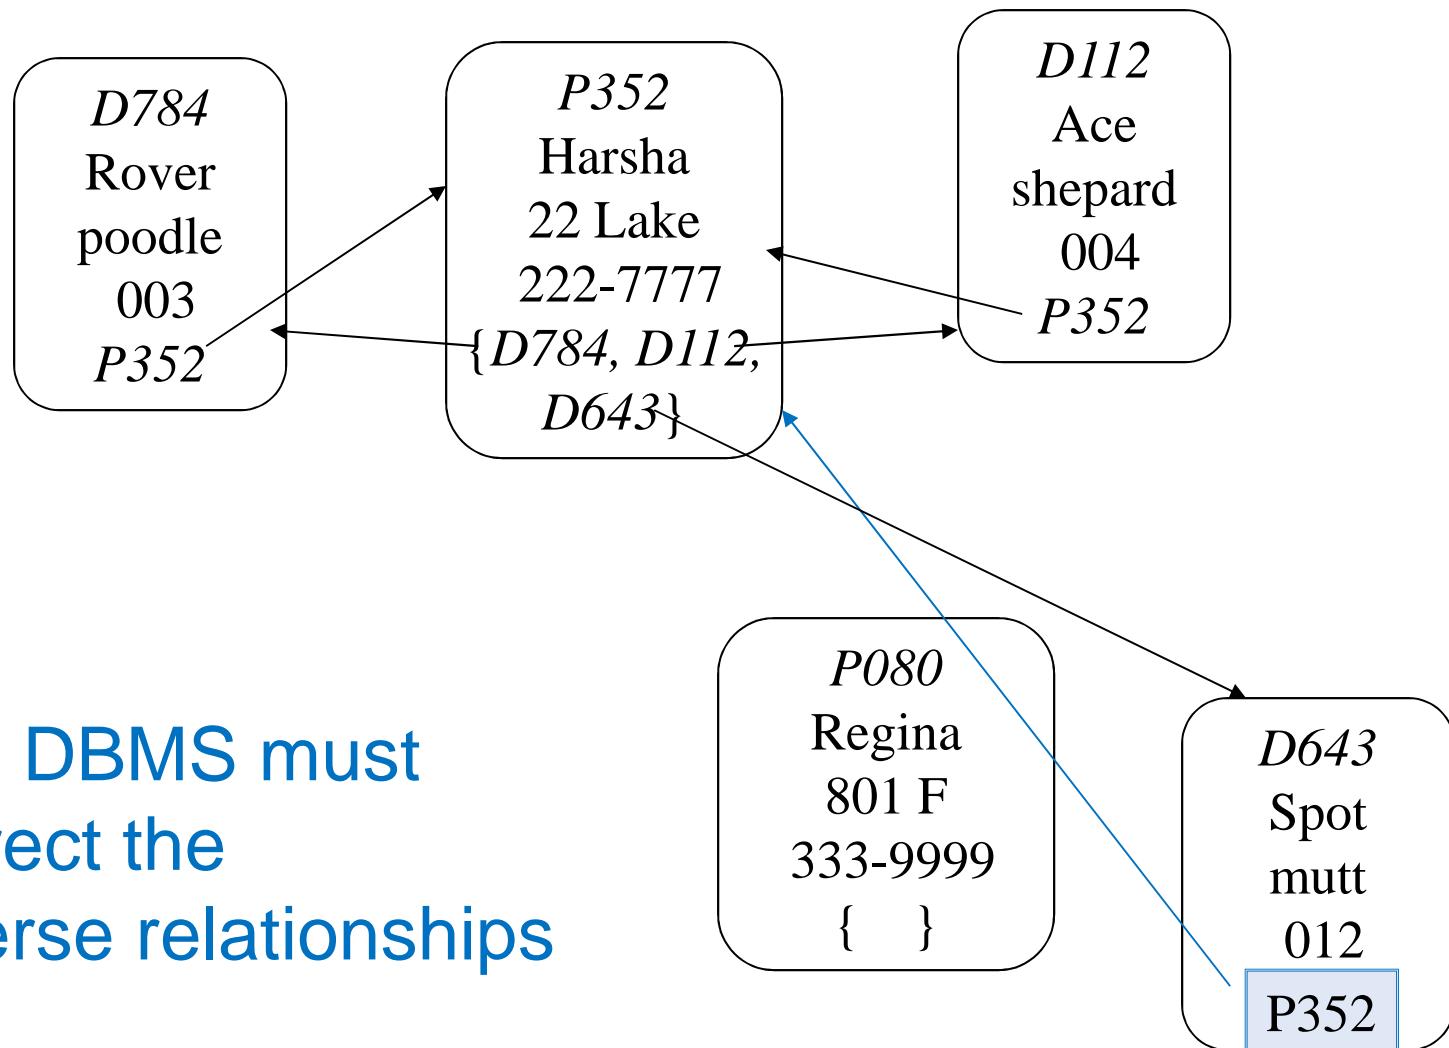
Relationship Maintenance



harsha->owns.insert_element(&spot);



Relationship Maintenance



OQL

■ Schema for OQL examples

```
class Professor
(extent professors, key uid)
{
    attribute string name;
    attribute short uid[9];
    relationship set<Student> advises
        inverse Student::advisor;
    relationship Department dept
        inverse Department::faculty;
    float avgGPAofAdvisees();
};

class Student
(extent students, key uid)
{
    attribute string name;
    attribute short uid[9];
    relationship Professor advisor
        inverse Professor::advises;
    float GPA();
};
```

OQL example

```
select s.name  
from Students s
```

- students is the *extent* of class Student
- s will refer to all Student instance (just as it would refer to all rows in a relation)
- answer is a set of strings

OQL example

```
select s.name  
from Students s  
where s.gpa() >= 3.0
```

- methods can be applied to objects

OQL example

```
select struct(name: s.name,  
              gpa: s.GPA())  
from Students s
```

- answer is a set of tuples
- build tuple structure in the query
(essentially the same as renaming attributes
in relational result)
- call to method in select clause

OQL example

```
select s.name  
from Professors p,  
      p.advises s  
where p.name = "Jill Pebble"  
and s.GPA >= 3.0
```

- equivalent to a join
- What does this query compute?

OQL – optimization issues

```
select p.name  
from Professors p,  
where p.avgGPAofAdvisees() >= 3.0
```

- What objects are accessed by this query?
- Can we estimate the cost of an arbitrary user-defined function?

ODMG Database Fucntionality

- d_Database and d_Transaction classes provide basic DBMS functionality
 - open/close database connection
 - start/commit/abort transactions
 - retrieve objects by name
(essential for persistence/reachability)

d_Database class

```
class d_Database
{
public:
    d_Database();

    enum access_status { not_open, read_write,
                        read_only, exclusive };

    void open(char *database_name,
              access_status _status = read_write);
    void close();

    void set_object_name(d_Ref_Any &theObject, char* theName);
    void rename_object(char *oldName, char *newName);
    d_Ref_Any lookup_object(char * name);

};
```

d_Transaction class

```
class d_Transaction
{
public:
    d_Transaction();
    ~d_Transaction();
    void begin();
    void commit();
    void abort();
    void checkpoint();
};
```

ODMG Extents

- ODL preprocessor generates Set objects that hold all instances of a class:

```
Set<Ref<PERSON> > persons;
```

```
Set<Ref<DOG> > pets;
```

- Extents can be accessed with iterators:

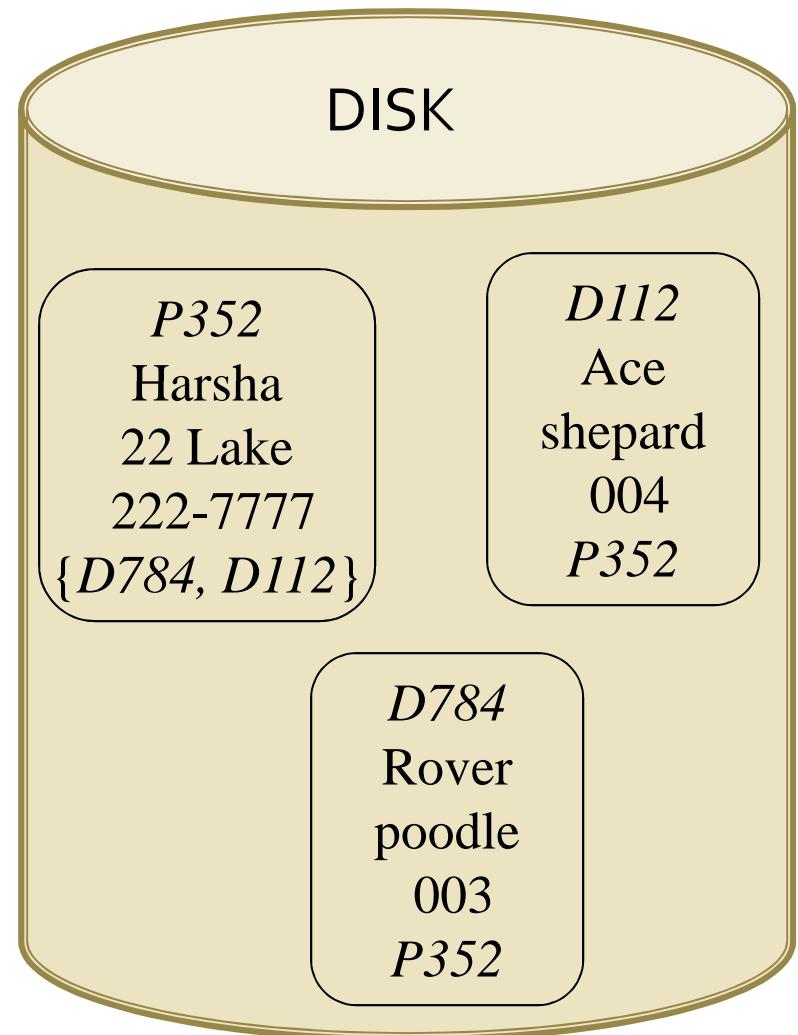
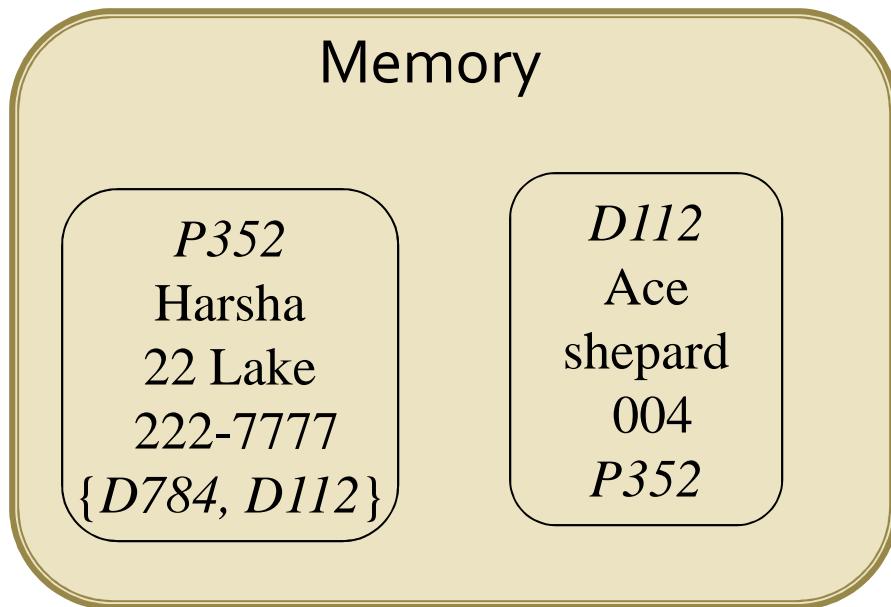
```
d_Iterator<d_Ref<DOG> > pet_iter =  
    pets.create_iterator();
```

```
d_Ref<DOG> p;  
while (pet_iter.next(p))  
    cout << p->name;
```

Persistence Capability

- All classes inherit from the system defined class **d_Object**
- Defines everything that DBMS needs make object behave as a database object.
 - object identifier (OID)
 - object size
 - modified or dirty flag: indicates if an object's state has changed from the value currently stored on disk
 - locking information: controls access to objects from multiple applications
 - object's class

Buffer Management



What happens if we iterate through Harsha's pets?

Smart Pointers

```
d_Iterator< d_Ref<DOG> > iter =  
    harsha.owns.create_iterator();  
d_Ref<DOG> p;  
while (pet_iter.next(p)) cout << p->name;
```

- at some point, p will refer to Rover, who is still on disk
- d_Ref will recognize this and read object from disk before returning pointer

What does this mean in terms of our discussions of disk access costs?