

COMP163

Database Management Systems

Query Processing: Chapter 19

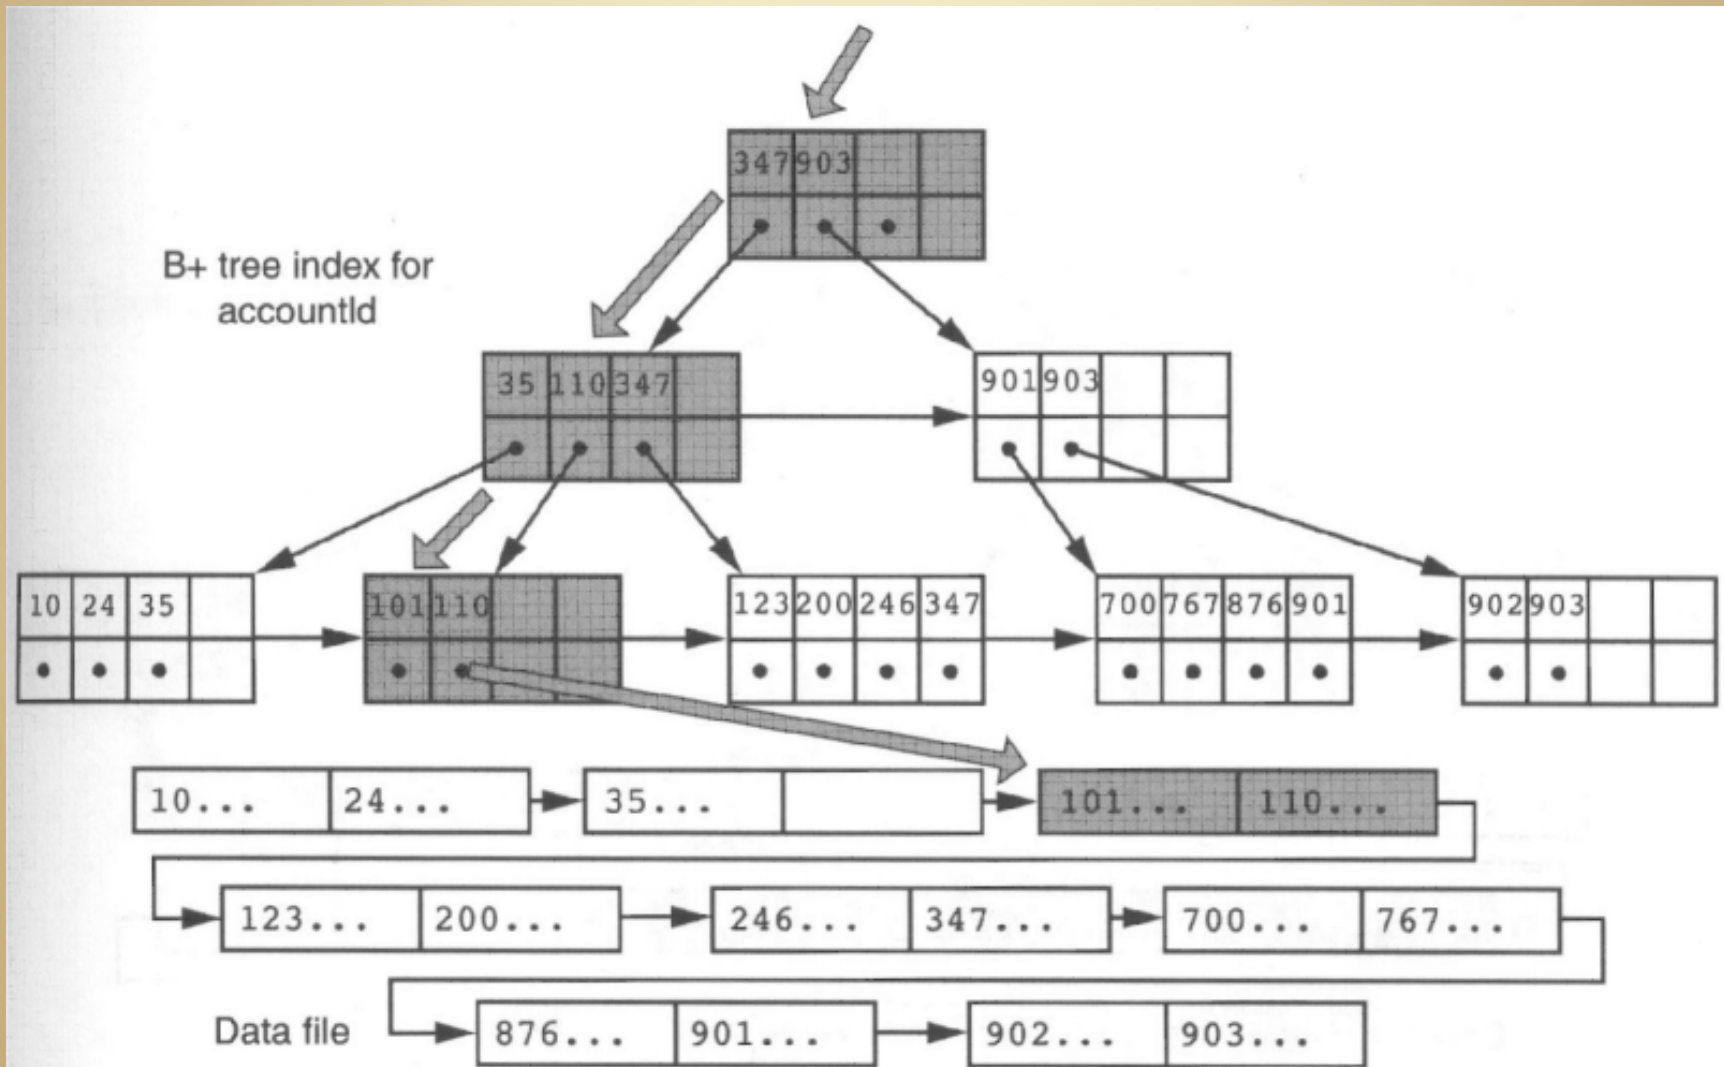
Using Indexes for Queries

Database Instance and Available Indexes:

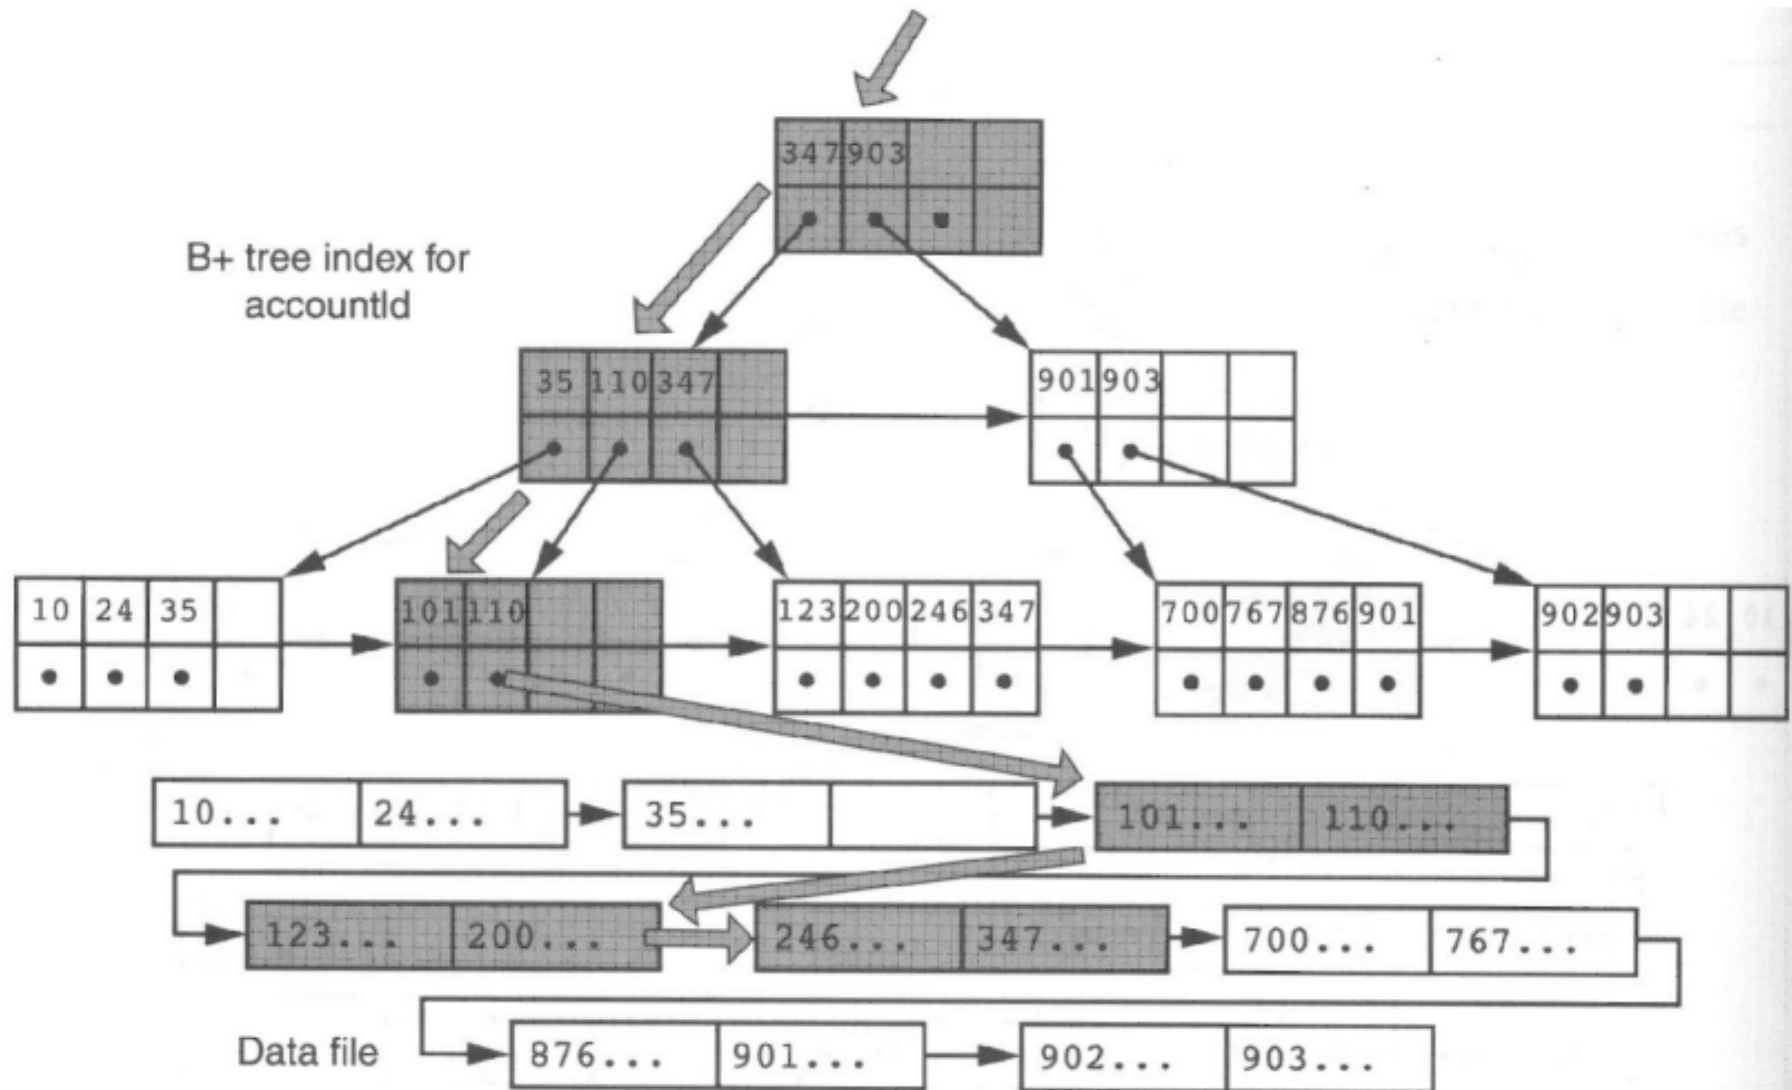
Table	Entries	Number of Entries per Block	Number of Blocks	Index Fields	Index Type	Keys per Node	Depth of B+ Tree
Customer	10,000	10	1,000	accountId	B+ tree and ordered sequential file	100	3
				lastName	B+ tree	50	3
				zipcode	hash	100	
Rental	1,000,000	100	10,000	accountId	B+ tree	100	3
				movieId	B+ tree	100	3
				date	B+ tree	100	2
Movie	10,000	20	500	movieId	B+ tree and ordered sequential file	100	3
				title	B tree	20	4
				genre	Hash	100	

Examples from: *Principles of Database Systems* by Greg Ricciardi, Addison Wesley, 2001

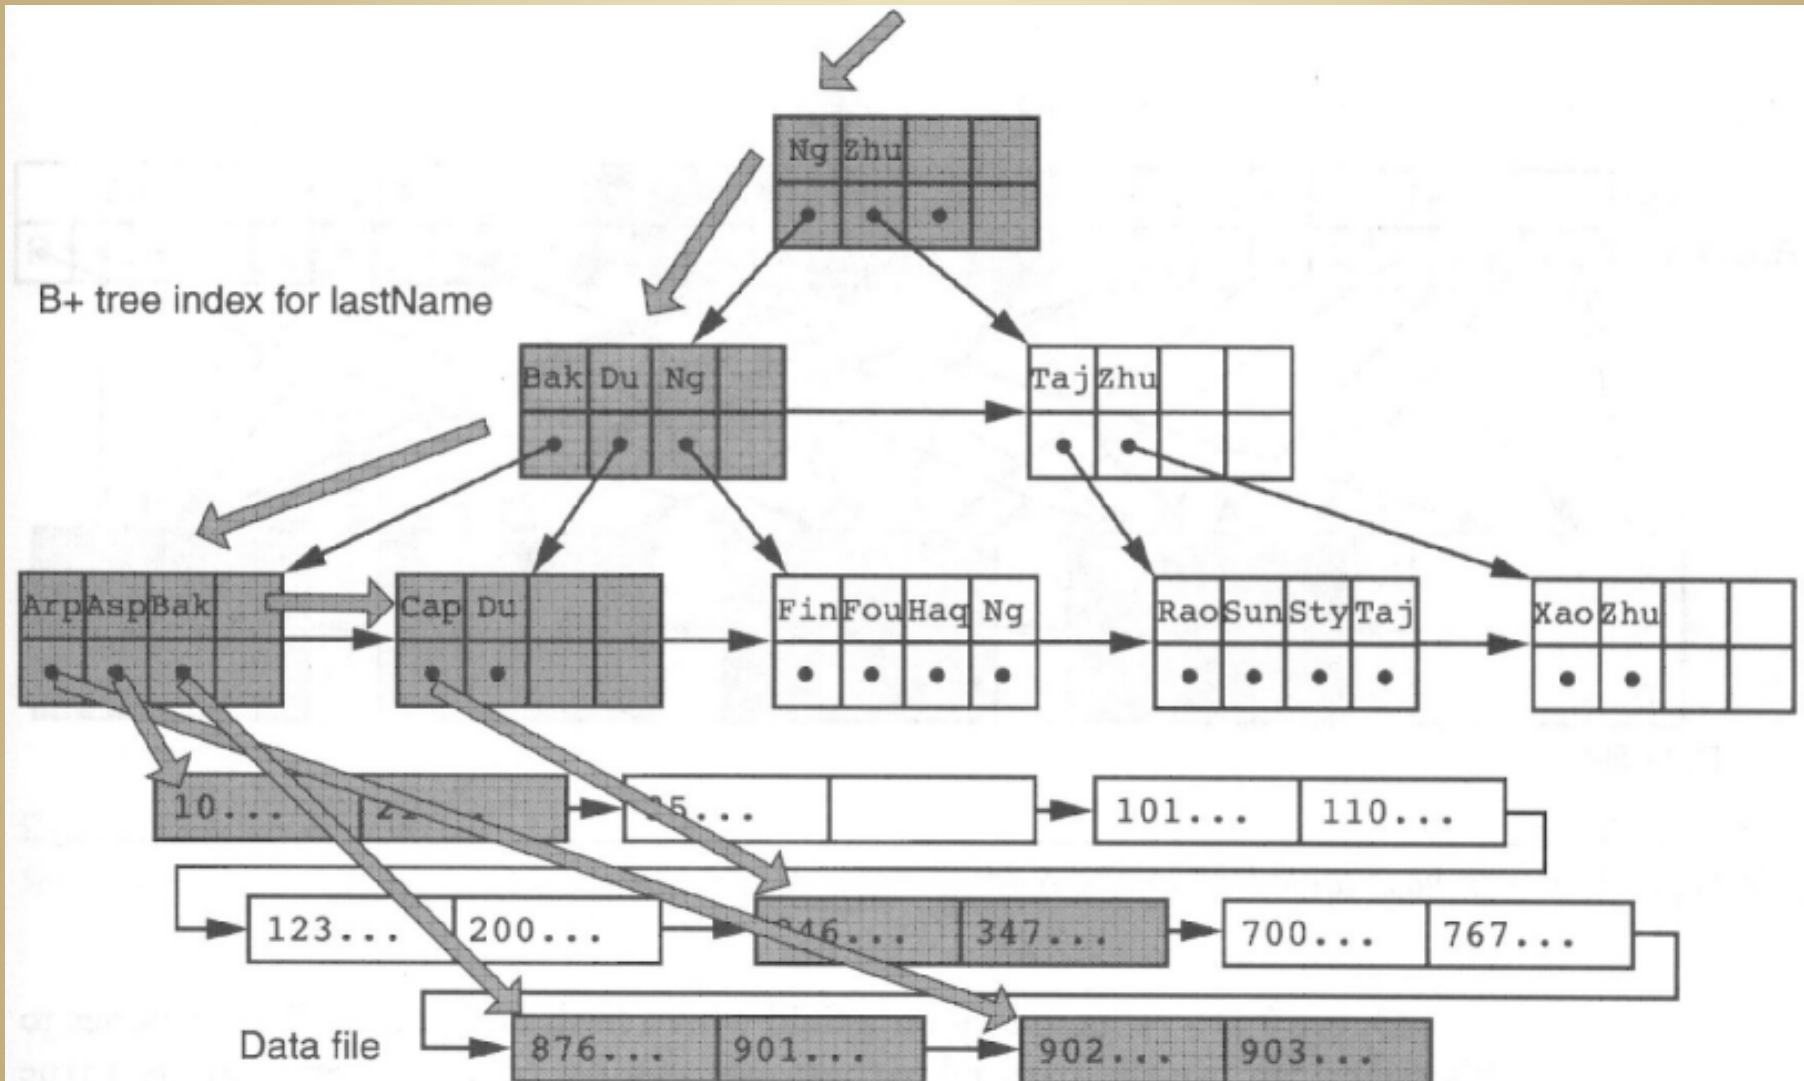
select * from Customer
where accountId = 101



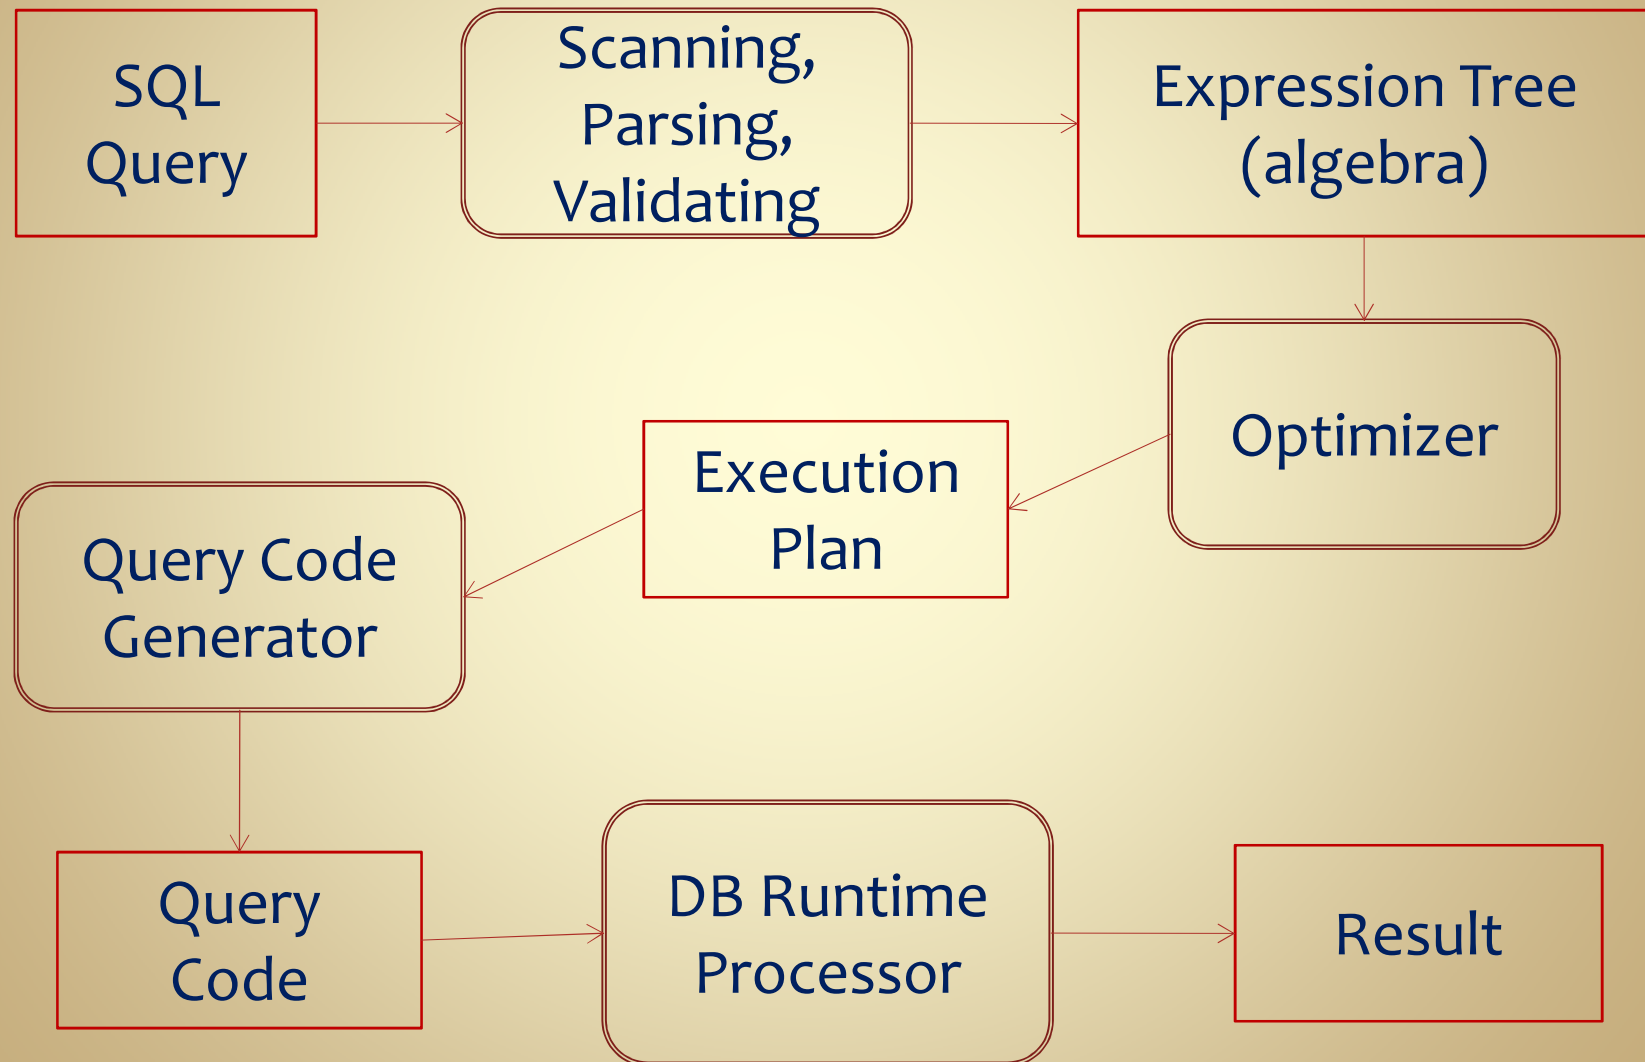
select * from Customer
where accountId >= 101 and accountId < 300



select * from Customer where lastName < 'D'



Query Processing



Query Optimization

- Query optimization determines the most efficient (or sufficiently efficient) process for executing the query
- Optimization reorganizes an expression tree for a query using algebraic transformation rules
- Information used:
 - implementation techniques for algebra operators
 - algebraic transformation rules
 - heuristic rules
 - cost estimates

Algebraic Transformations

- Commutative: $\sigma \bowtie x \cup \cap$
- Associative: $\bowtie x \cup \cap$
- Cascades of select:
 - $\sigma_{c1 \text{ AND } c2 \text{ AND } \dots \text{ AND } c_n}(R) \equiv \sigma_{c1}(\sigma_{c2}(\dots \sigma_{c_n}(R))\dots)$
 - selects can then be commuted
- Cascades of project:
 - $\pi_{L1}(\pi_{L2}(\dots \pi_{Ln}(R))\dots) \equiv \pi_{L1}(R)$
 - projects cannot be commuted

Algebraic Transformations

- Commuting select and project:

- $\pi_{A_1, A_2, \dots, A_n}(\sigma_c(R)) \equiv \sigma_c(\pi_{A_1, A_2, \dots, A_n}(R))$

valid when selection condition only involves attributes in projection list

- Converting select/cross-product into join:

- $\sigma_c(R \times S) \equiv R \bowtie_c S$

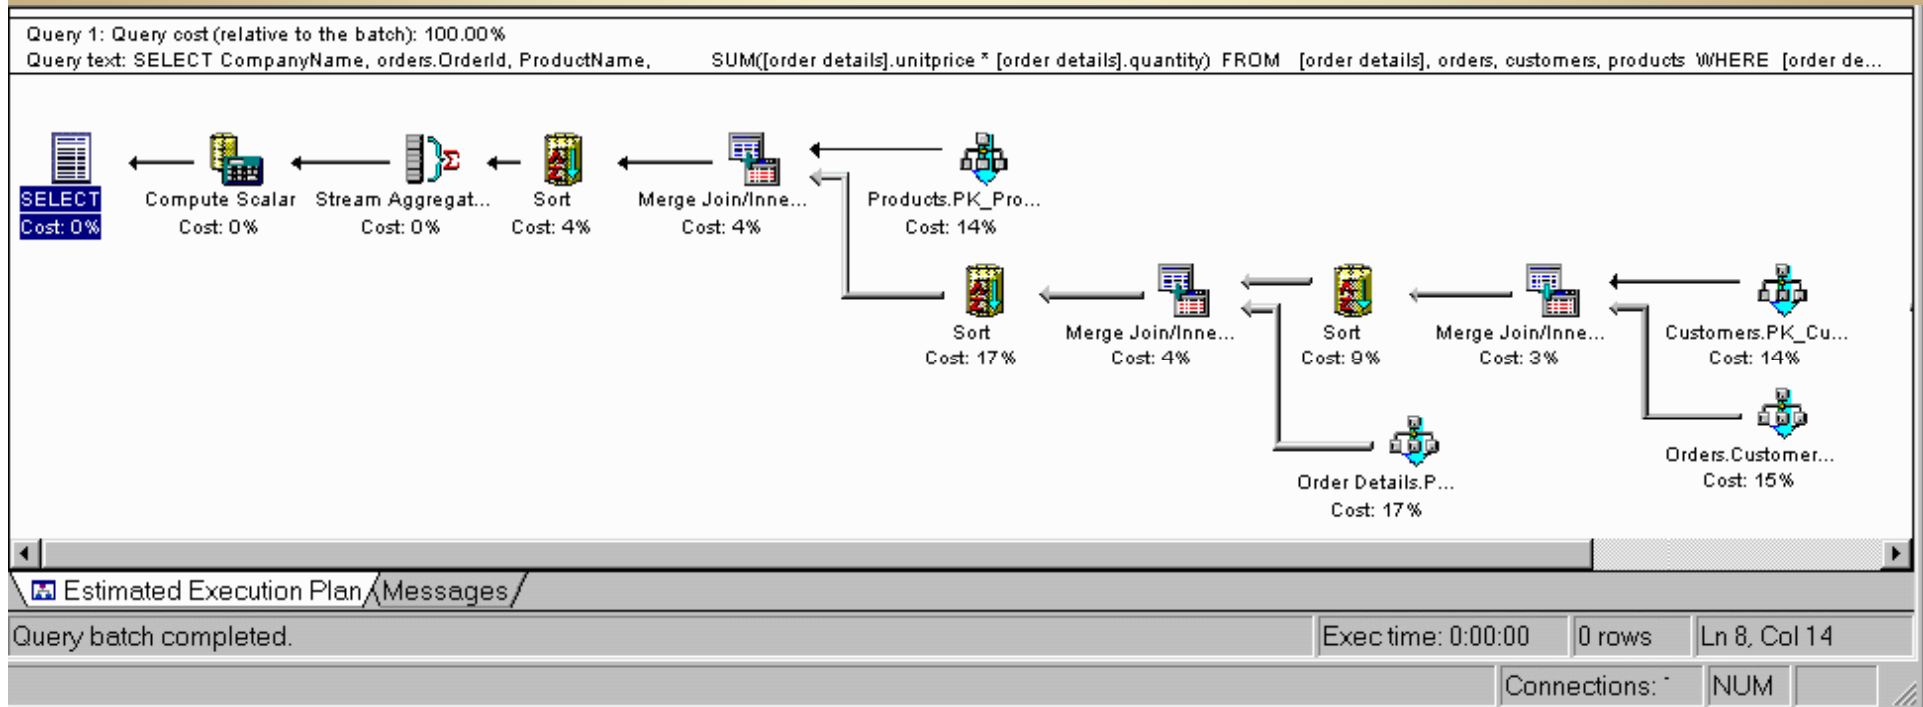
Algebraic Transformations

- Commuting select and join or cross-product:
 - $\sigma_c(R \bowtie S) \equiv \sigma_c(R) \bowtie S$
valid when selection condition involves only attributes in R
- Commuting project with join or cross-product:
 - $\pi_L(R \bowtie_c S) \equiv (\pi_{A_1, A_2, \dots, A_n}(R)) \bowtie_c (\pi_{B_1, B_2, \dots, B_n}(S))$
where $A_1, \dots, A_n \subseteq L$ are attributes of R,
and $B_1, \dots, B_n \subseteq L$ are attributes of S,
and $A_1, \dots, A_n, B_1, \dots, B_n$ are involved in C
(slightly more complicated if C contains attributes not in L)

Optimization Heuristics

- **Break up conjunctive select conditions**
into a cascade of selects
 - more flexibility in moving selects
- **Push selects as early as possible,**
using commutativity of select with other operators
- **Reorder sub-expressions such**
that most restrictive selects are done first
- **Combine select/cross-product into join**
- **Push projects as early as possible**
 - keep only necessary attributes

Example Query Plan



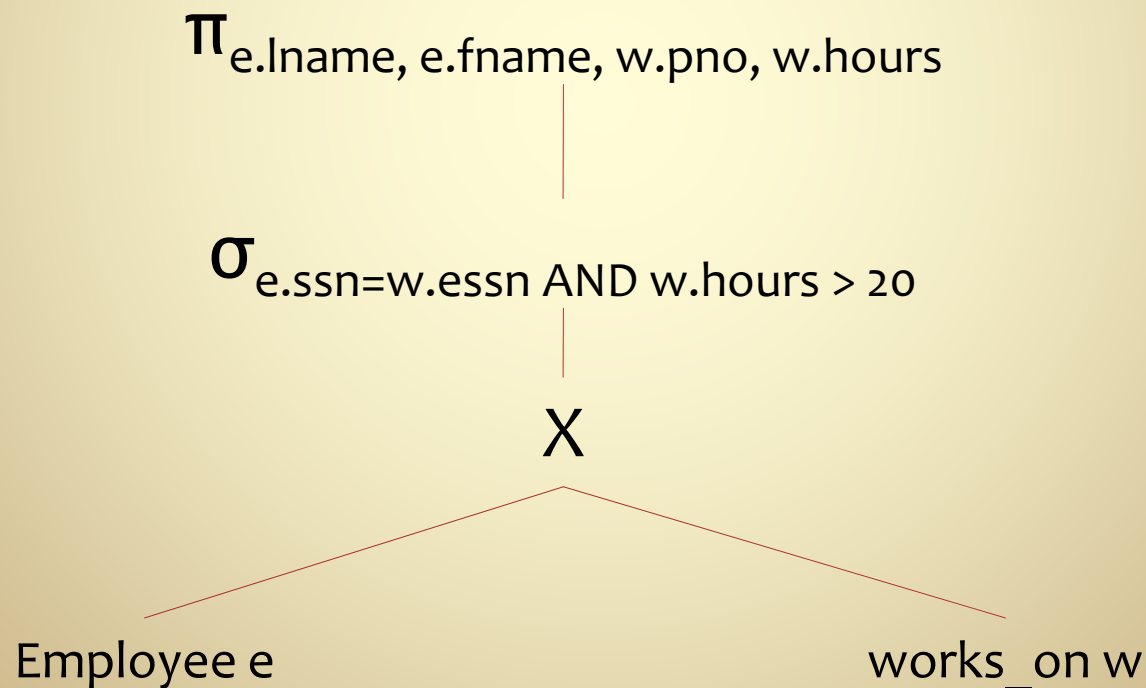
Microsoft SQL Server

Optimization Processor

- Enumerate the query plans
 - Apply algebraic transformations & heuristics to get all reasonable trees
- Estimate the cost of each plan
 - Account for size of relations, available indexes, information about file layout, info about value distributions ...
- Choose the best (fastest) plan

Example Query

```
select e.lname, e.fname, w.pno, w.hours  
from employee e, works_on w  
where e.ssn = w.essn and w.hours > 20;
```



Example Query

$\pi_{e.lname, e.fname, w.pno, w.hours}$

$\sigma_{w.hours > 20}$

$\sigma_{e.ssn=w.essn}$

X

Employee e

works_on w

Heuristic:
Conjunctive
select \rightarrow cascade
of selects

Note: selects
could commute

Example Query

$\pi_{e.lname, e.fname, w.pno, w.hours}$

Heuristic:
Combine select
and cross \rightarrow join

$\sigma_{w.hours > 20}$

$\bowtie_{e.ssn=w.essn}$

Employee e

works_on w

Example Query

$\pi_{e.lname, e.fname, w.pno, w.hours}$

Heuristic:
Push selects as
early as possible

$\bowtie_{e.ssn=w.essn}$

Employee e

$\sigma_{w.hours > 20}$

works_on w

Example Query

$\pi_{e.lname, e.fname, w.pno, w.hours}$

$\pi_{e.ssn, e.lname, e.fname, w.essn, w.pno, w.hours}$

Heuristic:
Push projects as
early as possible

$\bowtie_{e.ssn=w.essn}$

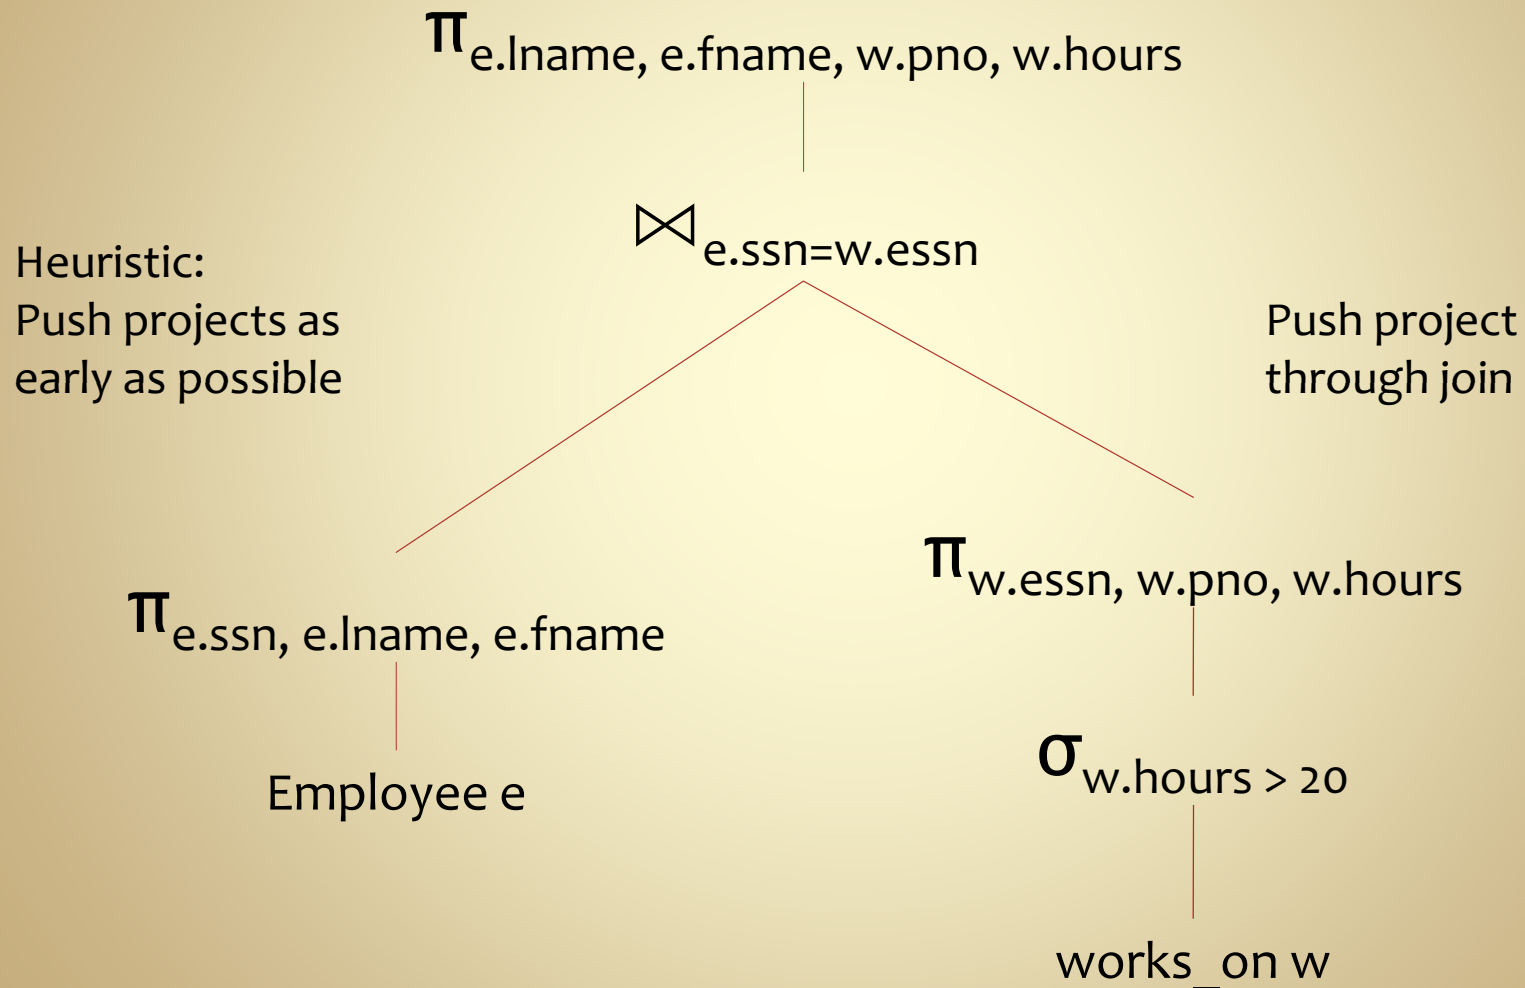
First cascade
projects to get
attributes
needed to push
through join

Employee e

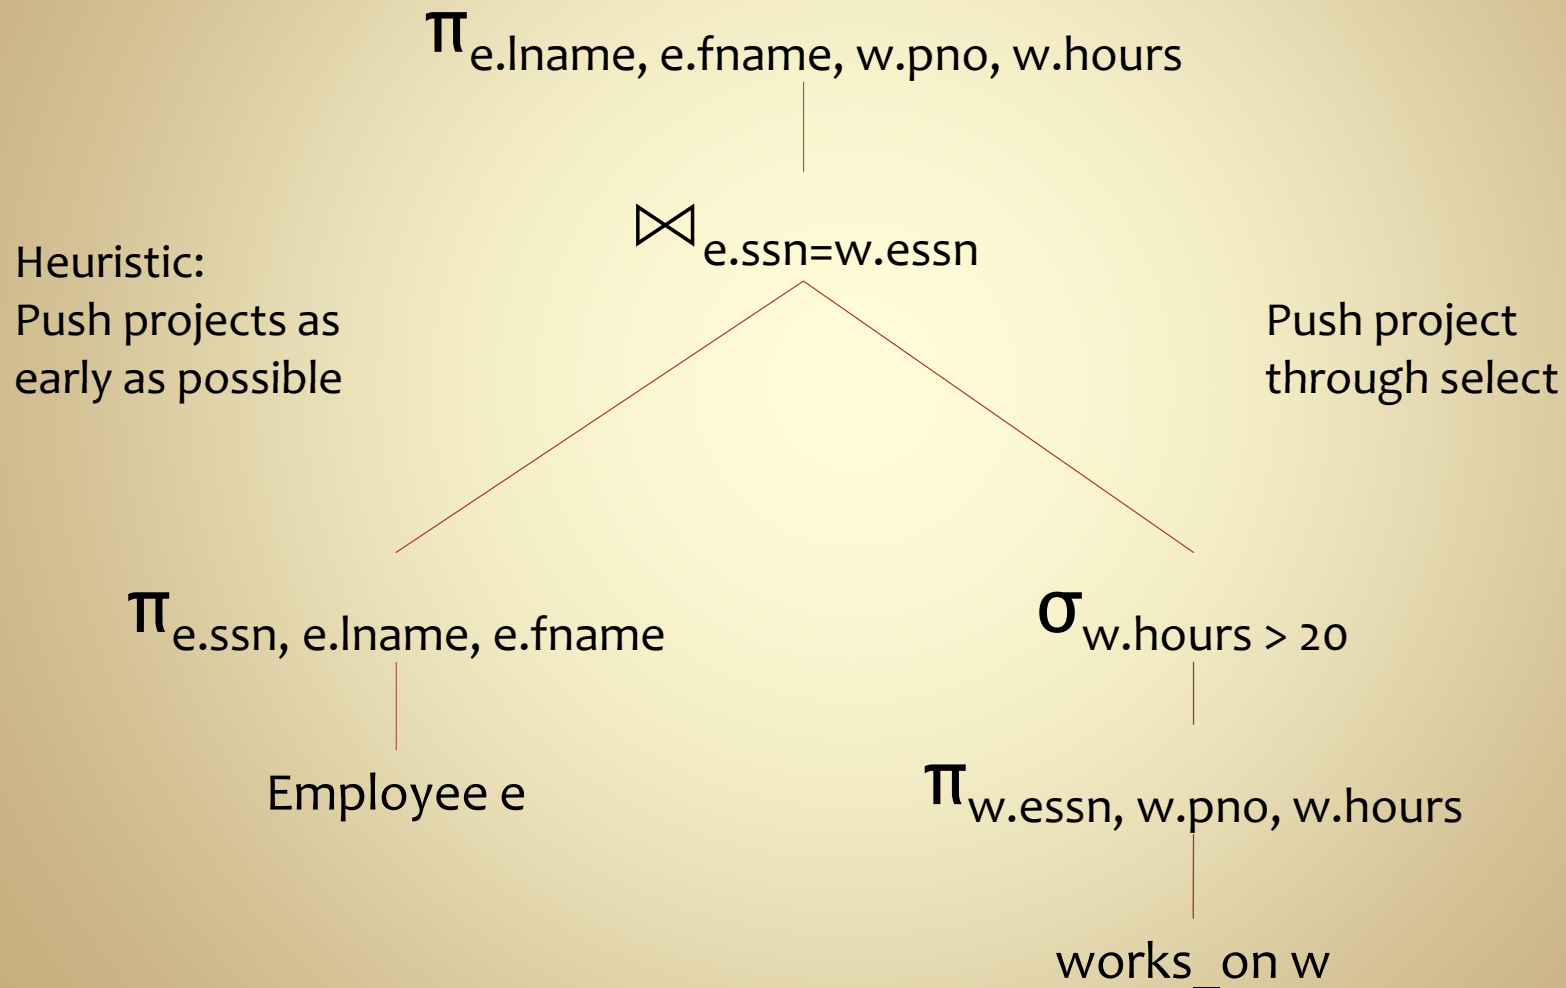
$\sigma_{w.hours > 20}$

works_on w

Example Query



Example Query



Example Query

- At this point we've generated seven different query plans for the same query
- Adding strategies/algorithms for implementing individual operations would add even more potential plans
- Next Step: Estimate the cost of executing each plan