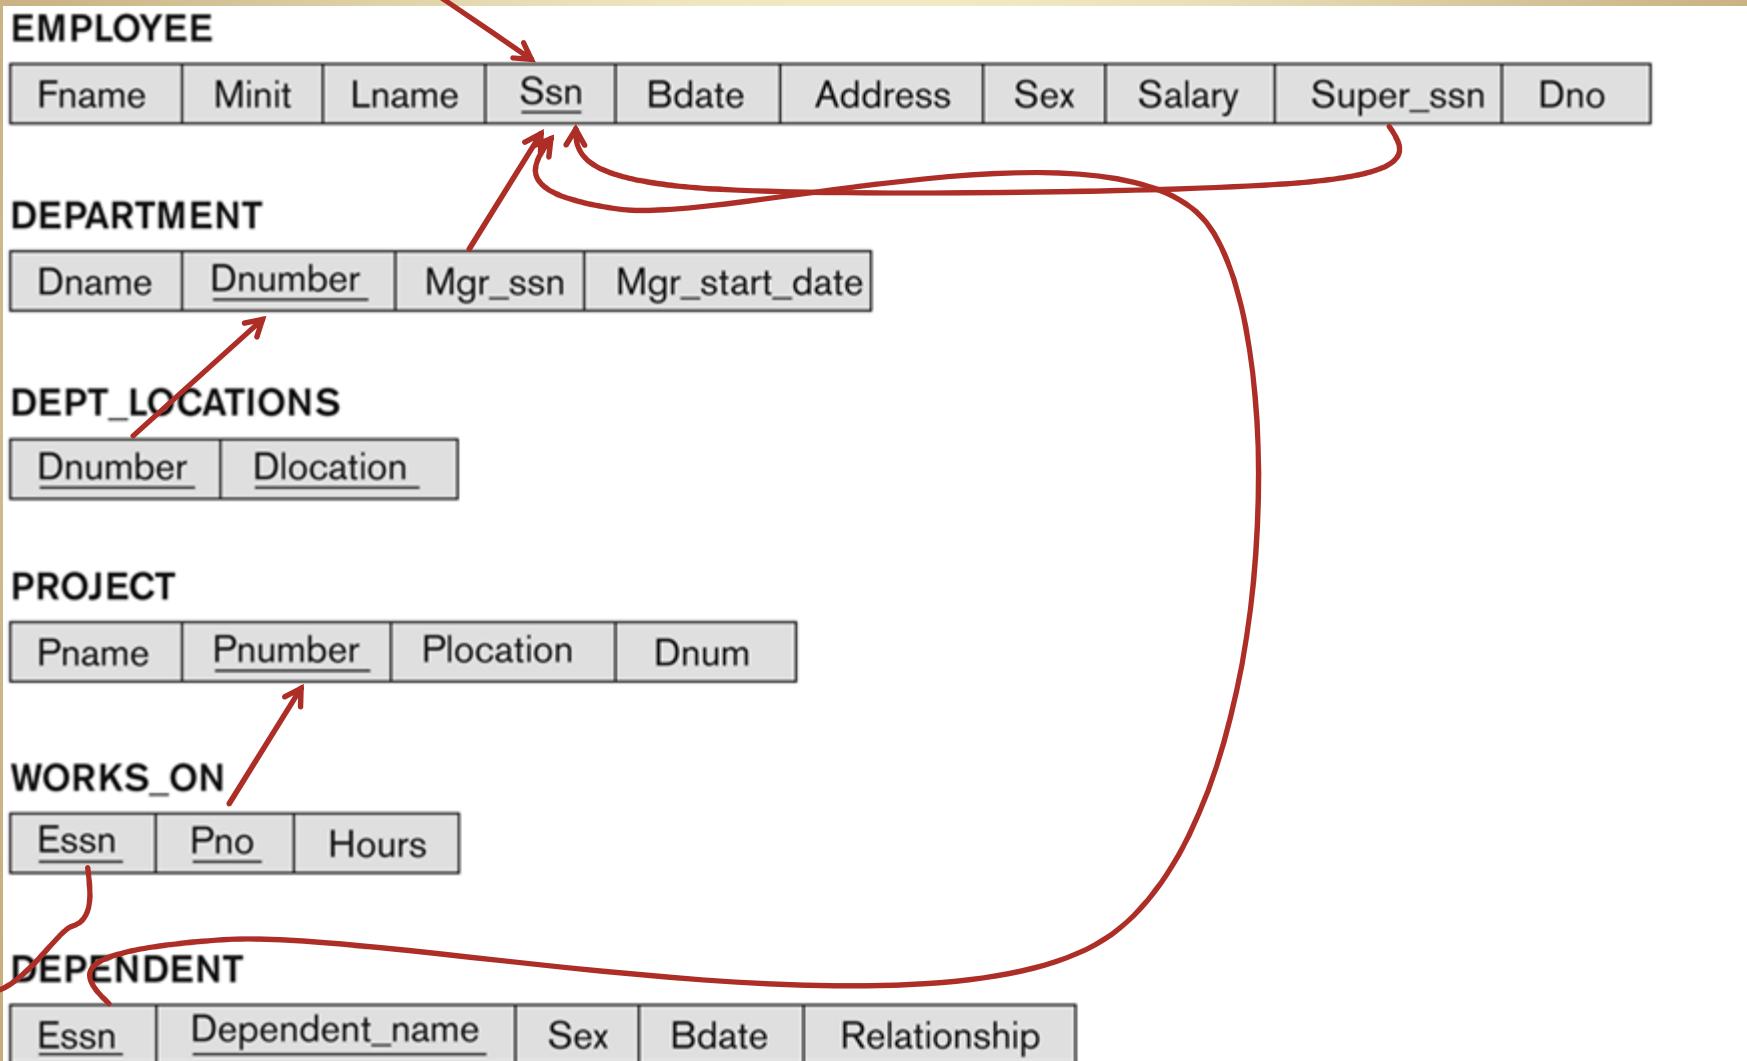


COMP163

Database Management Systems

Lecture 6 – Sections 6.1 and 6.3
The Relational Algebra

Example Schema



REVIEW

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

WORKS_ON

| Essn | Pno | Hours |
|-----------|-----|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | NULL |

PROJECT

| Pname | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

DEPENDENT

| Essn | Dependent_name | Sex | Bdate | Relationship |
|-----------|----------------|-----|------------|--------------|
| 333445555 | Alice | F | 1986-04-05 | Daughter |
| 333445555 | Theodore | M | 1983-10-25 | Son |
| 333445555 | Joy | F | 1958-05-03 | Spouse |
| 987654321 | Abner | M | 1942-02-28 | Spouse |
| 123456789 | Michael | M | 1988-01-04 | Son |
| 123456789 | Alice | F | 1988-12-30 | Daughter |
| 123456789 | Elizabeth | F | 1967-05-05 | Spouse |

Example State

State for in-class
exercise on
constraint violations

REVIEW

Queries

- Information is extracted from a database by writing *queries*
 - query inputs: relations
 - query result: relation
- queries are *read-only* operations
- The SQL command for queries select

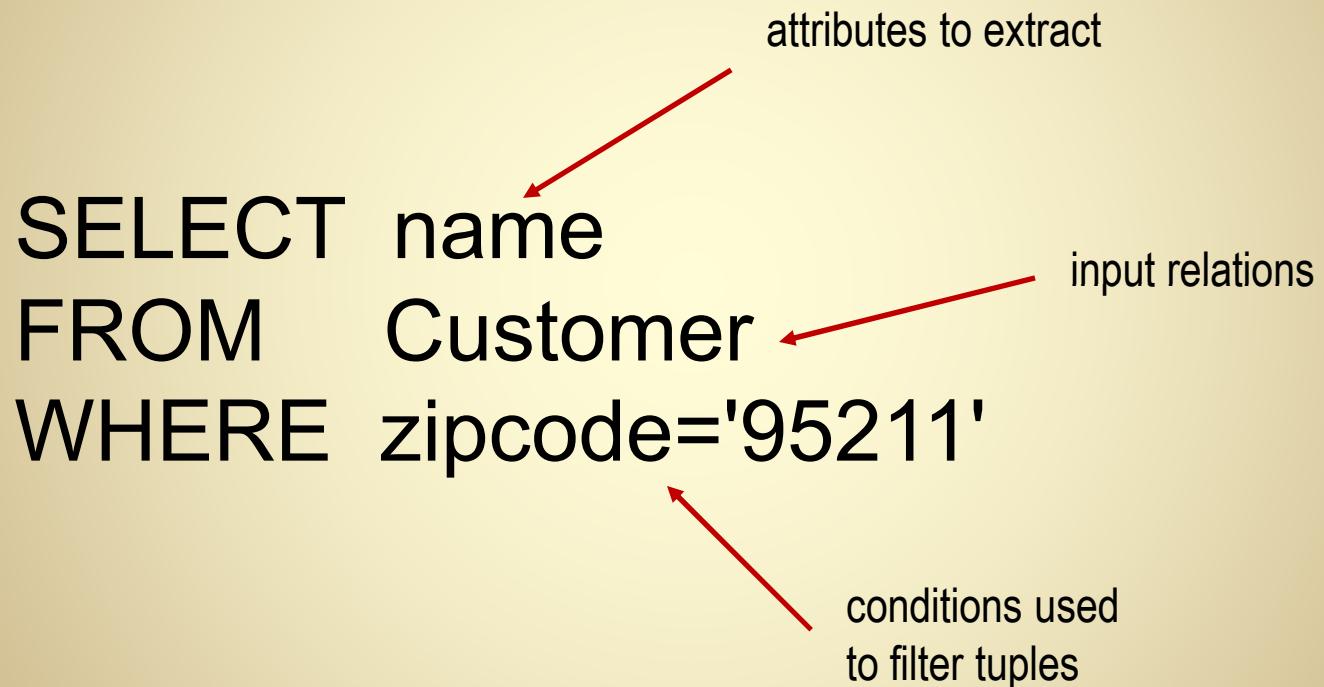
Example SQL Query

```
SELECT name  
FROM Customer  
WHERE zipcode='95211'
```

attributes to extract

input relations

conditions used
to filter tuples



The Relational Algebra

- The *relational algebra* defines mathematical operations on relations
 - provides a solid theory for dealing with queries
- These operations define the meaning of SQL queries
 - think in terms of the algebra operators
 - translate the algebra to SQL syntax
- Query optimization:
 - SQL queries are translated to expression trees containing algebraic operators
 - expression trees are reorganized to optimize the order of operations

Relational Algebra Operations

- set theoretic:
 - **union, intersection, difference, cross-product**
 - usual mathematical meaning
- relational database unique:
 - **select, project**: extraction from a single relation
 - **join**: combining two relations

The SELECT Operation

- SELECT extracts tuples from a relation
 - result has same relation schema as operand
- SELECT requires a *selection condition*
 - selection condition is a boolean expression to filter tuple values
- Syntax:

$$\sigma_{\text{<selection condition>}}(R)$$

- Selection condition may contain
AND, OR, NOT, =, <, ≤, >, ≥, ≠

SELECT Examples

$$r_2(\text{STORESTOCK}) = \left\{ \begin{array}{l} < "S002", "I065", 120 >, \\ < "S333", "I954", 198 >, \\ < "S047", "I099", 267 >, \\ < "S047", "I954", 300 > \end{array} \right\}$$

$$\sigma_{\text{StoreId} = "S047"}(\text{STORESTOCK}) \\ \left\{ \begin{array}{l} < "S047", "I099", 267 >, \\ < "S047", "I954", 300 > \end{array} \right\}$$

$$\sigma_{\text{quantity} < 200}(\text{STORESTOCK}) \\ \left\{ \begin{array}{l} < "S002", "I065", 120 >, \\ < "S333", "I954", 198 > \end{array} \right\}$$

The PROJECT Operation

- PROJECT extracts attributes from a relation
 - result schema attributes are a subset of the operand schema
- PROJECT requires a *attribute list*
- Syntax:

$$\pi_{\langle \text{attribute list} \rangle} (R)$$

- Duplicates are not kept in result
 - result is a relation, which is a set

PROJECT Examples

$$r_2(\text{STORESTOCK}) = \left\{ \begin{array}{l} < "S002", "I065", 120 >, \\ < "S333", "I954", 198 >, \\ < "S047", "I099", 267 >, \\ < "S047", "I954", 300 > \end{array} \right\}$$

$$\pi_{\text{StoreId}, \text{Item}}(\text{STORESTOCK})$$

$$\left\{ \begin{array}{l} < "S002", "I065" >, \\ < "S333", "I954" >, \\ < "S047", "I099" >, \\ < "S047", "I954" > \end{array} \right\}$$

PROJECT Examples

$r(STOCKITEM) =$

$$\left\{ \begin{array}{l} < "I075", "Ice Cream", \$1.49, \text{false} >, \\ < "I345", "Cupcakes", \$1.99, \text{false} >, \\ < "I333", "Twinkies", \$1.98, \text{false} > \end{array} \right\}$$

$\pi_{\text{Description}, \text{Price}}(STOREITEM)$

$$\left\{ \begin{array}{l} < "Ice Cream", \$1.49 >, \\ < "Cupcakes", \$1.99 >, \\ < "Twinkies", \$1.98 > \end{array} \right\}$$

Composing Operations

$r(STOCKITEM) =$

$$\left\{ \begin{array}{l} < "I075", "Ice Cream", \$1.49, \text{false} >, \\ < "I345", "Cupcakes", \$1.99, \text{false} >, \\ < "I333", "Twinkies", \$1.98, \text{false} > \end{array} \right\}$$

$\pi_{\text{Price}}(\sigma_{\text{Description}=\text{"Ice Cream"}}(\text{STOREITEM}))$

SELECT result:

$$\left\{ < "I075", "Ice Cream", \$1.49, \text{false} > \right\}$$

PROJECT result:

$$\left\{ < \$1.49 > \right\}$$

The JOIN Operation

- JOIN combines tuples from two tables based on values of related attributes (usually a FK)
- JOIN requires a join condition
 - boolean expression comparing attributes from each operand
- Syntax:

$$R \bowtie_{\text{<join condition>}} S$$

- The join condition may contain
AND, =, <, ≤, >, ≥, ≠

JOIN Examples

```
STORESTOCK( StoreId, Item, Quantity )
STOCKITEM( ItemId, Description, Price, Taxable )
```

$$r(STOCKITEM) = \left\{ \begin{array}{l} < "I075", "Ice Cream", \$1.49, \text{false} >, \\ < "I345", "Cupcakes", \$1.99, \text{false} >, \\ < "I333", "Twinkies", \$1.98, \text{false} > \end{array} \right\}$$
$$r(STORESTOCK) = \left\{ \begin{array}{l} < "S002", "I075", 120 >, \\ < "S047", "I333", 267 > \end{array} \right\}$$

STORESTOCK $\bowtie_{\langle \text{Item} = \text{ItemId} \rangle}$ STOCKITEM

JOIN Examples

RESULT(StoreId, Item, Quantity, ItemId, Description, Price, Taxable)

r(RESULT) =

$$\left\{ \begin{array}{l} < "S002", "I075", 120, "I075", "Ice Cream", \$1.49, \text{false} >, \\ < "S047", "I333", 267, "I333", "Twinkies", \$1.98, \text{false} > \end{array} \right\}$$

STORESTOCK $\bowtie_{\langle \text{Item} = \text{ItemId} \rangle}$ STOCKITEM

Writing Queries

STORESTOCK

| StoreId | Item | Quantity |
|---------|------|----------|
| S002 | I075 | 120 |
| S047 | I333 | 267 |
| S002 | I333 | 1200 |

Query:

Get the Description and Price
for all Items stocked
by Store S002

STOCKITEM

| ItemId | Description | Price | Taxable |
|--------|-------------|--------|---------|
| I075 | Ice Cream | \$1.49 | FALSE |
| I345 | Cup Cakes | \$1.99 | FALSE |
| I333 | Twinkies | \$1.98 | FALSE |

result only
includes tuples
with
certain ItemIds

result has these attributes

Writing Queries

STORESTOCK

| StoreId | Item | Quantity |
|---------|------|----------|
| S002 | I075 | 120 |
| S047 | I333 | 267 |
| S002 | I333 | 1200 |

STOCKITEM

| ItemId | Description | Price | Taxable |
|--------|-------------|--------|---------|
| I075 | Ice Cream | \$1.49 | FALSE |
| I345 | Cup Cakes | \$1.99 | FALSE |
| I333 | Twinkies | \$1.98 | FALSE |

Query:

Get the Description and Price
for all Items stocked
by Store S002

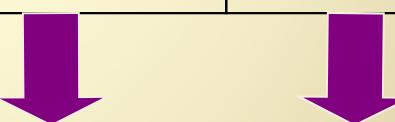
We need a join
to merge data
across relations



Writing Queries

STORESTOCK ⚡_{<Item = ItemId>} STOCKITEM

| StoreId | Item | Quantity | ItemId | Description | Price | Taxable |
|---------|------|----------|--------|-------------|--------|---------|
| S002 | I075 | 120 | I075 | Ice Cream | \$1.49 | FALSE |
| S047 | I333 | 267 | I333 | Twinkies | \$1.98 | FALSE |
| S002 | I333 | 1200 | I333 | Twinkies | \$1.98 | FALSE |



Query:

Get the Description and Price
for all Items stocked
by Store S002

Now we can select and project
to extract the information we want

Writing Queries

$R1 = \text{STORESTOCK} \bowtie_{\langle \text{Item} = \text{ItemId} \rangle} \text{STOCKITEM}$

| <u>StoreId</u> | <u>Item</u> | Quantity | <u>ItemId</u> | Description | Price | Taxable |
|----------------|-------------|----------|---------------|-------------|--------|---------|
| S002 | I075 | 120 | I075 | Ice Cream | \$1.49 | FALSE |
| S047 | I333 | 267 | I333 | Twinkies | \$1.98 | FALSE |
| S002 | I333 | 1200 | I333 | Twinkies | \$1.98 | FALSE |

$R2 = \sigma_{\text{StoreId} = "S002"} (R1)$

| <u>StoreId</u> | <u>Item</u> | Quantity | <u>ItemId</u> | Description | Price | Taxable |
|----------------|-------------|----------|---------------|-------------|--------|---------|
| S002 | I075 | 120 | I075 | Ice Cream | \$1.49 | FALSE |
| S002 | I333 | 1200 | I333 | Twinkies | \$1.98 | FALSE |

Writing Queries

$$R2 = \sigma_{StoreId = "Soo2"}(R1)$$

| StoreId | Item | Quantity | ItemId | Description | Price | Taxable |
|---------|------|----------|--------|-------------|--------|---------|
| Soo2 | I075 | 120 | I075 | Ice Cream | \$1.49 | FALSE |
| Soo2 | I333 | 1200 | I333 | Twinkies | \$1.98 | FALSE |


$$R3 = \pi_{\langle Description, Price \rangle}(R2)$$

| Description | Price |
|-------------|--------|
| Ice Cream | \$1.49 |
| Twinkies | \$1.98 |

Composing Queries

Since the operands and results of every query are relations, we can compose or chain queries.

$$R_1 = \text{STORESTOCK} \bowtie_{\langle \text{Item} = \text{ItemId} \rangle} \text{STOCKITEM}$$
$$R_2 = \sigma_{\text{StoreId} = "S002"} (R_1)$$
$$R_3 = \pi_{\langle \text{Description}, \text{Price} \rangle} (R_2)$$
$$\pi_{\langle \text{Description}, \text{Price} \rangle} (\sigma_{\text{StoreId} = "S002"} (\text{STORESTOCK} \bowtie_{\langle \text{Item} = \text{ItemId} \rangle} \text{STOCKITEM}))$$

SQL Queries

Relational algebra queries are easily translated to SQL queries
(more on this later)

$$\pi_{\langle \text{Description}, \text{Price} \rangle}(\sigma_{\text{StoreId} = "S002"} (\text{STORESTOCK} \bowtie_{\langle \text{Item} = \text{ItemId} \rangle} \text{STOCKITEM}))$$

```
SELECT Description, Price
FROM   STORESTOCK, STOCKITEM
WHERE StoreId='S002'
      AND STORESTOCK.Item = STOCKITEM.ItemId
```

EXERCISE 1: Schema

EMPLOYEE

| | | | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| | | | |
|-------|----------------|---------|----------------|
| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| | |
|----------------|-----------|
| <u>Dnumber</u> | Dlocation |
|----------------|-----------|

PROJECT

| | | | |
|-------|----------------|-----------|------|
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|

WORKS_ON

| | | |
|------|-----|-------|
| Essn | Pno | Hours |
|------|-----|-------|

DEPENDENT

| | | | | |
|------|----------------|-----|-------|--------------|
| Essn | Dependent_name | Sex | Bdate | Relationship |
|------|----------------|-----|-------|--------------|

EXERCISE 1: Queries

Express the following queries in the algebra:

1. First and last name of employees who have no supervisor.
2. First and last name of employees supervised by Franklin Wong.
3. Last name of employees who have dependents.
4. Last name of employees who have daughters.
5. Last name of employees in department 5 who work more than 10 hours/week on ProductX.
6. Last name of supervisors of employees in department 5 who work more than 10 hours/week on ProductX.

Section 6.2

Operations from Set Theory

Set Operations

- UNION: $R \cup S$
 - all tuples in either R or S
- INTERSECTION: $R \cap S$
 - all tuples in both R and S
- DIFFERENCE: $R - S$
 - all tuples in R but not in S
- CROSS-PRODUCT or CARTESIAN PRODUCT: $R \times S$
 - pair all tuples in R with all tuples in S

$\cup \cap -$
operands must be
union compatible
(same attribute types)

UNION: example

STORESTOCK

| StoreId | Item | Quantity |
|---------|------|----------|
| S002 | I075 | 120 |
| S047 | I333 | 267 |
| S002 | I333 | 267 |

WAREHOUSESTOCK

| StoreId | Item | Quantity |
|---------|------|----------|
| W998 | I075 | 1200 |
| W087 | I001 | 5000 |
| W222 | I188 | 11500 |
| W023 | I075 | 300 |

STORESTOCK \cup WAREHOUSESTOCK

| Id | Item | Quantity |
|------|------|----------|
| S002 | I075 | 120 |
| S047 | I333 | 267 |
| S002 | I333 | 267 |
| W998 | I075 | 1200 |
| W087 | I001 | 5000 |
| W222 | I188 | 11500 |
| W023 | I075 | 300 |

INTERSECTION: example

STORESTOCK

| StoreId | Item | Quantity |
|---------|------|----------|
| S002 | I075 | 120 |
| S047 | I333 | 267 |
| S002 | I333 | 267 |

WAREHOUSESTOCK

| StoreId | Item | Quantity |
|---------|------|----------|
| W998 | I075 | 1200 |
| W087 | I001 | 5000 |
| W222 | I188 | 11500 |
| W023 | I075 | 300 |

$$\pi_{\text{ItemID}}(\text{STORESTOCK}) \cap \pi_{\text{Item}}(\text{WAREHOUSESTOCK})$$

| Item |
|------|
| I075 |
| I333 |

∩

| Item |
|------|
| I075 |
| I001 |
| I188 |

=

| Item |
|------|
| I075 |

DIFFERENCE: example

STOCKITEM

| ItemId | Description | Price | Taxable |
|--------|-------------|--------|---------|
| I075 | Ice Cream | \$1.49 | FALSE |
| I345 | Cup Cakes | \$1.99 | FALSE |
| I333 | Twinkies | \$1.98 | FALSE |

STORESTOCK

| StoreId | Item | Quantity |
|---------|------|----------|
| S002 | I075 | 120 |
| S047 | I333 | 267 |
| S002 | I333 | 267 |

$$\pi_{\text{ItemId}}(\text{STOCKITEM}) - \pi_{\text{Item}}(\text{STORESTOCK})$$



CROSS PRODUCT: example

STOCKITEM

| ItemId | Description | Price | Taxable |
|--------|-------------|--------|---------|
| I075 | Ice Cream | \$1.49 | FALSE |
| I345 | Cup Cakes | \$1.99 | FALSE |
| I333 | Twinkies | \$1.98 | FALSE |

STORESTOCK

| StoreId | Item | Quantity |
|---------|------|----------|
| S002 | I075 | 120 |
| S047 | I333 | 267 |

STOCKITEM × STORESTOCK

| ItemId | Description | Price | Taxable | StoreId | Item | Quantity |
|--------|-------------|--------|---------|---------|------|----------|
| I075 | Ice Cream | \$1.49 | FALSE | S002 | I075 | 120 |
| I345 | Cup Cakes | \$1.99 | FALSE | S002 | I075 | 120 |
| I333 | Twinkies | \$1.98 | FALSE | S002 | I075 | 120 |
| I075 | Ice Cream | \$1.49 | FALSE | S047 | I333 | 267 |
| I345 | Cup Cakes | \$1.99 | FALSE | S047 | I333 | 267 |
| I333 | Twinkies | \$1.98 | FALSE | S047 | I333 | 267 |

CROSS PRODUCT

- CROSS PRODUCT is also called CROSS JOIN
 - a JOIN with no join condition
- XPROD brings together all possible information from two relations
- XPROD can result in very large relations
 - if R has i tuples and S has j tuples,
then $R \times S$ has $i*j$ tuples
- In real applications, avoid XPROD whenever possible
 - Database equivalent of a brute-force nested loop
 - massive unnecessary memory usage

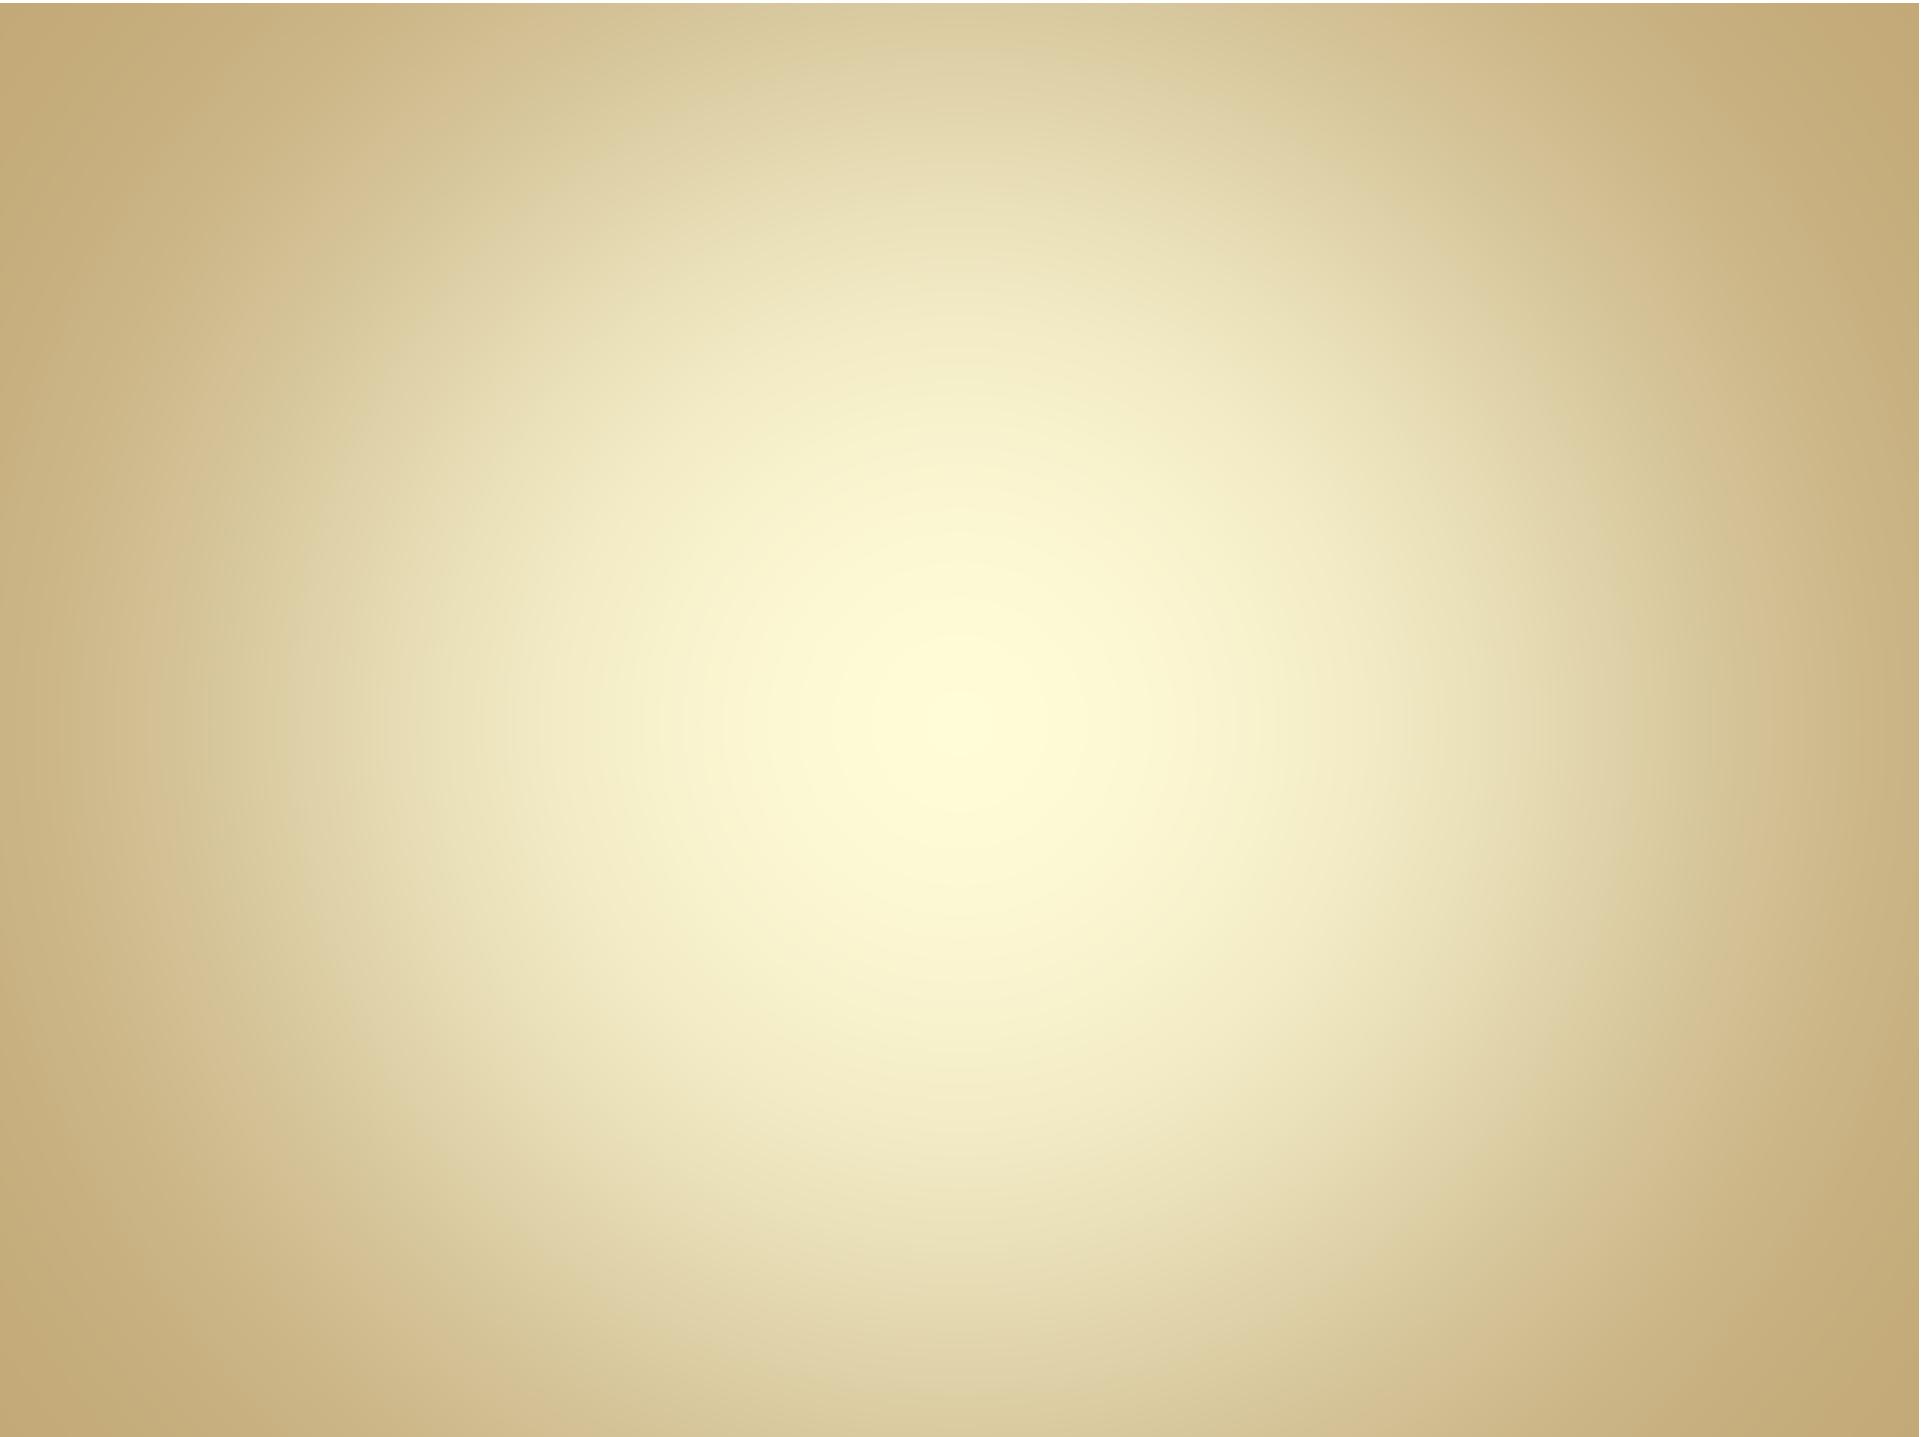
JOIN = XPROD and SELECT

STORESTOCK $\bowtie_{\langle \text{Item} = \text{ItemId} \rangle}$ STOCKITEM

| ItemId | Description | Price | Taxable | StoreId | Item | Quantity |
|--------|-------------|--------|---------|---------|------|----------|
| I075 | Ice Cream | \$1.49 | FALSE | S002 | I075 | 120 |
| I333 | Twinkies | \$1.98 | FALSE | S047 | I333 | 267 |

$\sigma_{\text{ItemId}=\text{Item}}$ (STOCKITEM \times STORESTOCK)

| ItemId | Description | Price | Taxable | StoreId | Item | Quantity |
|--------|-------------|--------|---------|---------|------|----------|
| I075 | Ice Cream | \$1.49 | FALSE | S002 | I075 | 120 |
| I345 | Cup Cakes | \$1.99 | FALSE | S002 | I075 | 120 |
| I333 | Twinkies | \$1.98 | FALSE | S002 | I075 | 120 |
| I075 | Ice Cream | \$1.49 | FALSE | S047 | I333 | 267 |
| I345 | Cup Cakes | \$1.99 | FALSE | S047 | I333 | 267 |
| I333 | Twinkies | \$1.98 | FALSE | S047 | I333 | 267 |



Declarative Programming

Queries are **declarative**, not **procedural**.

These queries specify the information we want, in terms of mathematical operations ... but not how to compute it.

Does not say “compute a cross-product, then select ...”, rather give me something “equivalent to cross-product, select ...”

$$\pi_{\langle \text{Description}, \text{Price} \rangle}(\sigma_{\text{StoreId} = "S002"} (\text{STORESTOCK} \bowtie_{\langle \text{Item} = \text{ItemId} \rangle} \text{STOCKITEM}))$$

SELECT Description, Price
FROM STORESTOCK, STOCKITEM
WHERE StoreId='S002'
AND STORESTOCK.Item = STOCKITEM.ItemId

π

σ condition

\bowtie condition

Complete Set of Operations

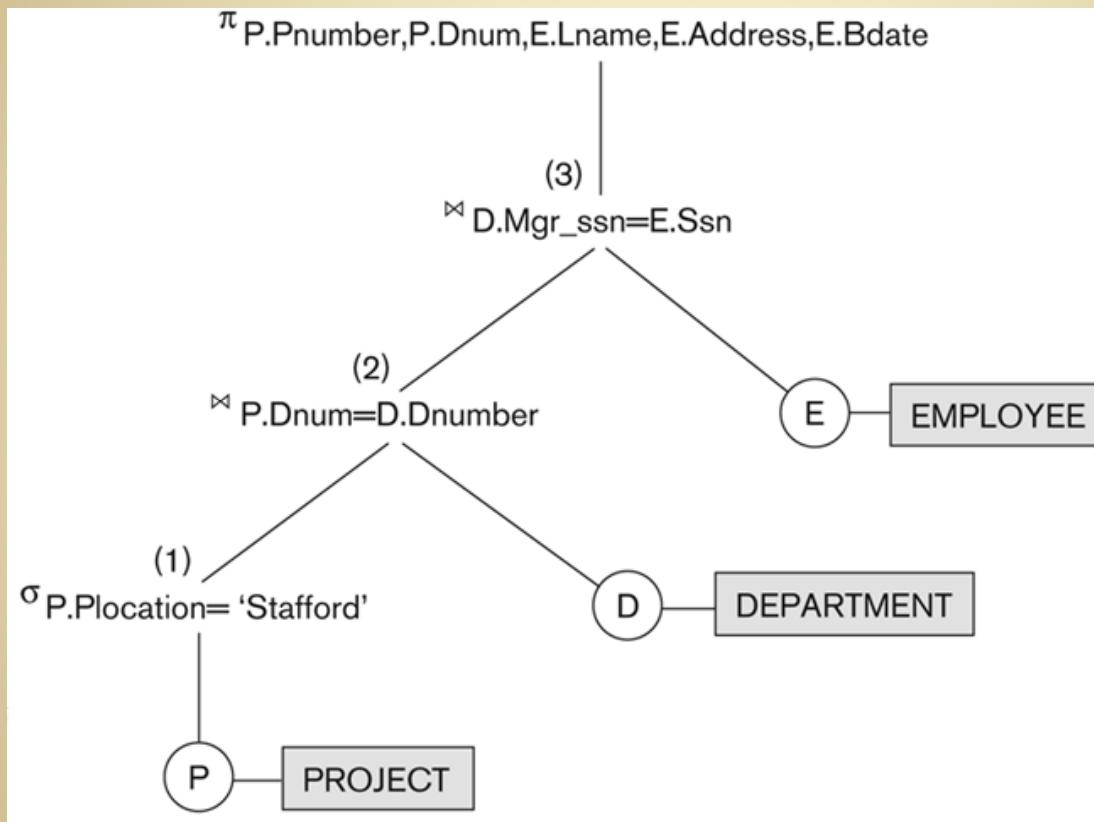
- A complete set of operations is a subset of all operations that is sufficient to express all queries expressible by all operators
 - any operators outside a complete set can be expressed as combinations of operators in the complete set
- Complete Set: $\{ \sigma, \pi, U, -, \times \}$
- Operators outside a complete set are not necessary, but they are generally convenient
 - JOIN: $A \bowtie_C B = \sigma_C(A \times B)$
 - INTERSECTION: $A \cap B = (A \cup B) - ((A - B) \cup (B - A))$

Additional Operators

- There are some queries in commercial languages (i.e. SQL) that cannot be expressed in the basic relational algebra
- A few additional operations handle most of these
 - **aggregation:** apply a function to a collection of values
 - **OUTER JOINS** and **OUTER UNIONs:** null values are used to keep data that would otherwise be discarded

| Operation | Purpose | Notation |
|---------------------------|--|--|
| SELECT → | Selects all tuples that satisfy the selection condition from a relation R . | $\sigma_{\text{selection condition}}(R)$ |
| PROJECT → | Produces a new relation with only some of the attributes of R , and removes duplicate tuples. | $\pi_{\text{attribute list}}(R)$ |
| THETA JOIN → | Produces all combinations of tuples from R_1 and R_2 that satisfy the join condition. | $R_1 \bowtie_{\text{join condition}} R_2$ |
| EQUIJOIN | Produces all the combinations of tuples from R_1 and R_2 that satisfy a join condition with only equality comparisons. | $R_1 \bowtie_{\text{join condition}} R_2$, OR $R_1 \bowtie_{(\text{join attributes 1}), (\text{join attributes 2})} R_2$ |
| NATURAL JOIN → | Same as EQUIJOIN except that the join attributes of R_2 are not included in the resulting relation; if the join attributes have the same names, they do not have to be specified at all. | $R_1 *_{\text{join condition}} R_2$, OR $R_1 *_{(\text{join attributes 1}), (\text{join attributes 2})} R_2$ OR $R_1 * R_2$ |
| UNION → | Produces a relation that includes all the tuples in R_1 or R_2 or both R_1 and R_2 ; R_1 and R_2 must be union compatible. | $R_1 \cup R_2$ |
| INTERSECTION → | Produces a relation that includes all the tuples in both R_1 and R_2 ; R_1 and R_2 must be union compatible. | $R_1 \cap R_2$ |
| DIFFERENCE → | Produces a relation that includes all the tuples in R_1 that are not in R_2 ; R_1 and R_2 must be union compatible. | $R_1 - R_2$ |
| CARTESIAN PRODUCT → | Produces a relation that has the attributes of R_1 and R_2 and includes as tuples all possible combinations of tuples from R_1 and R_2 . | $R_1 \times R_2$ |
| DIVISION | Produces a relation $R(X)$ that includes all tuples $t[X]$ in $R_1(Z)$ that appear in R_1 in combination with every tuple from $R_2(Y)$, where $Z = X \cup Y$. | $R_1(Z) \div R_2(Y)$ |

Query Trees



Query trees are representations of queries that can be manipulated by **query optimizers**, according to mathematical properties of the operators.
(Chapter 15)

EXERCISE 2: Schema

EMPLOYEE

| | | | | | | | | | |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|
| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|-------|---------|-----|--------|-----------|-----|

DEPARTMENT

| | | | |
|-------|----------------|---------|----------------|
| Dname | <u>Dnumber</u> | Mgr_ssn | Mgr_start_date |
|-------|----------------|---------|----------------|

DEPT_LOCATIONS

| | |
|----------------|-----------|
| <u>Dnumber</u> | Dlocation |
|----------------|-----------|

PROJECT

| | | | |
|-------|----------------|-----------|------|
| Pname | <u>Pnumber</u> | Plocation | Dnum |
|-------|----------------|-----------|------|

WORKS_ON

| | | |
|------|-----|-------|
| Essn | Pno | Hours |
|------|-----|-------|

DEPENDENT

| | | | | |
|------|----------------|-----|-------|--------------|
| Essn | Dependent_name | Sex | Bdate | Relationship |
|------|----------------|-----|-------|--------------|

EXERCISE 2: Queries

1. First and last names
of all department managers.

2. Salaries of all employees
who have worked on the Reorganization project.

3. SSN of all employees who have worked on a project
that is controlled by a department different
than the department that they are assigned to.

4. Last name of all employees who are not married.

Exercise 3: Schema

AIRPORT

| | | | |
|--------------|------|------|-------|
| Airport_code | Name | City | State |
|--------------|------|------|-------|

FLIGHT

| | | |
|---------------|---------|----------|
| Flight_number | Airline | Weekdays |
|---------------|---------|----------|

FLIGHT_LEG

| | | | | | |
|---------------|------------|------------------------|--------------------------|----------------------|------------------------|
| Flight_number | Leg_number | Departure_airport_code | Scheduled_departure_time | | |
| | | | | Arrival_airport_code | Scheduled_arrival_time |

LEG_INSTANCE

| | | | | | | |
|---------------|------------|------------------------|---------------------------|----------------------|--------------|--|
| Flight_number | Leg_number | Date | Number_of_available_seats | Airplane_id | | |
| | | Departure_airport_code | Departure_time | Arrival_airport_code | Arrival_time | |

FARE

| | | | |
|---------------|-----------|--------|--------------|
| Flight_number | Fare_code | Amount | Restrictions |
|---------------|-----------|--------|--------------|

AIRPLANE_TYPE

| | | |
|--------------------|-----------|---------|
| Airplane_type_name | Max_seats | Company |
|--------------------|-----------|---------|

CAN_LAND

| | |
|--------------------|--------------|
| Airplane_type_name | Airport_code |
|--------------------|--------------|

AIRPLANE

| | | |
|-------------|-----------------------|---------------|
| Airplane_id | Total_number_of_seats | Airplane_type |
|-------------|-----------------------|---------------|

SEAT_RESERVATION

| | | | | | |
|---------------|------------|------|-------------|---------------|----------------|
| Flight_number | Leg_number | Date | Seat_number | Customer_name | Customer_phone |
|---------------|------------|------|-------------|---------------|----------------|

EXERCISE 3: Queries

1. List all airplane types that can land at any airport in San Francisco.
2. List the ids and number of seats for all airplanes that can land at any airport in Chicago.
3. List the name and phone number of all customers with a seat reserved on a flight that leaves Chicago O'Hara airport (ORD) on October 31, 2010.
4. List all airlines that have seats available for flights leaving Los Angeles (LAX) on September 25, 2010.
5. List all airlines that operate at San Jose International Airport (SJC).

RENAME

- Rename the attributes of a relation or change the relation name
- The general RENAME operation ρ :
 - $\rho_{S(B_1, B_2, \dots, B_n)}(R)$ changes both:
 - the relation name to S , and
 - the column (attribute) names to B_1, B_1, \dots, B_n
 - $\rho_S(R)$ changes:
 - the *relation name* only to S
 - $\rho_{(B_1, B_2, \dots, B_n)}(R)$ changes:
 - the *column (attribute) names* only to B_1, B_1, \dots, B_n

RENAME (shorthand)

- shorthand for renaming attributes :

$$R1 \leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{SALARY}} (\text{DEP5_EMPS})$$

R1 has same attribute names as DEP5_EMPS

$$R2 \leftarrow \rho_{(F,L,S)} (\pi_{\text{FNAME}, \text{LNAME}, \text{SALARY}} (\text{DEP5_EMPS}))$$

R2 has attributes renamed to F, L and S

→ R3 (F,L,S) $\leftarrow \pi_{\text{FNAME}, \text{LNAME}, \text{SALARY}} (\text{DEP5_EMPS})$

R3 also has attributes renamed to F, L and S

Aggregate Functions

- Aggregate functions apply mathematical operations to a collection of values.
- Examples:
 - compute average salary for all employees
 - compute maximum salary for all managers
 - count number of airports
- Common aggregate functions:
 - SUM
 - AVERAGE
 - MAX
 - MIN
 - COUNT

Aggregate Function

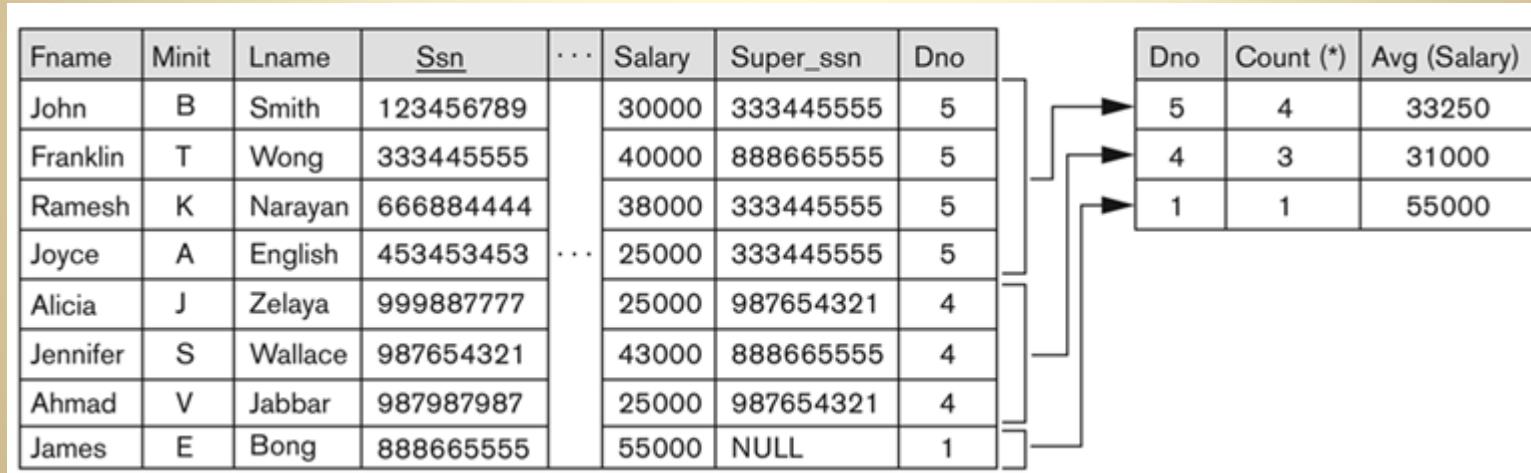
- Symbol for aggregate functions is \mathcal{F}
 - $\mathcal{F}_{\text{MAX Salary}}(\text{EMPLOYEE})$
retrieves the maximum salary value from the EMPLOYEE relation
 - $\mathcal{F}_{\text{MIN Salary}}(\text{EMPLOYEE})$
retrieves the minimum Salary value from the EMPLOYEE relation
 - $\mathcal{F}_{\text{SUM Salary}}(\text{EMPLOYEE})$
retrieves the sum of the Salary from the EMPLOYEE relation
 - $\mathcal{F}_{\text{COUNT SSN}, \text{AVERAGE Salary}}(\text{EMPLOYEE})$
computes the number of employees and their average salary

Grouping with Aggregation

- Grouping can be combined with Aggregate Functions
- Example:
 - For each department, retrieve the DNO, COUNT SSN, and AVERAGE SALARY
 - $\text{DNO } \mathcal{F}_{\text{COUNT SSN, AVERAGE Salary}} (\text{EMPLOYEE})$
- Group employees by DNO,
then compute aggregate functions on each group

Grouping with Aggregation

DNO \mathcal{F} COUNT SSN, AVERAGE Salary (EMPLOYEE)



JOIN VARIANTS

- Equijoin
- Natural join
- Outer join
- Outer union

THETA JOIN

- Theta Join is the general case

$$\text{STORESTOCK} \bowtie_{\langle \text{Item} = \text{ItemId} \rangle} \text{STOCKITEM}$$

- The join condition is called *theta*
- Join condition can be any Boolean expression on the attributes of R and S
 - $R.A_i < S.B_j \text{ AND } (R.A_k = S.B_l \text{ OR } R.A_p < S.B_q)$
- Most join conditions involve one or more equality conditions, connected with AND
 - $R.A_i = S.B_j \text{ AND } R.A_k = S.B_l \text{ AND } R.A_p = S.B_q$

EQUIJOIN

- The most common use of join involves join conditions with *equality comparisons* only
- When the only comparison operator used is $=$, we call it an EQUIJOIN
 - In the result of an EQUIJOIN we always have one or more pairs of attributes (whose names need not be identical) that have identical values in every tuple.

NATURAL JOIN

- NATURAL JOIN removes the superfluous attributes in an EQUIJOIN
 - The standard definition of natural join requires that the two join attributes, or each pair of corresponding join attributes, have *the same name* in both relations
 - If this is not the case, a renaming operation is applied first.
- Operator for NATURAL JOIN is *

R * S joins R and S with equality tests on all common attributes, then removes duplicate attributes

NATURAL JOIN

- $Q \leftarrow R(A,B,C,D) * S(C,D,E)$
 - The implicit join condition includes each pair of attributes with the same name, connected by AND
 - $R \bowtie_{R.C = S.C \text{ AND } R.D = S.D} S$
 - Result keeps only one attribute of each such pair:
 - $Q(A,B,C,D,E)$

OUTER JOIN

- THETA JOIN, NATURAL JOIN and EQUIJOIN, eliminate tuples without a *matching* (related) tuple
 - Tuples with null in the join attributes are also eliminated
 - This amounts to loss of information.
- OUTER JOINs keep the unmatched tuples and match them with NULLs
 - LEFT OUTER JOIN ($=\bowtie$): keep all tuples from left relation
 - RIGHT OUTER JOIN ($\bowtie=$): keep all tuples from right relation
 - FULL OUTER JOIN ($=\bowtie=$): keep all tuples from both relations

OUTER JOIN EXAMPLE

S

| A | B |
|---|---|
| 1 | 4 |
| 2 | 5 |
| 3 | 6 |

T

| C | D |
|---|----|
| 5 | 8 |
| 5 | 9 |
| 7 | 10 |

$R_1 = S \bowtie_{B=C} T$ (left join)

| A | B | C | D |
|---|---|------|------|
| 1 | 4 | NULL | NULL |
| 2 | 5 | 5 | 8 |
| 2 | 5 | 5 | 9 |
| 3 | 6 | NULL | NULL |

$R_2 = S \bowtie_{B=C} T$ (right join)

| A | B | C | D |
|------|------|---|----|
| 2 | 5 | 5 | 8 |
| 2 | 5 | 5 | 9 |
| NULL | NULL | 7 | 10 |

$R_3 = S \bowtie_{B=C} T$ (full join)

| A | B | C | D |
|------|------|------|------|
| 1 | 4 | NULL | NULL |
| 2 | 5 | 5 | 8 |
| 2 | 5 | 5 | 9 |
| 3 | 6 | NULL | NULL |
| NULL | NULL | 7 | 10 |

notice that $R_3 = R_1 \cup R_2$

OUTER UNION

- OUTER UNION computes the union of tuples from two relations when the relations are not type compatible
 - Tuples are included for all tuples from either relation that do not match a tuple in the other relation
 - Tuples are also included for pairs of tuples from each attribute that do match on common attributes
 - Remaining attribute values from both relations are kept
 - Missing values are replaced with NULL

OUTER UNION

- Example: OUTER UNION of
STUDENT(Name, SSN, Department, Advisor) and
INSTRUCTOR(Name, SSN, Department, Rank)
 - Tuples from the two relations are matched based on having the same combination of values of the shared attributes:
Name, SSN, Department
 - The result relation will have the following attributes:
STUDENT_OR_INSTRUCTOR (Name, SSN, Department, Advisor, Rank)
 - If a student is also an instructor, both Advisor and Rank will have a value; otherwise, one of these two attributes will be null

Recursive Closure

- Example:

- Find all employees who work under (directly or indirectly) James Borg

- This operation is applied to a **recursive relationship**.

- simplified schema:

- EMPLOYEE(name, ID, superID)

- superId is the ID of an employees supervisor
(may be NULL)

Recursive Closure

- Example:

- Find all employees who work under Miro

EMPLOYEE

| Name | ID | superID |
|--------|----|---------|
| Ang | 1 | 5 |
| Wassim | 2 | 3 |
| Iria | 3 | NULL |
| Jung | 4 | 6 |
| Miro | 5 | 3 |
| Sally | 6 | 5 |

RECURSIVE CLOSURE

- Solution: Retrieve employees at each level and build solution with union

$T_1 = \pi_{ID}(\sigma_{Name='Miro'} EMPLOYEE)$

Repeat until T_1 does not change:

$T_2 = \pi_{ID}(EMPLOYEE \bowtie_{EMPLOYEE.superID=T_1.ID} T_1)$

$T_1 = T_1 \cup T_2$

- cannot, in general, specify recursive closure without iteration (loops)
 - The SQL3 standard includes syntax for recursive closure.

DIVISION

- Example:
 - Retrieve names of employees who work on *all* projects that John Smith works on.
- $R = S \div T$
 $x \in R$, iff for every $y \in T$, $\langle x, y \rangle \in S$
- The tuples in R come from the tuples in S that match every tuple in T

DIVISION Example

| R | S | T |
|----|----|---|
| A | A | B |
| a1 | b1 | |
| a2 | b1 | |
| a3 | b1 | |
| a4 | b1 | |
| a1 | b2 | |
| a3 | b2 | |
| a2 | b3 | |
| a3 | b3 | |
| a4 | b3 | |
| a1 | b4 | |
| a2 | b4 | |
| a3 | b4 | |

$$T = R \div S$$

T has all attributes of R that are not in S.

Find tuples in R that match every tuple in S.
Project the attributes that are not in S.

DIVISION Example

SSN_PNOS

| Essn | Pno |
|-----------|-----|
| 123456789 | 1 |
| 123456789 | 2 |
| 666884444 | 3 |
| 453453453 | 1 |
| 453453453 | 2 |
| 333445555 | 2 |
| 333445555 | 3 |
| 333445555 | 10 |
| 333445555 | 20 |
| 999887777 | 30 |
| 999887777 | 10 |
| 987987987 | 10 |
| 987987987 | 30 |
| 987654321 | 30 |
| 987654321 | 20 |
| 888665555 | 20 |

SMITH_PNOS

| Pno |
|-----|
| 1 |
| 2 |

SSNS

| Ssn |
|-----------|
| 123456789 |
| 453453453 |

$$SSNS = SSN_PNOS \div SMITH_PNOS$$

Retrieve names of employees
who work on
all projects that John Smith
works on.

DIVISION

- $R = S \div T$ can be computed as

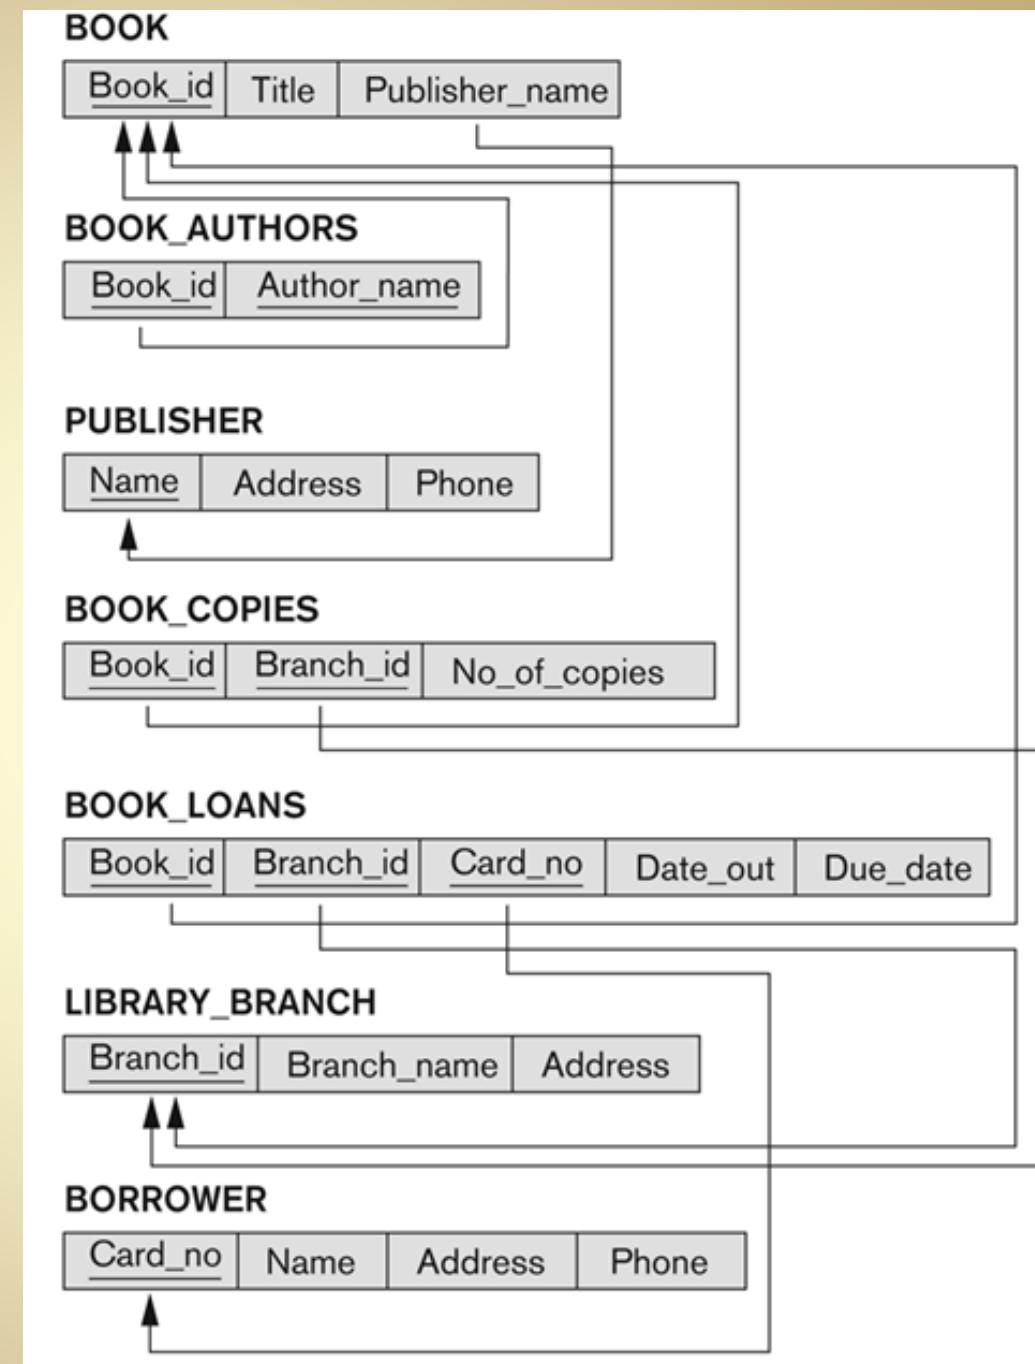
$$R_1 = \pi_Y(S)$$

$$R_2 = \pi_Y((T \times R_1) - S)$$

$$R = R_1 - R_2$$

EXERCISE 4:

Schema



EXERCISE 4: Queries

1. Count the number of overdue books.
2. How many books by author Harry Crews are in the database?
3. Determine the number of library cards assigned to each borrower phone number.
4. Find names of all borrowers who do not have any book loans.
5. Do any library branches have every book?