

COMP163

Database Management Systems

Lecture 7 – Sections 4.3, 5.1
SQL Queries

SQL

- **SQL** (Structured Query Language)
is the standard language for commercial DBMSs
- **SEQUEL** (Structured English QUery Language)
was originally defined by IBM for **SYSTEM R**
 - mid 1970s
 - unofficial pronunciation (**see**-kwuhl) still sometimes used
- SQL is more than a query language:
it includes a DDL, DML and admin commands

SQL commands

- Administration:
 - CREATE DATABASE
 - CREATE SCHEMA
 - SET ROLE
 - GRANT PRIVILEGES
- Data Definition:
 - CREATE TABLE
 - ALTER TABLE
 - DROP TABLE
 - CREATE VIEW
- Data Modification:
 - INSERT
 - DELETE
 - UPDATE
- Queries:
 - SELECT

48 commands listed in
SQL in a Nutshell

SQL Queries

- Queries in SQL are variations of the SELECT command
- Basic SQL queries correspond to the following relational algebra operations:
 - select σ
 - project π
 - cross product X
 - joins must be expressed as σ and X

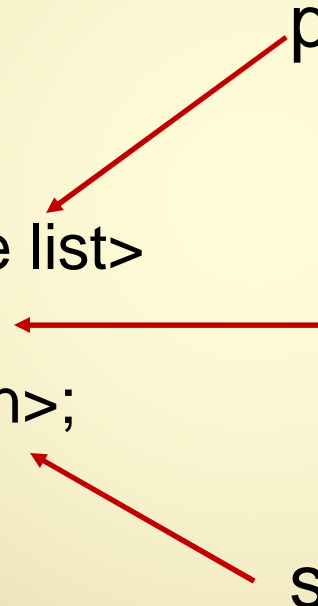
Basic SELECT Command

SELECT <attribute list>
FROM <table list>
WHERE <condition>;

project π

cross product X

select σ



Single Table Queries (σ and π)

$\pi_{Ssn}(\sigma_{Salary > 60000}(\text{EMPLOYEE}))$

```
SELECT Ssn  
FROM EMPLOYEE  
WHERE Salary > 60000;
```

$\pi_{City, State}(\sigma_{Airport_code = 'SFO'}(\text{AIRPORT}))$

```
SELECT City, State  
FROM AIRPORT  
WHERE Airport_code = 'SFO';
```

Join as Select & Cross

- In the basic SELECT/FROM/WHERE form, joins must be expressed as using σ and \times

$$\pi_{Lname, Dname} (EMPLOYEE \bowtie_{Ssn=Mgr_ssn} DEPARTMENT)$$

$$\pi_{Lname, Dname} (\sigma_{Ssn=Mgr_ssn} (EMPLOYEE \times DEPARTMENT))$$

```
SELECT Lname, Dname  
FROM EMPLOYEE, DEPARTMENT  
WHERE Ssn = Mgr_ssn;
```


Basic SQL Queries

- Retrieve the name and address of all employees who work for the 'Research' department.

```
SELECT  FNAME, LNAME, ADDRESS  
FROM    EMPLOYEE, DEPARTMENT  
WHERE   DNAME='Research' AND DNUMBER=DNO
```

selection
condition

join
condition

Basic SQL Queries

- For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```
SELECT  PNUMBER, DNUM, LNAME, BDATE, ADDRESS
```

```
FROM    PROJECT, DEPARTMENT, EMPLOYEE
```

```
WHERE   DNUM=DNUMBER AND MGRSSN=SSN
```

```
AND PLOCATION='Stafford'
```

join
PROJECT and
DEPARTMENT



select



join
DEPARTMENT
and EMPLOYEE



Tuple Variables (Aliases)

- We can give names to the tuples coming from each of the input relations

```
SELECT      E.Lname, D.Dname  
FROM        EMPLOYEE E, DEPARTMENT D  
WHERE       E.Ssn = D.Mgr_ssn;
```

- This can disambiguate common attribute names and improve readability

Renaming Attributes

- Attributes can also be renamed in the FROM clause
 - similar to alternate rename syntax in the algebra

SELECT	Fn, Ln
FROM	EMPLOYEE E(Fn, Mi, Ln, Bd, Ad, Sx, Sl, Sssn, Dn)
WHERE	Dn = 4;

Self Join

- For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

```
SELECT      E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM        EMPLOYEE E, EMPLOYEE S
WHERE       E.SUPERSSN=S.SSN
```

- Aliases are necessary for this query
- Think of E and S as two different copies of EMPLOYEE
 - E represents employees in role of *supervisees* and
S represents employees in role of *supervisors*

Aliases: alternate syntax

- Can also use the AS keyword to specify aliases

```
SELECT      E.FNAME, E.LNAME, S.FNAME, S.LNAME
FROM        EMPLOYEE AS E, EMPLOYEE AS S
WHERE       E.SUPERSSN=S.SSN
```

- Can also simply use the relation names
(when non-ambiguous)

```
SELECT      EMPLOYEE.Lname, DEPARTMENT.Dname
FROM        EMPLOYEE, DEPARTMENT
WHERE       EMPLOYEE.Ssn = DEPARTMENT.Mgr_ssn;
```

No $\sigma \rightarrow$ No WHERE

- If there are no selection (or join) conditions, the WHERE clause can be omitted

SELECT Ssn
FROM EMPLOYEE

π_{Ssn} EMPLOYEE

- Two or more relations in FROM clause with no join is a *CROSS PRODUCT*

SELECT Lname, Dname
FROM EMPLOYEE, DEPARTMENT

$\pi_{Lname, Dname} (EMPLOYEE \times DEPARTMENT)$

No $\pi \rightarrow *$

- To retrieve all the attribute values of the selected tuples, use *, which stands for *all the attributes*

```
SELECT      *  
FROM        EMPLOYEE  
WHERE       DNO=5
```

```
SELECT      *  
FROM        EMPLOYEE, DEPARTMENT  
WHERE       DNAME='Research' AND  
            DNO=DNUMBER
```


Tables as Sets → DISTINCT

- SQL does not treat a relation as a set; duplicate tuples can appear
- To eliminate duplicate tuples in a query result, the keyword **DISTINCT** is used

SELECT SALARY
FROM EMPLOYEE

may contain duplicates

SELECT **DISTINCT** SALARY
FROM EMPLOYEE

duplicates eliminated



Set Operations

- union operation (UNION)
intersection (INTERSECT)
set difference (MINUS, sometimes called EXCEPT)
 - some implementations of SQL do not support all set operations
- Set operation results **are sets** of tuples
duplicate tuples are eliminated from the result
- The set operations apply only to *union compatible relations*:
the two relations must have the same attributes and
the attributes must appear in the same order

Set Operations: Example

- List project numbers for all projects that involve an employee whose last name is 'Smith' as a worker or as a manager of the department that controls the project.

```
(SELECT      PNAME
FROM          PROJECT, DEPARTMENT, EMPLOYEE
WHERE        DNUM=DNUMBER AND MGRSSN=SSN AND LNAME='Smith')
```

UNION

```
(SELECT      PNAME
FROM          PROJECT, WORKS_ON, EMPLOYEE
WHERE        PNUMBER=PNO AND ESSN=SSN AND NAME='Smith')
```

Multiset Operations

- UNION ALL, INTERSECT ALL, EXCEPT ALL
- Multiset operation results are multisets of tuples
duplicate tuples are not eliminated

```
(SELECT      PNAME
FROM        PROJECT, DEPARTMENT, EMPLOYEE
WHERE       DNUM=DNUMBER AND MGRSSN=SSN AND LNAME='Smith')
      UNION ALL
(SELECT      PNAME
FROM        PROJECT, WORKS_ON, EMPLOYEE
WHERE       PNUMBER=PNO AND ESSN=SSN AND NAME='Smith')
```

WHERE Clause

- WHERE clause is a general boolean expression
- Boolean operators:
AND, OR, NOT
- Comparison operators:
=, <, <=, >, >=, <>
- String comparison operators:
LIKE
- Parentheses can be used to set precedence
- String literals can be enclosed in "... " or '...'


String Comparison

- The **LIKE** comparison operator is used to compare partial strings
- Two wildcard characters are used:
 - '%' replaces an arbitrary number of characters
 - '_' replaces a single arbitrary character

String Comparison Example

- Retrieve all employees whose address is in Houston, Texas.
- The value of the ADDRESS attribute must contain the substring “Houston, TX”.

```
SELECT    FNAME, LNAME  
FROM      EMPLOYEE  
WHERE     ADDRESS LIKE '%Houston, TX%'
```



zero or more
characters, before
and after substring

String Comparison Example

- Retrieve all employees who were born during the 1960s.
 - '6' must be the 3rd character of the 10 character date string

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       BDATE LIKE  '__6_____'
```

- Following would also work:

```
SELECT      FNAME, LNAME
FROM        EMPLOYEE
WHERE       BDATE LIKE  '__6 %'
```

assumes date format is YYYY-MM-DD

Arithmetic Operation

- The standard arithmetic operators '+', '-', '*', and '/' can be applied to numeric values in an SQL query result
- Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

```
SELECT      FNAME, LNAME, 1.1*SALARY
FROM        EMPLOYEE, WORKS_ON, PROJECT
WHERE       SSN=ESSN AND PNO=PNUMBER
            AND PNAME='ProductX'
```

Aggregate Functions

- Aggregate functions are applied to result attributes **COUNT, SUM, MAX, MIN, and AVG**
- Find the maximum salary, the minimum salary, and the average salary among all employees.

```
SELECT      MAX(Salary), MIN(Salary), AVG(Salary)
FROM        EMPLOYEE
```

- Find the total salary paid to employees who work for the 'Research' department.

```
SELECT      SUM(Salary)
FROM        EMPLOYEE, DEPARTMENT
WHERE       Dno=Dnumber AND Dname='Research'
```

Aggregate Functions

- Retrieve the total number of employees in the company and the number of employees in the Research' department.

```
SELECT      COUNT (*)  
FROM        EMPLOYEE
```

```
SELECT      COUNT (*)  
FROM        EMPLOYEE, DEPARTMENT  
WHERE       DNO=DNUMBER AND DNAME='Research'
```

Join as \bowtie and σ

```
mysql> SELECT * FROM r;
```

x	y
3	4
5	6
7	8
9	6

```
mysql> SELECT * FROM s;
```

a	b
2	3
4	7

```
mysql> SELECT * FROM r,s WHERE y=a;
```

x	y	a	b
3	4	4	7

$$R \bowtie_{y=a} S$$
$$\sigma_{y=a}(R \times S)$$

Explicit Join

- Joins can be explicitly stated in the FROM clause.

```
SELECT * FROM (r JOIN s ON y=a);
```

	x		y	
	a		b	
	3		4	
	4		7	

$R \bowtie_{y=a} S$

Left/Right Outer Join

```
SELECT * FROM (r LEFT JOIN s ON y=a);
```

x	y	a	b
3	4	4	7
5	6	NULL	NULL
7	8	NULL	NULL
9	6	NULL	NULL

```
SELECT * FROM (r RIGHT JOIN s ON y=a);
```

x	y	a	b
NULL	NULL	2	3
3	4	4	7

Full Outer Join

```
SELECT * FROM r FULL OUTER JOIN s ON y=a;
```

mysql doesn't support full outer join,
so we'll substitute an equivalent query:

```
(SELECT * FROM r LEFT JOIN s ON y=a) UNION  
(SELECT * FROM r RIGHT JOIN s ON y=a);
```

x	y	a	b
3	4	4	7
5	6	NULL	NULL
7	8	NULL	NULL
9	6	NULL	NULL
NULL	NULL	2	3

Ordering Results

- An ORDER BY clause can be added to order the result tuples

SELECT * FROM t;	SELECT * FROM t ORDER BY j;	SELECT * FROM t ORDER BY i;
+-----+-----+	+-----+-----+	+-----+-----+
i j	i j	i j
+-----+-----+	+-----+-----+	+-----+-----+
10 ten	11 eleven	4 four
11 eleven	4 four	10 ten
20 twenty	10 ten	11 eleven
4 four	20 twenty	20 twenty
+-----+-----+	+-----+-----+	+-----+-----+

ORDER BY Examples

- order by Lname first,
then by Fname if Lname is the same:

```
SELECT Lname, Fname  
FROM Employee  
WHERE salary > 60000  
ORDER BY Lname, Fname
```

- order by Lname in ascending order,
then by salary in descending order

```
SELECT Lname, salary  
FROM Employee  
WHERE salary > 60000  
ORDER BY Lname ASC, salary DESC
```

Grouping

- Forms groups (subsets) of result tuples before applying aggregate functions
- Example: count the number of employees in each department
(group employees by DNO, then count tuples in each group)

```
SELECT Dno, COUNT(*)  
FROM Employee  
GROUP BY Dno
```

$\text{DNO } \mathcal{F}_{\text{COUNT}} * (\text{EMPLOYEE})$

+-----+-----+	
Dno	COUNT
+-----+-----+	
8	120
22	238
7	82
20	169
+-----+-----+	

GROUP BY Example

- For each project, get the project name, project number and the number of employees working on that project

```
SELECT Pnumber, Pname, COUNT(*)  
FROM PROJECT, WORKS_ON  
WHERE Pnumber = Pno  
GROUP BY Pnumber, Pname
```

Attributes in SELECT clause must be aggregates
or must appear in the GROUP BY clause

Filtering Groups: HAVING

- We can throw away some groups by adding a condition in a HAVING clause
- example:
for each project *that has more than two employees*,
get the project name, project number and
the number of employees working on that project

```
SELECT Pnumber, Pname, COUNT(*)  
FROM PROJECT, WORKS_ON  
WHERE Pnumber = Pno  
GROUP BY Pnumber, Pname  
HAVING COUNT(*) > 2
```

GROUP BY Examples

SELECT * FROM e;

eid	salary	dept
E01	65000	ADMIN
E12	58400	ENGR
E08	76900	ENGR
E23	63800	ADMIN
E07	56900	ADMIN
E27	76400	ENGR
E14	48000	TEST

SELECT COUNT(*) FROM e
GROUP BY dept;

count(*)
3
3
1

SELECT dept, COUNT(*)
FROM e GROUP BY dept;

dept	count(*)
ADMIN	3
ENGR	3
TEST	1

GROUP BY Examples

```
SELECT dept, COUNT(*)  
FROM e  
GROUP BY dept  
HAVING COUNT(*) > 1;
```

dept	count(*)
ADMIN	3
ENGR	3

```
SELECT dept, AVG(salary)  
FROM e  
GROUP BY dept  
HAVING COUNT(*) > 1;
```

dept	AVG(salary)
ADMIN	61900
ENGR	70566.66667

Nested Queries

- Nested queries can be used as set values in the WHERE clause
- Set comparison operators
 - IN – set membership (“is in”, \in)
 - EXISTS – set not empty (\exists)
 - ALL – applies to all set members (\forall)
 - ANY – applies to any set member
 - CONTAINS – proper superset

Nested Queries

- find all employees who work on a project with John Smith

```
SELECT Lname, Fname
FROM EMPLOYEE E1, WORKS_ON W1
WHERE E1.SSN = W1.ESSN
      AND W1.Pno IN (SELECT Pno
                      FROM EMPLOYEE E2, WORKS_ON W2
                      WHERE E2.SSN = W2.ESSN
                        AND E2.Fname = "John"
                        AND E2.Lname = "Smith")
```

Nested Queries

- find the highest paid employee in department 5

```
SELECT Lname, Fname
FROM EMPLOYEE E1
WHERE E1.Dno=5
      AND E1.Salary >= ALL (SELECT E2.Salary
                             FROM EMPLOYEE E2
                             WHERE E2.Dno=5)
```

Nested Queries

- List names of managers who have dependents

```
SELECT Lname, Fname
FROM EMPLOYEE E1
WHERE EXISTS (SELECT *
              FROM DEPENDENT D1
              WHERE E1.Ssn = D1.Essn)
AND
EXISTS (SELECT *
       FROM DEPARTMENT D2
       WHERE E1.Ssn = D2.Mgr_ssn)
```

This is an example of a *correlated nested query*,
since the nested queries refer to the relations in the outer query.

Nested Queries


- List names of employees who work on all projects controlled by department 5

```
SELECT Lname, Fname
FROM EMPLOYEE E
WHERE (SELECT W.Pno
        FROM WORKS_ON W
        WHERE E.Ssn = W.Essn)
CONTAINS
(SELECT P.Pnumber
    FROM PROJECT P
    WHERE P.Dnum=5)
```

Nested Queries

- List names of all projects controlled by department 5 or department 7

```
SELECT P.Pname  
FROM PROJECT P  
WHERE P.Dnum IN (5,7)
```



explicit set of values