

# COMP163

Database Management Systems

**Lecture 09 – Sections 5.2, 5.3, 26.1**  
**Views, Semantic Constraints**  
**and Triggers**

# Views

# Views in SQL

- A ***view*** is a ***virtual table***  
that is derived from other tables
  - Virtual → computed, not stored
- Allows full query operations
  - Behaves as a table when queried
- Allows for limited update operations
  - Updates must be mapped down to tables that contribute to the view

# Specification of Views

- SQL command: **CREATE VIEW**
  - a table (view) name
  - a possible list of attribute names  
(for example, when arithmetic operations are specified or when we want the names to be different from the attributes in the base relations)
  - a query to specify the table contents

# SQL Views: An Example

- Specify a different WORKS\_ON table

```
CREATE VIEW WORKS_ON_BYNAME AS  
  SELECT FNAME, LNAME, PNAME, HOURS  
  FROM EMPLOYEE, PROJECT, WORKS_ON  
  WHERE SSN=ESSN AND PNO=PNUMBER;
```

# Using a Virtual Table

- We can specify SQL queries on a newly created table (view):

```
SELECT FNAME, LNAME  
      FROM WORKS_ON_BYNAME  
      WHERE PNAME='Seena';
```

- When no longer needed, a view can be dropped:

```
DROP WORKS_ON_NEW;
```

# Efficient View Implementation

- Query modification:
  - Represent the view query in terms of a query on the underlying base tables
- Disadvantage:
  - Inefficient for views defined via complex queries
    - Especially if additional queries are to be applied to the view within a short time period



# Efficient View Implementation

- View materialization:
  - Involves physically creating and keeping a temporary table
- Assumption:
  - Other queries on the view will follow
- Concerns:
  - Maintaining correspondence between the base table and the view when the base table is updated
- Strategy:
  - Incremental update



# Updating Views

- Update on a view of a single table without aggregate operations:
  - Update may map to an update on the underlying base table
- Views involving joins:
  - An update *may* map to an update on the underlying base relations
    - (This is not always possible)

# Non-updatable Views

- Views defined using GROUP BY and aggregate functions are not updateable
- Views defined on multiple tables using joins are generally not updateable
- View on tables that have CHECK constraints: the CHECK must also be added to the definition of a view if the view is to be updated
  - To allow check for updatability and to plan for an execution strategy

# Constraints

# Previous Constraints

```
CREATE TABLE Teaching (  
  ProfId   INTEGER,  
  CrsCode CHAR (6),  
  Semester CHAR (6),  
  PRIMARY KEY (CrsCode, Semester),
```

```
  FOREIGN KEY (ProfId) REFERENCES Professor (Id)  
    ON DELETE NO ACTION  
    ON UPDATE CASCADE,
```

```
  FOREIGN KEY (CrsCode) REFERENCES Course (CrsCode)  
    ON DELETE SET NULL  
    ON UPDATE CASCADE )
```

# Table Semantic Constraints

- Used for application dependent conditions
- *Example:* limit attribute values

```
CREATE TABLE Transcript (  
    StudId INTEGER,  
    CrsCode CHAR(6),  
    Semester CHAR(6),  
    Grade CHAR(1),  
    CHECK (Grade IN ('A', 'B', 'C', 'D', 'F')),  
    CHECK (StudId > 0 AND StudId < 1000000000) )
```

- Each row in table must satisfy condition

# User-Defined Domains

- Possible attribute values can be specified
  - Using a CHECK constraint, or
  - Creating a new domain
- Domain can be used in several declarations
- Domain is a schema element

```
CREATE DOMAIN Grades CHAR (1)
    CHECK (VALUE IN ('A', 'B', 'C', 'D', 'F'))
CREATE TABLE Transcript (
    ...,
    Grade: Grades,
    ... )
```

# Table Constraint Example

- Ensure that managers are paid more than their employees.

```
CREATE TABLE Employee (  
    Id INTEGER,  
    Name CHAR(20),  
    Salary INTEGER,  
    MngrSalary INTEGER,  
    CHECK ( MngrSalary > Salary )
```



# Constraints – Problems

- **Problem 1:**

An empty table always satisfies all CHECK constraints (an idiosyncrasy of the SQL standard)

```
CREATE TABLE Employee (  
    Id INTEGER,  
    Name CHAR(20),  
    Salary INTEGER,  
    MngrSalary INTEGER,  
    CHECK ( 0 < (SELECT COUNT (*) FROM Employee)) )
```

- If Employee is empty, there are no rows on which to evaluate the CHECK condition.

# Constraints – Problems

- **Problem 2:**

Inter-relational constraints should be symmetric

```
CREATE TABLE Employee (  
  Id INTEGER,  
  Name CHAR(20),  
  Salary INTEGER,  
  MngrSalary INTEGER,  
  CHECK ((SELECT COUNT (*) FROM Manager) <  
         (SELECT COUNT (*) FROM Employee)) )
```

- Why should constraint be in Employee, rather than Manager?
- What if Employee is empty?

# Assertions

- Assertions are schema elements
- Symmetrically specifies an inter-relational constraint
- Applies to entire database (not just the individual rows of a single table)
  - hence it works even if Employee is empty

```
CREATE ASSERTION DontFireEveryone  
CHECK (0 < SELECT COUNT (*) FROM Employee)
```

# Designing Assertions

- Specify a query that violates the condition; include that inside a NOT EXISTS clause
- Query result must be empty
  - if the query result is not empty, the assertion has been violated

```
CREATE ASSERTION KeepEmployeeSalariesDown  
CHECK (NOT EXISTS(  
    SELECT * FROM Employee E  
    WHERE E.Salary > E.MngrSalary))
```

# Assertion Example

- The salary of an employee must not be greater than the salary of the manager of the department that the employee works for.

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK (NOT EXISTS
      (SELECT *
        FROM EMPLOYEE E, EMPLOYEE M,
             DEPARTMENT D
        WHERE E.SALARY > M.SALARY AND
              E.DNO=D.NUMBER AND
              D.MGRSSN=M.SSN) )
```

# Assertion Example

```
CREATE ASSERTION NoEmptyCourses
CHECK (NOT EXISTS (
    SELECT * FROM Teaching T
    WHERE NOT EXISTS (
        SELECT * FROM Transcript R
        WHERE T.CrsCode = R.CrsCode
        AND T.Semester = R.Semester)
    ))
```

note the double negative logic:

It is **not true** that there are courses taught that **do not** have students.

# Triggers



# Triggers

- Triggers are active statements that specify responses to specific conditions.
- A trigger is a schema element

```
CREATE TRIGGER CrsChange  
  AFTER UPDATE OF CrsCode, Semester ON Transcript  
  WHEN (Grade IS NOT NULL)  
  ROLLBACK
```

# Trigger Overview

- Element of the database schema
- General form:
  - ON <event> IF <condition> THEN <action>
  - *Event*- request to execute database operation
  - *Condition* - predicate evaluated on database state
  - *Action* – execution of procedure that might involve database updates
- Example:

ON updating maximum course enrollment

IF number registered > new max enrollment limit  
THEN deregister students using LIFO policy

# Possible Trigger Semantics

- **Activation** - Occurrence of the event
- **Consideration** - The point, after activation, when the *condition* is evaluated
  - **Immediate:** evaluate condition as soon as the event occurs
  - **Deferred:** wait to evaluate the condition at the end of the transaction
  - The *condition* may refer to the database state both before and after the triggering event

# Possible Trigger Semantics

- **Execution** – point at which the *action* occurs
  - With deferred consideration, execution is also deferred
  - With immediate consideration, execution can occur immediately after consideration or it can be deferred
    - If execution is immediate, execution can occur before, after, or instead of triggering event.
    - Before triggers adapt naturally to maintaining integrity constraints: violation results in rejection of event.

# Possible Trigger Semantics

- **Granularity**

- *Row-level granularity*: change of a single row is an event (a single UPDATE statement might result in multiple events)
- *Statement-level granularity*: events are statements (a single UPDATE statement that changes multiple rows is a single event).

# Possible Trigger Semantics

- **Multiple Triggers**

- How should multiple triggers activated by a single event be handled?
  - Evaluate one condition at a time and if true immediately execute action or
  - Evaluate all conditions, then execute actions
- The execution of an action can affect the truth of a subsequently evaluated condition so the choice is significant.



# Triggers in SQL

- **Events:**  
INSERT, DELETE, or UPDATE statements or changes to individual rows caused by these statements
- **Condition:**  
Anything that is allowed in a WHERE clause
- **Action:**  
An individual SQL statement or a program written in the language of Procedural Stored Modules (PSM)  
(which may contain embedded SQL statements)



# Triggers in SQL

- **Consideration:** *Immediate*
  - Condition can refer to the state of the affected row or table before *and* after the event occurs
- **Execution:** *Immediate* – can be before or after the execution of the triggering event
  - Action of a before trigger cannot modify the database
- **Granularity:** Both *row-level* and *statement-level*

# Trigger Syntax

```
CREATE TRIGGER trigger-name
  { BEFORE | AFTER }
  { INSERT | DELETE | UPDATE [OF column-name-list] }
  ON table-name
  [ REFERENCING { OLD AS old-tuple-name |
                  NEW AS new-tuple-name |
                  OLD TABLE AS old-table-name |
                  NEW TABLE AS new-table-name } ]
  [ FOR EACH { ROW | STATEMENT } ]
  [ WHEN (precondition) ]
  statement-list
```

# Before Trigger (row granularity)

```
CREATE TRIGGER Max_EnrollCheck
  BEFORE INSERT ON Transcript
    REFERENCEING NEW AS N  --row to be added
  FOR EACH ROW
  WHEN
    ((SELECT COUNT (T.StudId) FROM Transcript T
      WHERE T.CrsCode = N.CrsCode
        AND T.Semester = N.Semester)
    >=
    (SELECT C.MaxEnroll FROM Course C
      WHERE C.CrsCode = N.CrsCode ))
  ABORT TRANSACTION
```

Check that  
enrollment  $\leq$  limit

# After Trigger (row granularity)

```
CREATE TRIGGER LimitSalaryRaise
AFTER UPDATE OF Salary ON Employee
REFERENCING OLD AS O
              NEW AS N
FOR EACH ROW
WHEN (N.Salary - O.Salary > 0.05 * O.Salary)
  UPDATE Employee
  SET Salary = 1.05 * O.Salary
  WHERE Id = O.Id
```

No salary raises  
greater than 5%

Note: The action itself is a triggering event (but in this case a chain reaction is not possible)

# After Trigger (stmt granularity)

```
CREATE TRIGGER RecordNewAverage  
AFTER UPDATE OF Salary ON Employee  
FOR EACH STATEMENT  
  INSERT INTO Log  
  VALUES (CURRENT_DATE,  
           SELECT AVG (Salary)  
           FROM Employee)
```

Keep track of salary  
averages in the log

# Trigger Example

- A trigger to compare an employee's salary to his/her supervisor during insert or update operations:

```
CREATE TRIGGER INFORM_SUPERVISOR
BEFORE INSERT OR UPDATE OF
    SALARY, SUPERVISOR_SSN ON EMPLOYEE
FOR EACH ROW
WHEN
    (NEW.SALARY > (SELECT SALARY FROM EMPLOYEE
                    WHERE SSN=NEW.SUPERVISOR_SSN))
INFORM_SUPERVISOR(NEW.SUPERVISOR_SSN,NEW.SSN);
```

INFORM\_SUPERVISOR is a *stored procedure*.